
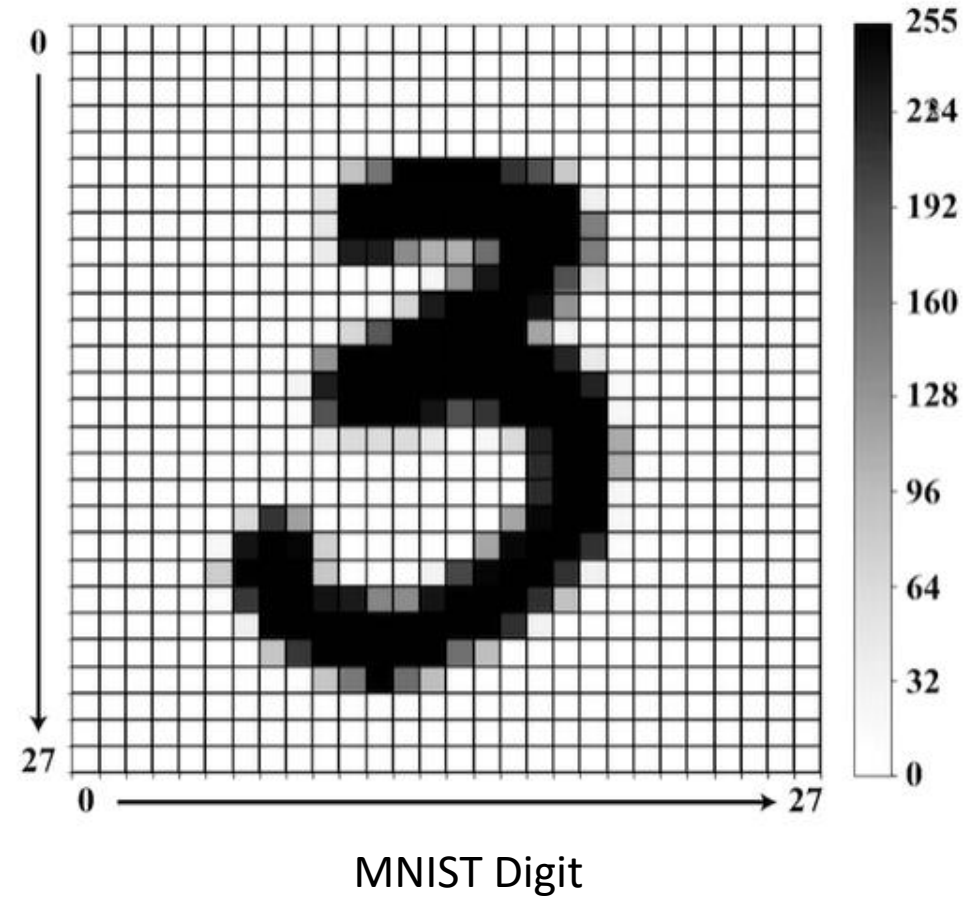


Neural Networks


Sample MNIST Digits



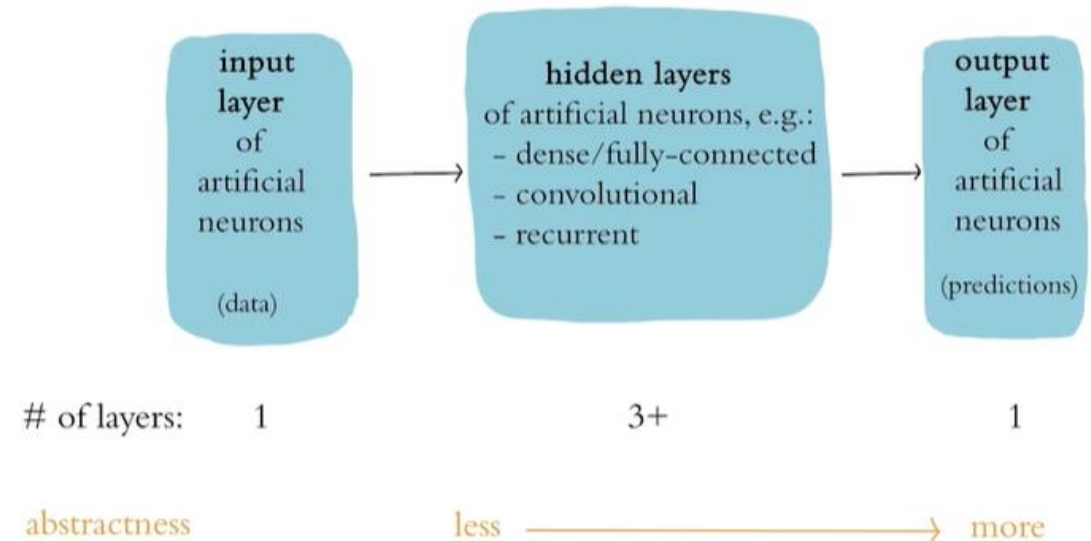
MNIST is Modified National Institute of Standards and Technology database




Shallow Neural Net that we will be building soon




A Generalized Deep Learning Model

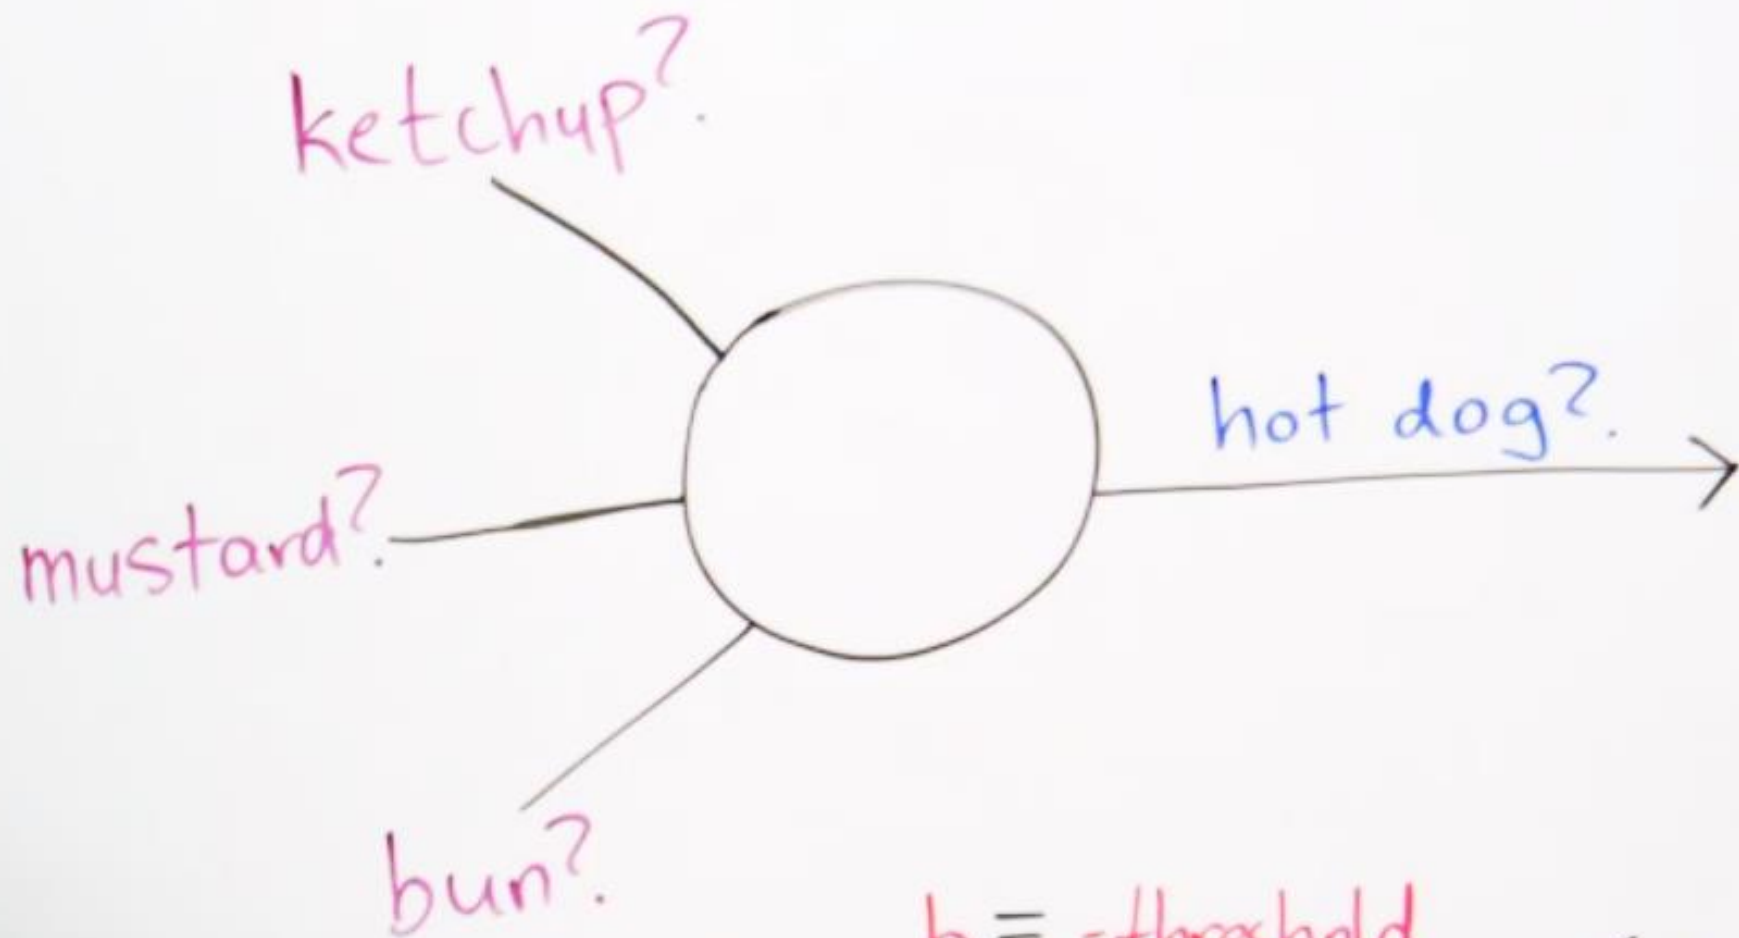


A Biological Neuron



Frank Rosenblatt – The Perceptron

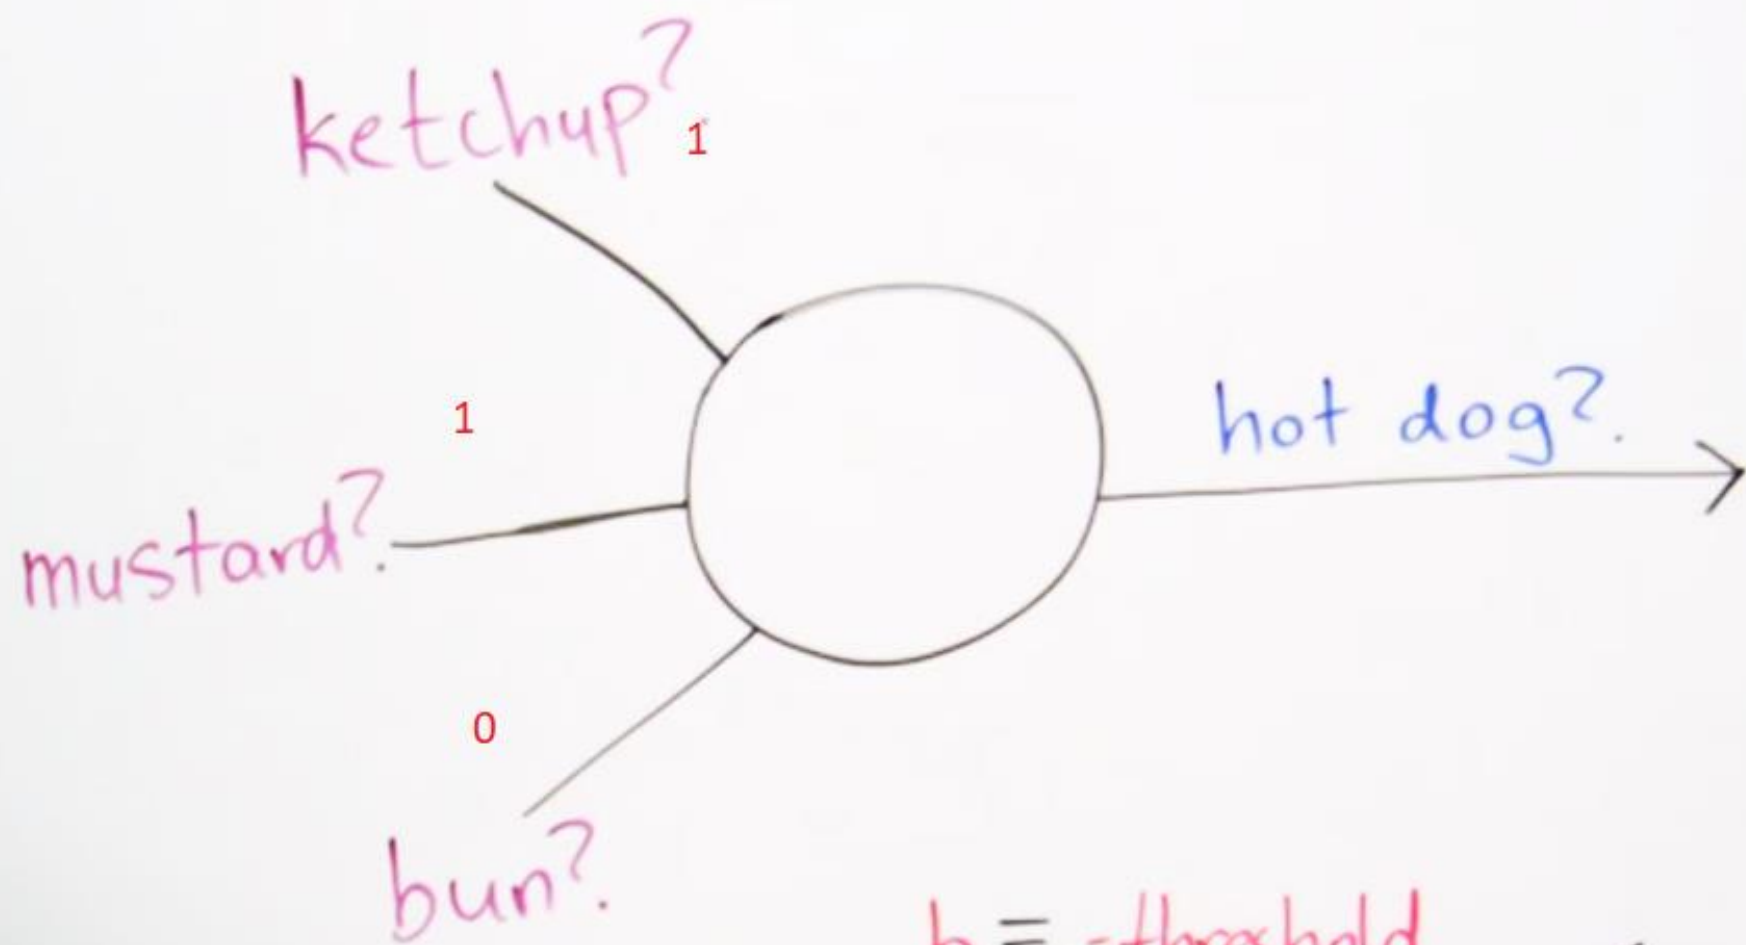




$$\sum_{i=1}^n w_i x_i = w \cdot x$$

$b \equiv$ -threshold

$$\text{output} \begin{cases} 1 & \text{if } w \cdot x + b > 0 \\ 0 & \text{otherwise} \end{cases}$$

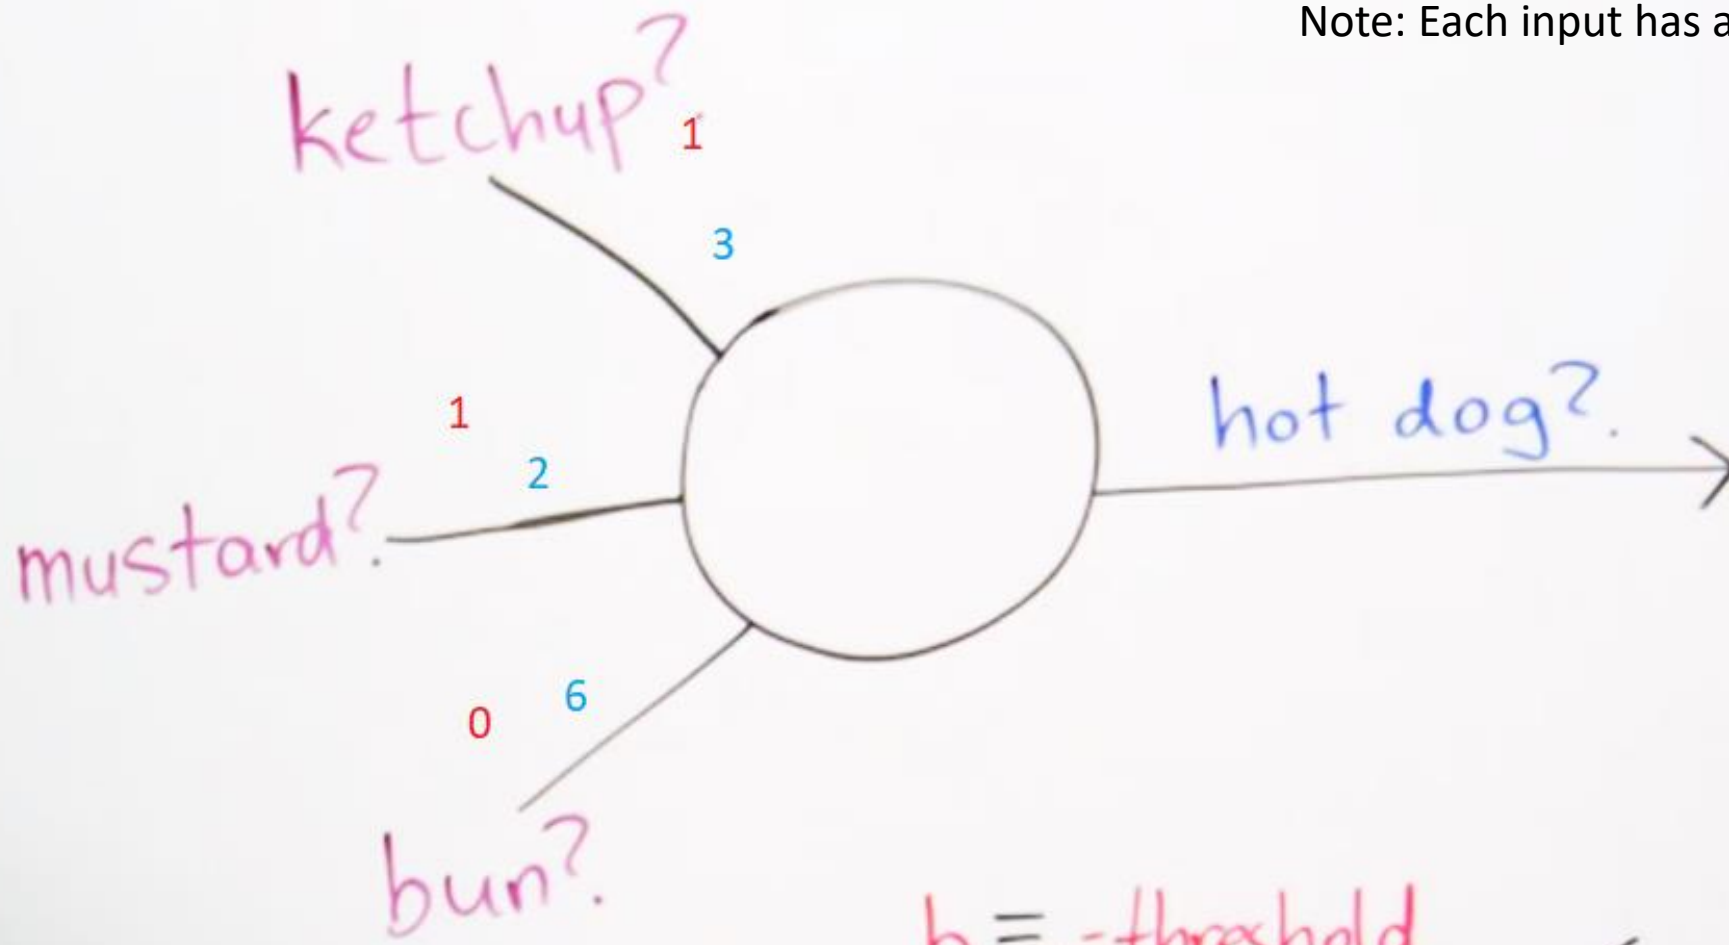


$$\sum_{i=1}^n w_i x_i = w \cdot x$$

$b \equiv$ -threshold

$$\text{output} \begin{cases} 1 & \text{if } w \cdot x + b > 0 \\ 0 & \text{otherwise} \end{cases}$$

Note: Each input has a particular weight



$$\sum_{i=1}^n w_i x_i = w \cdot x$$

$b \equiv$ -threshold

$$\text{output} \begin{cases} 1 & \text{if } w \cdot x + b > 0 \\ 0 & \text{otherwise} \end{cases}$$

ketchup?

$\left. \begin{matrix} 3 \\ 2 \\ 0 \end{matrix} \right\} 5$

3

1

2

mustard?

0

6

bun?

hot dog?

$b \equiv$ -threshold

$$\sum_{i=1}^n$$

$$w_i x_i = w \cdot x$$

$$\text{output} \begin{cases} 1 & \text{if } w \cdot x + b > 0 \\ 0 & \text{otherwise} \end{cases}$$

ketchup?

1

$$\left. \begin{matrix} 3 \\ 2 \\ 0 \end{matrix} \right\} 5 > 4$$

3

1

2

4

hot dog?

0

6

bun?

$b \equiv -\text{threshold}$

$$\sum_{i=1}^n$$

$$w_i x_i = w \cdot x$$

$$\text{output} \begin{cases} 1 & \text{if } w \cdot x + b > 0 \\ 0 & \text{otherwise} \end{cases}$$

ketchup?

1

$$\left. \begin{matrix} 3 \\ 2 \\ 0 \end{matrix} \right\} 5 > 4$$

3

4

hot dog?

1

mustard?

1

2

0

6

bun?

$b \equiv -\text{threshold}$

$$\sum_{i=1}^n$$

$$w_i x_i = w \cdot x$$

$$\text{output} \begin{cases} 1 & \text{if } w \cdot x + b > 0 \\ 0 & \text{otherwise} \end{cases}$$

ketchup?

0

$$\left. \begin{matrix} 0 \\ 2 \\ 0 \end{matrix} \right\} 5 > 4$$

3

1

2

mustard?

4

hot dog?

1

0

6


bun?

$b \equiv$ -threshold

$$\sum_{i=1}^n$$

$$w_i x_i = w \cdot x$$

$$\text{output} \begin{cases} 1 & \text{if } w \cdot x + b > 0 \\ 0 & \text{otherwise} \end{cases}$$



ketchup?

0

$$\left. \begin{matrix} 0 \\ 2 \\ 0 \end{matrix} \right\} 2 < 4$$

3

1

2

mustard?

4

hot dog?

0

0

6

bun?

$b \equiv$ -threshold

$$\sum_{i=1}^n$$

$$w_i x_i = w \cdot x$$

$$\text{output} \begin{cases} 1 & \text{if } w \cdot x + b > 0 \\ 0 & \text{otherwise} \end{cases}$$

ketchup?

$$\left. \begin{matrix} 0 \\ 0 \\ 6 \end{matrix} \right\} 6 > 4$$

3

4

hot dog?

1

mustard?

0

2

1

6


bun?

$b \equiv$ -threshold

$$\sum_{i=1}^n$$

$$w_i x_i = w \cdot x$$

$$\text{output} \begin{cases} 1 & \text{if } w \cdot x + b > 0 \\ 0 & \text{otherwise} \end{cases}$$



Design neural network architecture


```
In [17]: ▶ model = Sequential()  
          model.add(Dense(64, activation='sigmoid', input_shape=(784,)))  
          model.add(Dense(10, activation='softmax'))
```


```
In [18]: ▶ model.summary()
```

Model: "sequential"

| Layer (type) | Output Shape | Param # |
|-----------------|--------------|---------|
| dense (Dense) | (None, 64) | 50240 |
| dense_1 (Dense) | (None, 10) | 650 |

=====
Total params: 50,890
Trainable params: 50,890
Non-trainable params: 0
=====





ketchup?

$$\begin{Bmatrix} 0 \\ 2 \\ 0 \end{Bmatrix} 2 < 4$$

3

1

2

4

hot dog?

0

mustard?

0

6

bun?

$b \equiv$ -threshold

output $\begin{cases} 1 & \text{if } w \cdot x + b > 0 \\ 0 & \text{otherwise} \end{cases}$

$$\sum_{i=1}^n$$

$$w_i x_i = w \cdot x$$

$$\text{sigmoid } \sigma(z) = \frac{1}{1 + e^{-z}}$$

Design neural network architecture


```
In [17]: ▶ model = Sequential()  
          model.add(Dense(64, activation='sigmoid', input_shape=(784,)))  
          model.add(Dense(10, activation='softmax'))
```


```
In [18]: ▶ model.summary()
```


Model: "sequential"


| Layer (type) | Output Shape | Param # |
|-----------------|--------------|---------|
| dense (Dense) | (None, 64) | 50240 |
| dense_1 (Dense) | (None, 10) | 650 |

=====
Total params: 50,890
Trainable params: 50,890
Non-trainable params: 0
=====









ketchup?

$$\left. \begin{matrix} 0 \\ 0 \\ 6 \end{matrix} \right\} 6 > 4$$

3

4

hot dog?

1

mustard?

0

2

1

6

bun?




$b \equiv$ -threshold

$$\text{output} \begin{cases} 1 & \text{if } w \cdot x + b > 0 \\ 0 & \text{otherwise} \end{cases}$$

$$\sum_{i=1}^n w_i x_i = w \cdot x$$

$$\text{sigmoid } \sigma(z) = \frac{1}{1 + e^{-z}}$$

$$\text{ReLU } \sigma(z) = \max(0, z)$$



Key Concepts

- parameters:
 - weight w
 - bias b
- activation a
- neurons:
 - sigmoid
 - tanh
 - ReLU

input

28 x 28 = 784



hidden

64 sigmoid neurons



output

10 softmax neurons


Design neural network architecture

```
In [17]: model = Sequential()  
          model.add(Dense(64, activation='sigmoid', input_shape=(784,)))  
          model.add(Dense(10, activation='softmax'))
```

```
In [18]: model.summary()
```


Model: "sequential"


| Layer (type) | Output Shape | Param # |
|--------------------------|--------------|---------|
| dense (Dense) | (None, 64) | 50240 |
| dense_1 (Dense) | (None, 10) | 650 |
| Total params: 50,890 | | |
| Trainable params: 50,890 | | |
| Non-trainable params: 0 | | |



Hidden Layers


- Dense Layer/ Fully Connected Layer







| | | | | |
|--------------|---|---|---|---|
| layer | 1 | 2 | 3 | 4 |
| hidden layer | | 1 | 2 | |


forward propagation








$$w \cdot x + b$$



$$w \cdot x + b$$

forward prop.



layer 1 2 3 4
hidden layer 1 2

forward propagation

Design neural network architecture


```
In [17]: ▶ model = Sequential()  
model.add(Dense(64, activation='sigmoid', input_shape=(784,)))  
model.add(Dense(10, activation='softmax'))
```

```
In [18]: ▶ model.summary()
```

Model: "sequential"

| Layer (type) | Output Shape | Param # |
|-----------------|--------------|---------|
| dense (Dense) | (None, 64) | 50240 |
| dense_1 (Dense) | (None, 10) | 650 |

=====
Total params: 50,890
Trainable params: 50,890
Non-trainable params: 0
=====



Design neural network architecture

```
In [17]: model = Sequential()  
model.add(Dense(64, activation='sigmoid', input_shape=(784,)))  
model.add(Dense(10, activation='softmax'))
```

```
In [18]: model.summary()
```

Model: "sequential"

| Layer (type) | Output Shape | Param # |
|--------------------------|--------------|---------|
| ===== | ===== | ===== |
| dense (Dense) | (None, 64) | 50240 |
| dense_1 (Dense) | (None, 10) | 650 |
| ===== | ===== | ===== |
| Total params: 50,890 | | |
| Trainable params: 50,890 | | |
| Non-trainable params: 0 | | |

input

28 x 28 = 784




hidden

64 sigmoid neurons



output

10 softmax neurons



layer 1 2 3 4
hidden layer 1 2
forward propagation

Design neural network architecture

```
In [17]: model = Sequential()  
model.add(Dense(64, activation='sigmoid', input_shape=(784,)))  
model.add(Dense(10, activation='softmax'))
```

```
In [18]: model.summary()
```

Model: "sequential"

| Layer (type) | Output Shape | Param # |
|--------------------------|--------------|---------|
| ===== | ===== | ===== |
| dense (Dense) | (None, 64) | 50240 |
| dense_1 (Dense) | (None, 10) | 650 |
| ===== | ===== | ===== |
| Total params: 50,890 | | |
| Trainable params: 50,890 | | |
| Non-trainable params: 0 | | |

input

28 x 28 = 784



hidden

64 sigmoid neurons




output

10 softmax neurons

Key Concepts

- parameters:
 - weight w
 - bias b
- activation a
- neurons:
 - sigmoid
 - tanh
 - ReLU
- layer types:
 - dense/FC
 - softmax
- input layer
- hidden layer
- output layer
- forward prop.






$$C = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

$$w \cdot x + b$$



forward prop.

Neuron Saturation







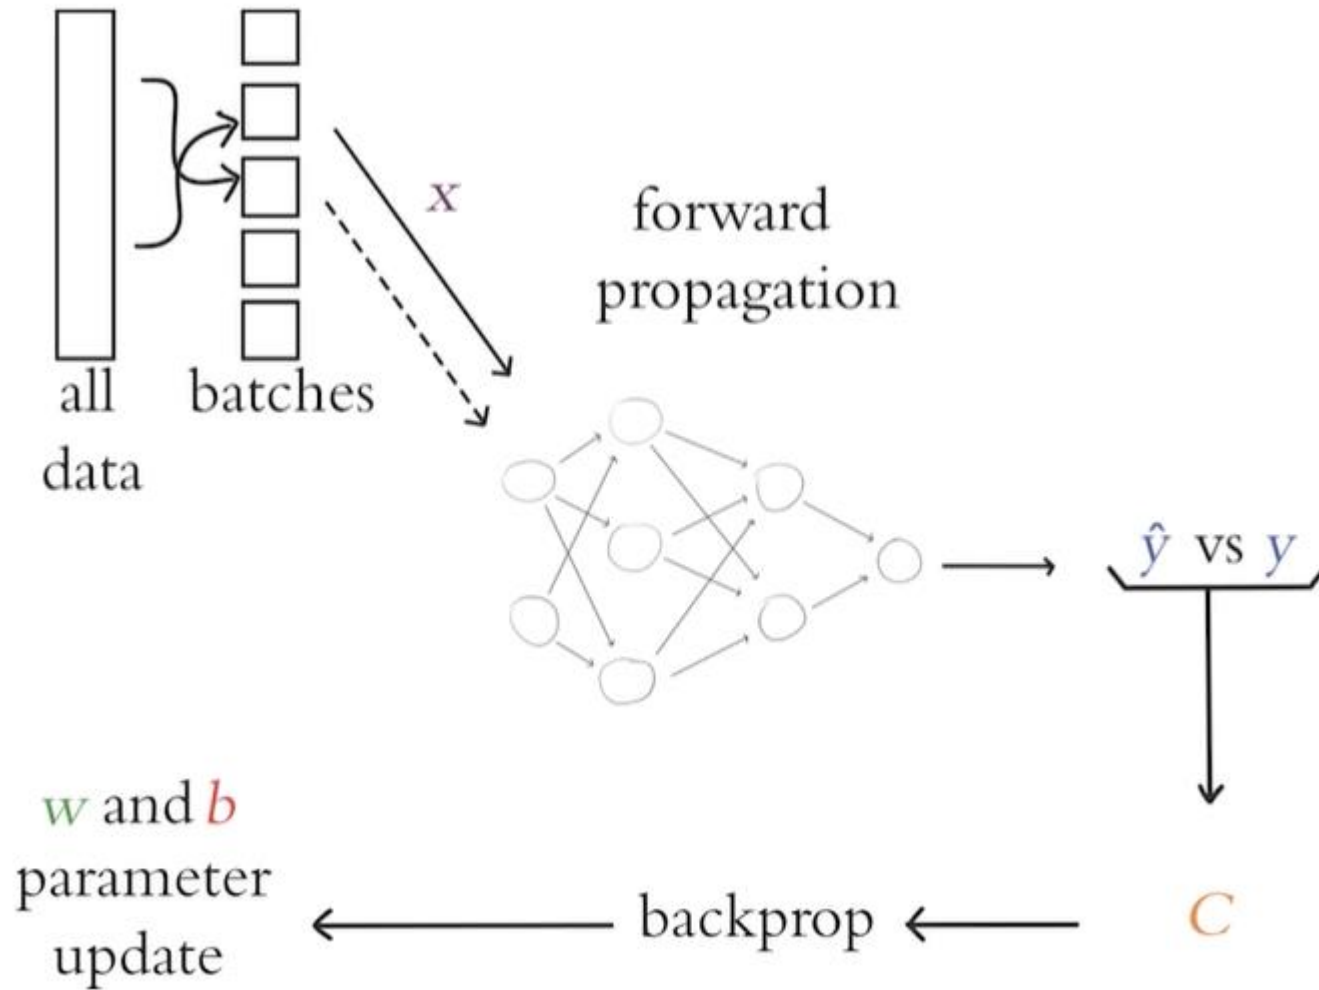
Gradient Descent



Learning Rate



Stochastic Gradient Descent



Configure model

```
In [22]: model.compile(loss='mean_squared_error', optimizer=SGD(lr=0.01), metrics=['accuracy'])
```

Train!



```
In [23]: model.fit(X_train, y_train, batch_size=128, epochs=200, verbose=1, validation_data=(X_valid, y_valid))
```

```
Epoch 1/200
469/469 [=====] - 14s 15ms/step - loss: 0.0921 - accuracy: 0.1042 - val_loss: 0.0913 - val_accu
acy: 0.0996
Epoch 2/200
469/469 [=====] - 5s 11ms/step - loss: 0.0911 - accuracy: 0.1008 - val_loss: 0.0906 - val_accu
cy: 0.0990
Epoch 3/200
469/469 [=====] - 6s 12ms/step - loss: 0.0905 - accuracy: 0.1018 - val_loss: 0.0901 - val_accu
cy: 0.1043
Epoch 4/200
469/469 [=====] - 6s 12ms/step - loss: 0.0900 - accuracy: 0.1094 - val_loss: 0.0897 - val_accu
cy: 0.1204
Epoch 5/200
469/469 [=====] - 6s 12ms/step - loss: 0.0896 - accuracy: 0.1248 - val_loss: 0.0893 - val_accu
cy: 0.1433
Epoch 6/200
469/469 [=====] - 6s 12ms/step - loss: 0.0892 - accuracy: 0.1473 - val_loss: 0.0890 - val_accu
cy: 0.1728
Epoch 7/200
469/469 [=====] - 6s 12ms/step - loss: 0.0889 - accuracy: 0.1682 - val_loss: 0.0887 - val_accu
cy: 0.1933
```

$$60000/128 = 468.75$$





Round of Training:

1. Sample a mini-batch of x values
2. Forward propagate x through network to estimate y with \hat{y}
3. Calculate cost C by comparing y and \hat{y}
4. Descend gradient of C to adjust w and b , enabling x to better predict y





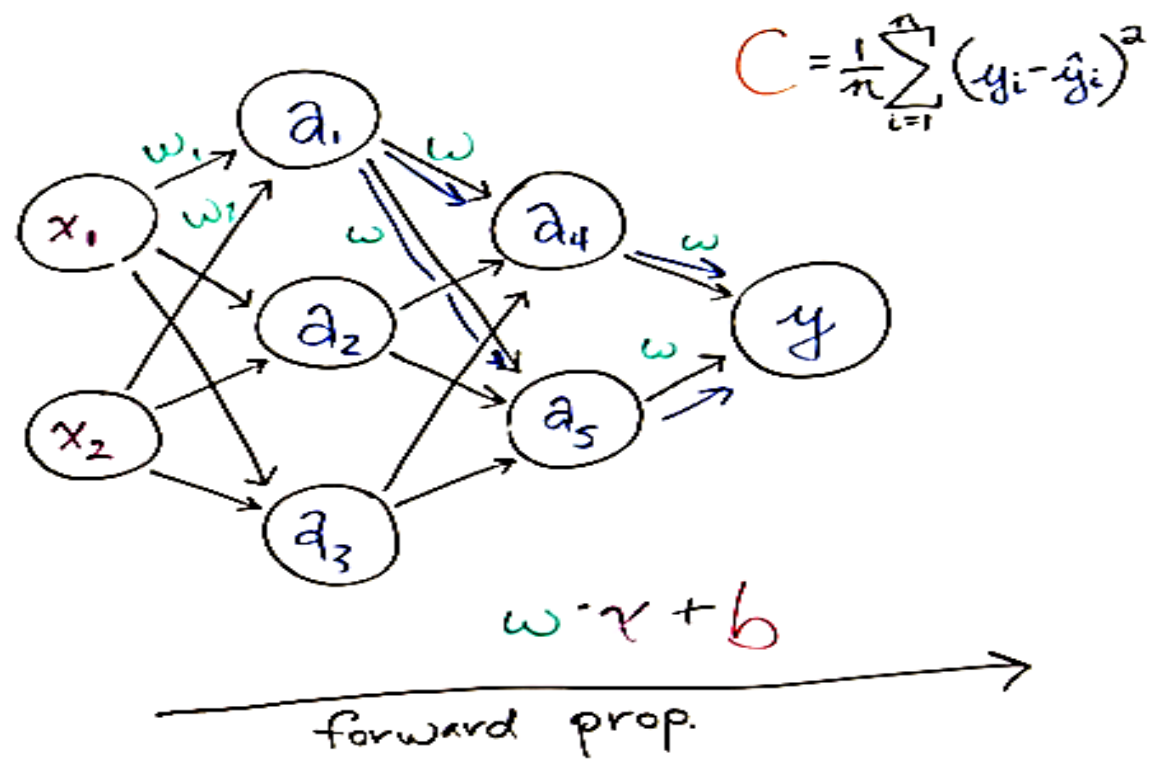
Escaping the local minimum of cost

To avoid getting trapped in local minima, a batch size of larger than 128 is not recommended.




Learning Speed by layer





Key Concepts

- parameters:
 - weight w
 - bias b
 - activation a
 - neurons:
 - sigmoid
 - tanh
 - ReLU
 - input layer
 - hidden layer
 - output layer
 - forward/backprop.
- layer types:
 - dense/FC
 - softmax
 - cost function C
 - quadratic
 - cross-entropy
 - optimizers:
 - SGD
 - optimizer hyperparams:
 - learning rate η
 - batch size



Train!

```
In [23]: model.fit(X_train, y_train, batch_size=128, epochs=200, verbose=1, validation_data=(X_valid, y_valid))
```

| Epoch | 469/469 | Time | Step | Loss | Accuracy | Val Loss | Val Accuracy |
|-------------|---------|------|-----------|--------|----------|----------|--------------|
| Epoch 1/200 | [=====] | 14s | 15ms/step | 0.0921 | 0.1042 | 0.0913 | 0.0996 |
| Epoch 2/200 | [=====] | 5s | 11ms/step | 0.0911 | 0.1008 | 0.0906 | 0.0990 |
| Epoch 3/200 | [=====] | 6s | 12ms/step | 0.0905 | 0.1018 | 0.0901 | 0.1043 |
| Epoch 4/200 | [=====] | 6s | 12ms/step | 0.0900 | 0.1094 | 0.0897 | 0.1204 |
| Epoch 5/200 | [=====] | 6s | 12ms/step | 0.0896 | 0.1248 | 0.0893 | 0.1433 |
| Epoch 6/200 | [=====] | 6s | 12ms/step | 0.0892 | 0.1473 | 0.0890 | 0.1728 |
| Epoch 7/200 | [=====] | 6s | 12ms/step | 0.0889 | 0.1682 | 0.0887 | 0.1933 |

Key Concepts

- parameters:

- weight w
- bias b

- activation a

- neurons:

- sigmoid
- tanh
- ReLU

- input layer

- hidden layer

- output layer

- forward/backprop.

- layer types:

- dense/FC
- softmax

- cost function C

- quadratic
- cross-entropy

- optimizers:

- SGD

- optimizer hyperparams:

- learning rate η
- batch size

