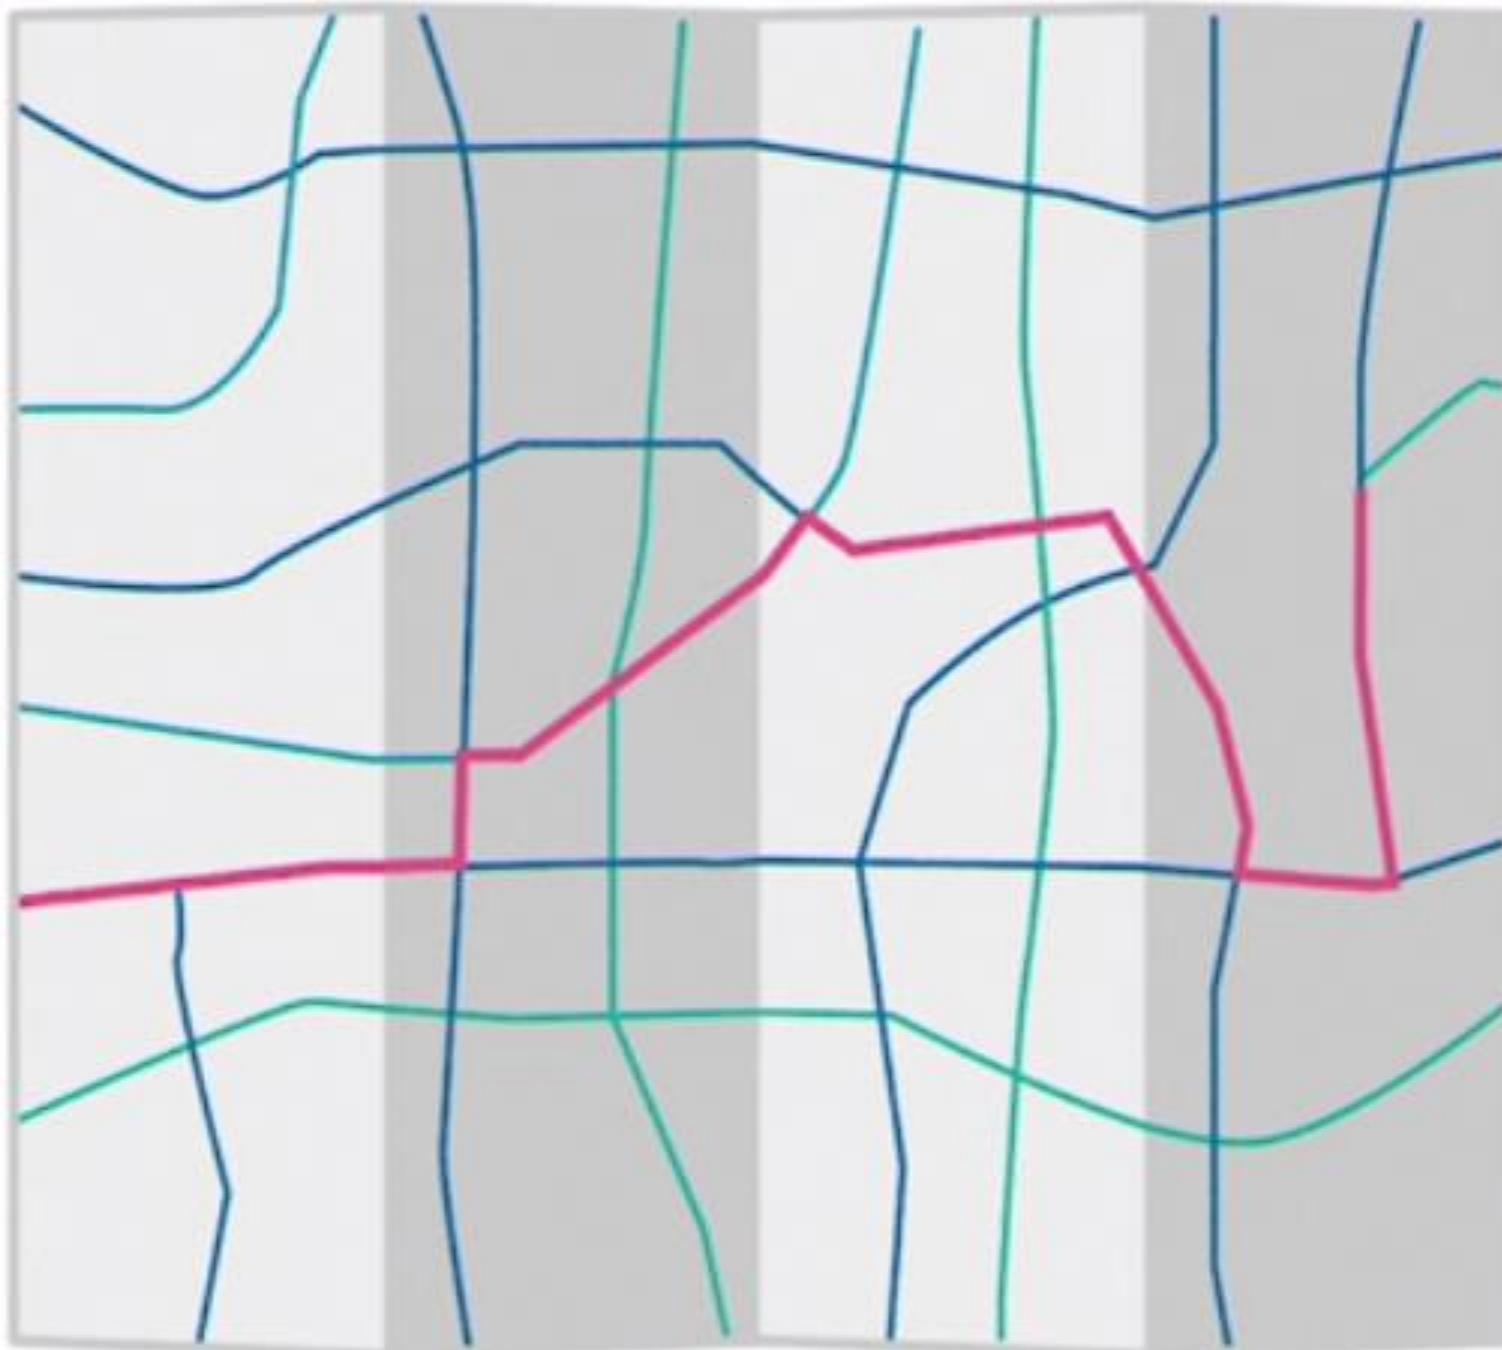
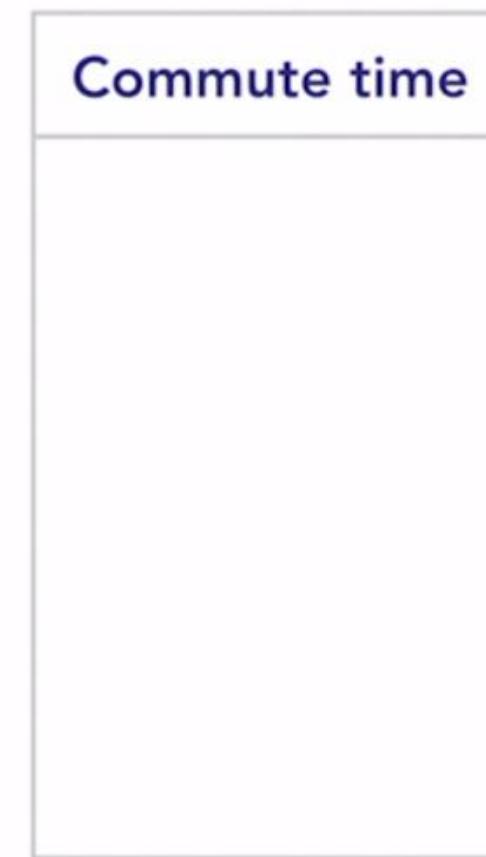
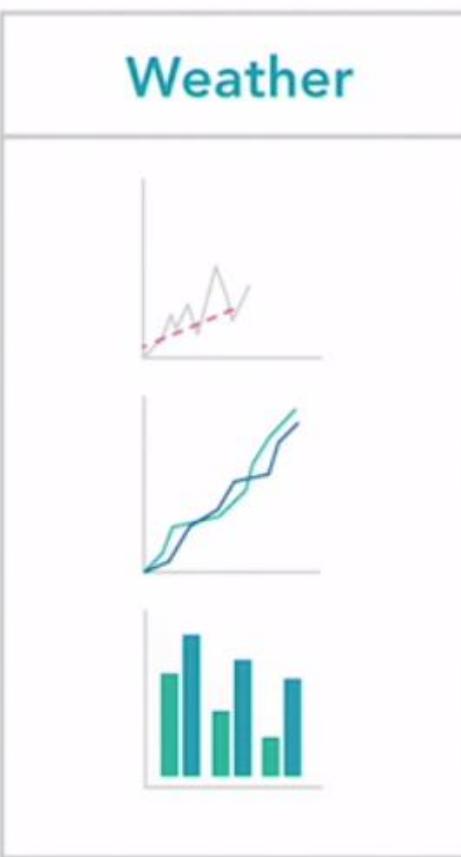


Ready

Done



Test Data

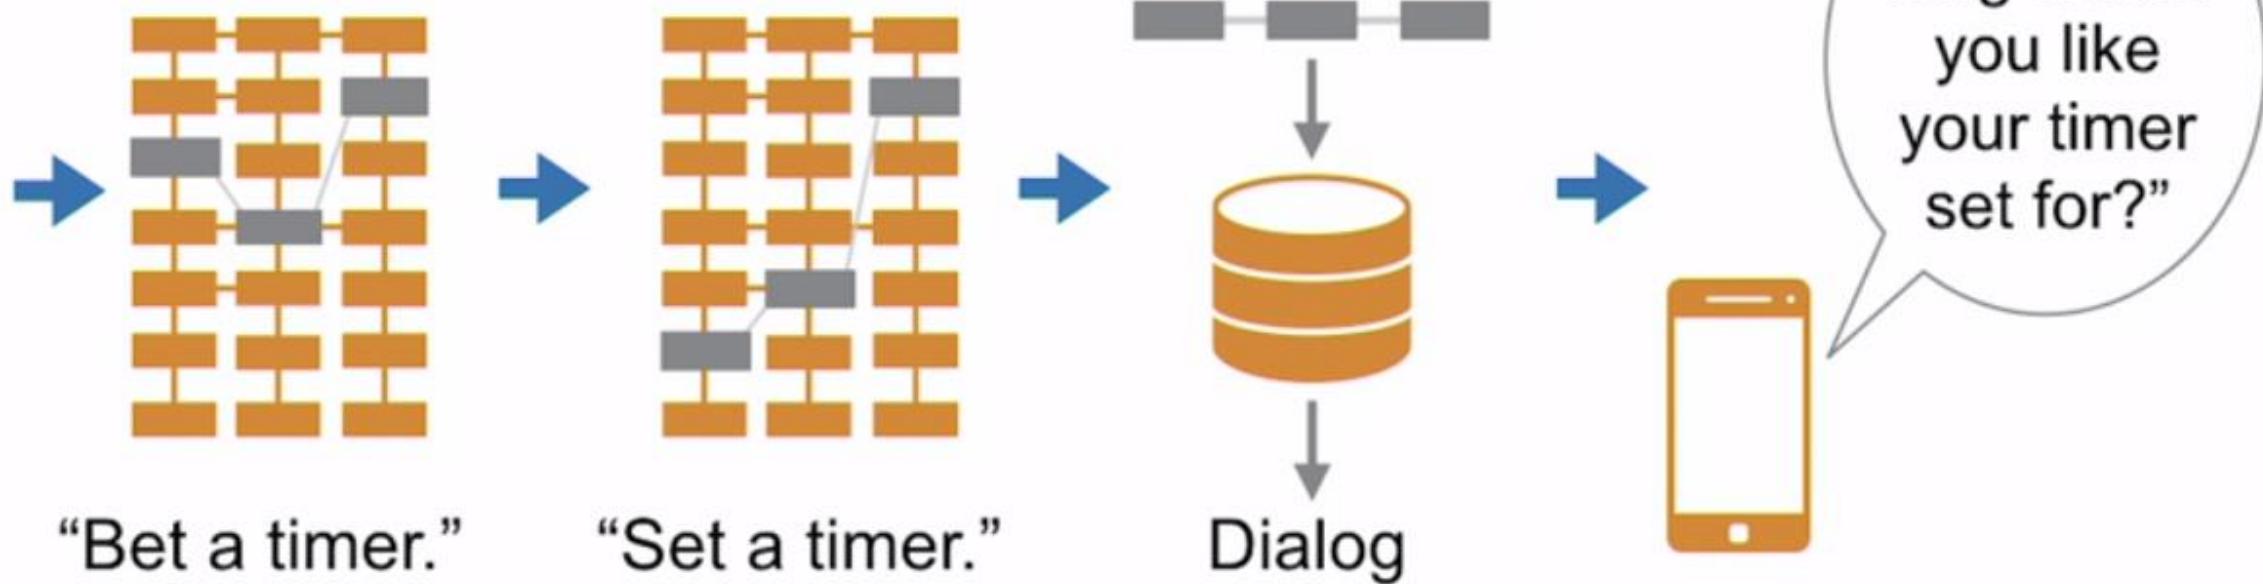


Machine Learning Fundamentals

Machine Learning is all around you!

Personal assistants – Siri, Google Home, Amazon Echo

SIRI Teardown

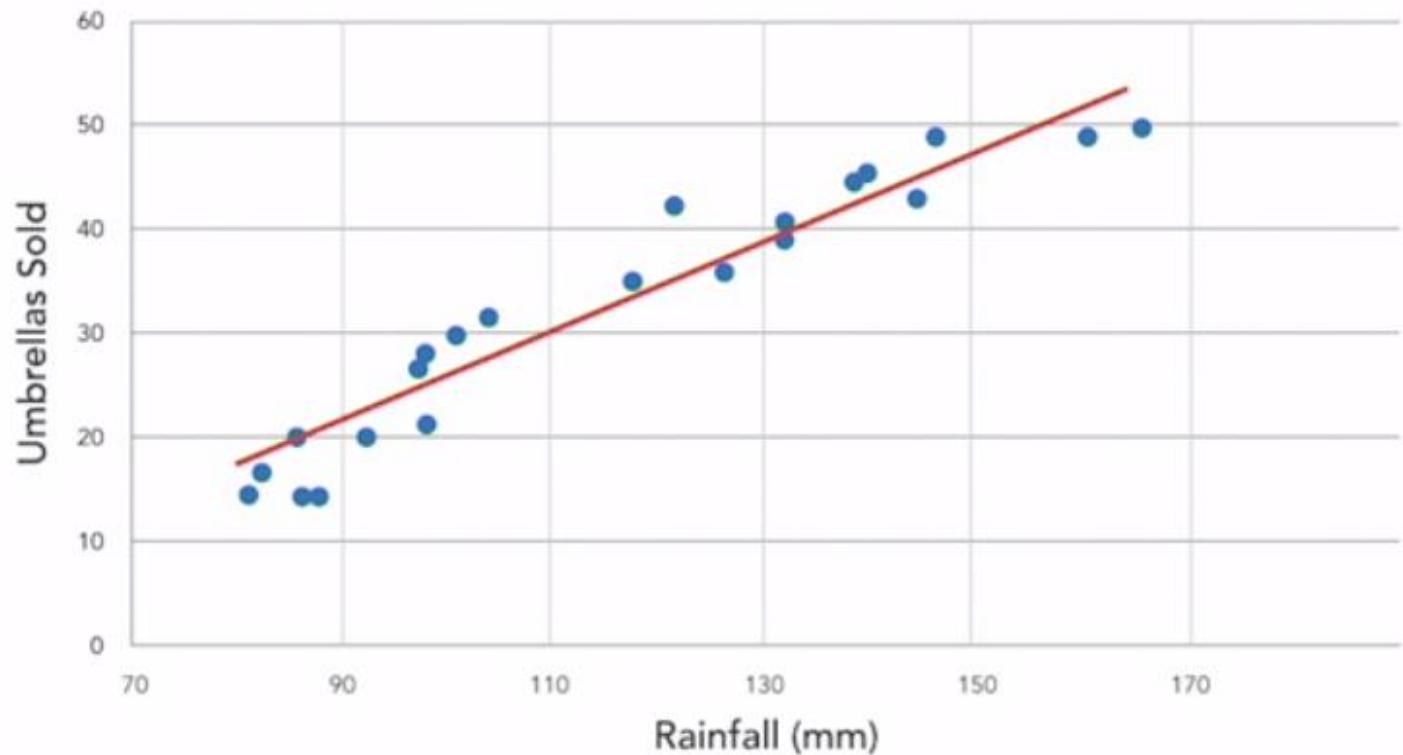


What is Machine Learning ?

- Pattern matching

$$y = mx + b$$

Linear Regression





EXIT

27cn4X0007U1901

9c0U43X09209728

2X13722n5478373



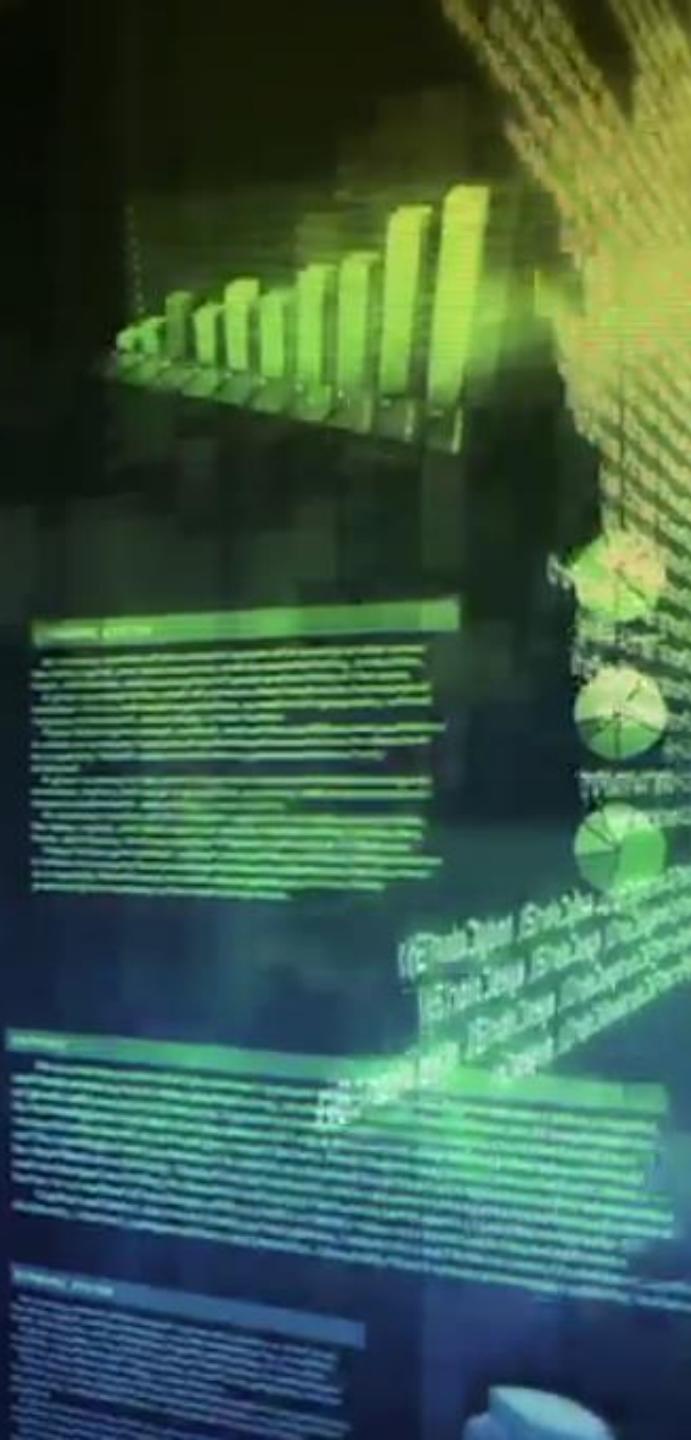
Z99440ha1921291

711y70591F8t910

```
        end
        self.selectedObj:accOn();
        if dir == "left" then
            inv = true
            changeAcc()
        elseif dir == "right" then
            inv = false
            changeAcc()
        end
        elseif selectedObj == randomRect then
            randomMonster()
        elseif selectedObj == backRect then
            if dir == "left" then
                inv = true
                changeBackground()
            elseif dir == "right" then
                inv = false
                changeBackground()
            end
        end
    end
    display.getCurrentStage():setFocus(nil)
end
```

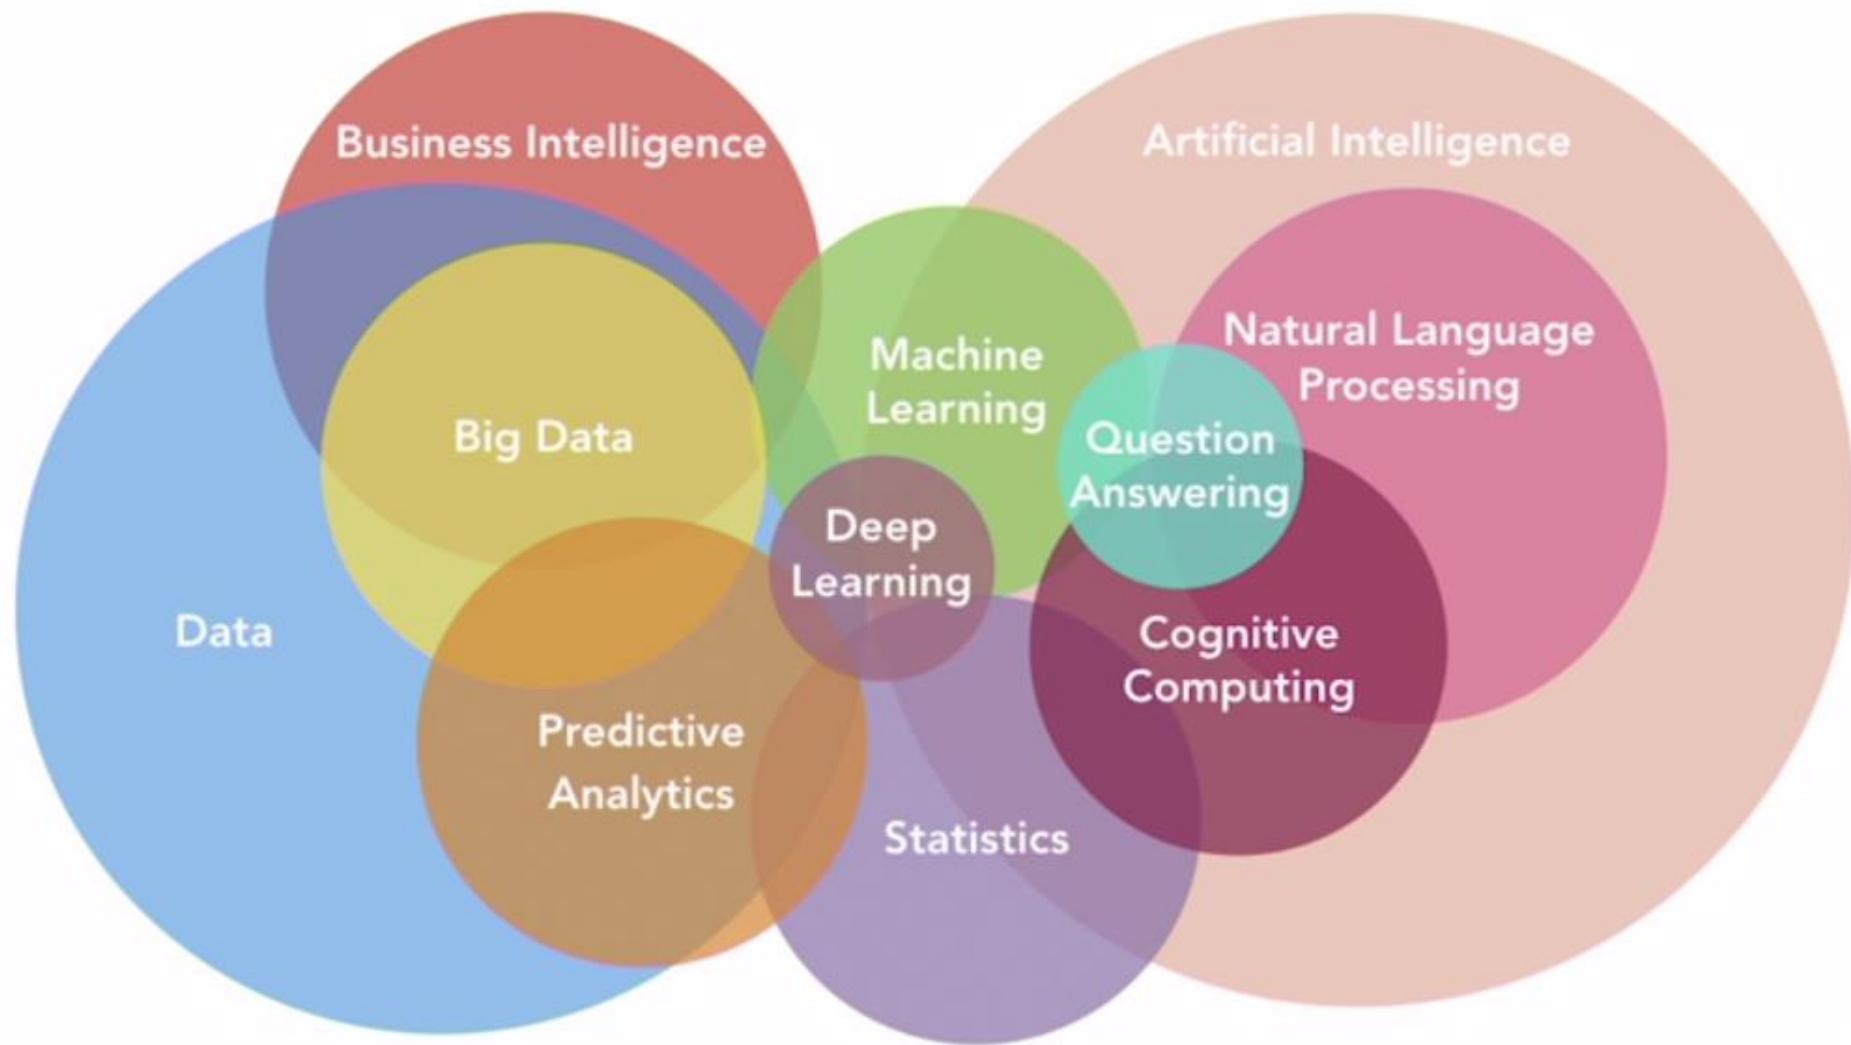
....BURNING....

```
local hideBlink = function()
    eyes.isVisible = true
    eye1.isVisible = true
    burning.isHidden = true
    burningPart.isHidden = true
end
```



3900
GARAGE

002	183	26	-2.2%	123.7
307	163	117.52	-5.98%	321.1
101	0.91	29.18	-4.15%	163.13
194	15.36	156.11	5.98%	31.18
253	1.42	28.58	1.72%	31.18
364	159	97.71	-1.74%	97.08
364	3.07	72.88	-3.41%	69.01
07	15	110.33	17.56%	109.42
133	143	45.55	-4.51%	
257	227	224.2		
144	6.02			
35				



This lecture will cover some of the foundations of machine learning like:

- exploratory data analysis
- cleaning your data
- fitting robust models
- tuning hyperparameters
- and finally, evaluating a model to ensure that it generalizes to unseen examples.



What You Should Know?

- Basic Python Knowledge
- Basic Experience with Numpy, Pandas and Scikit-learn



Good to Have, But not Required!

- Basic Statistics
- Entry Level Data Analysis Experience
- Familiarity with common Machine Learning Terms



Tools we would be using

- Python 3.4
- Jupyter Notebook

Jupyter Notebooks is a fantastic tool that allows you to combine:

- live code,
- explanatory text,
- visualizations,
- and equations all in one concise package.

"Field of study that gives computers the ability to learn without being explicitly programmed."

—Arthur Samuel, IBM, 1959

"Machine learning algorithms can figure out how to perform important tasks by generalizing from examples."

—Pedro Domingos, University of Washington

Fitting a function to examples and using that
function to generalize and make predictions
about new examples

What is Machine Learning?

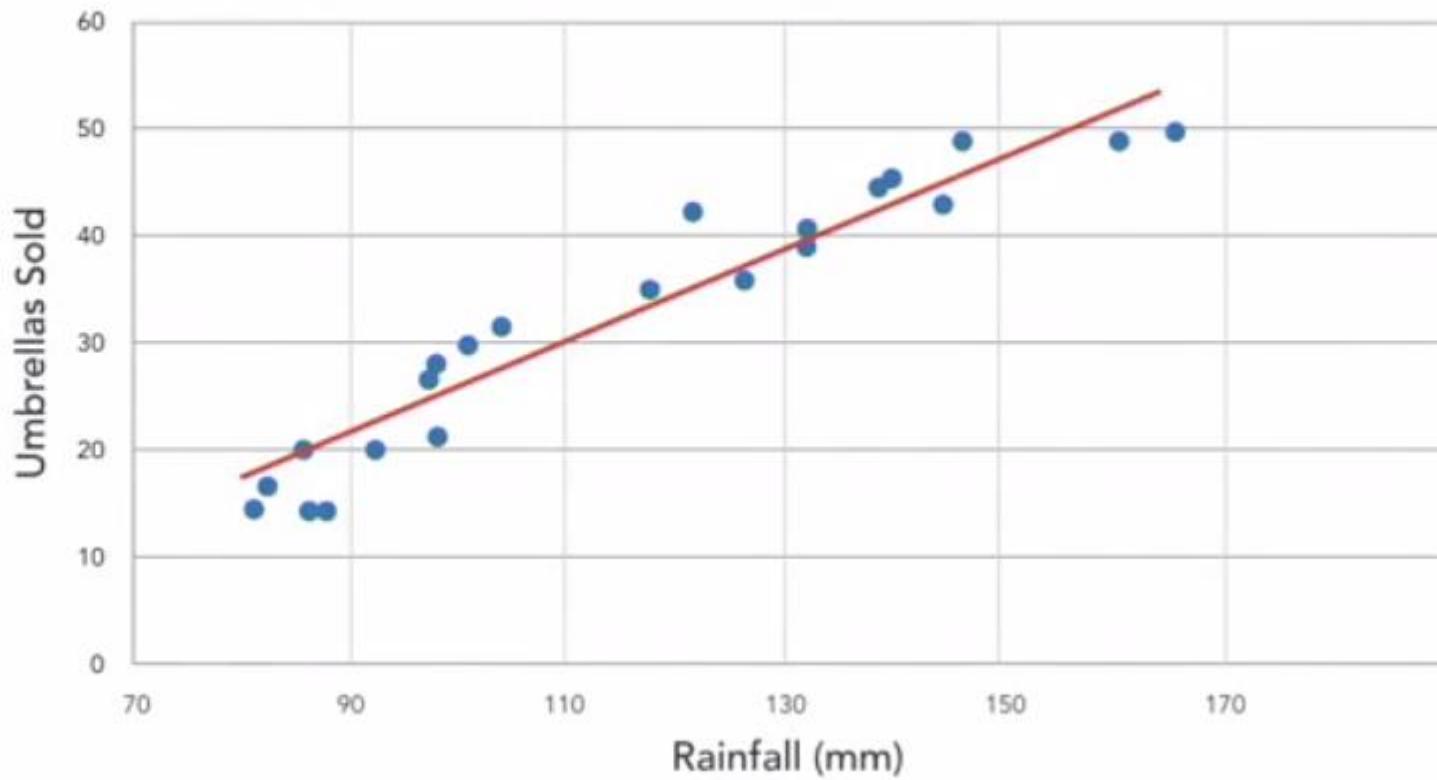
- Pattern Matching

What is Machine Learning?

- Pattern matching

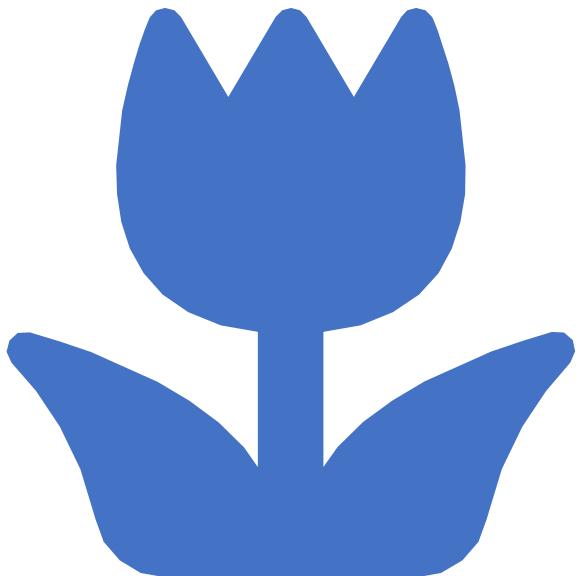
$$y = mx + b$$

Linear Regression



150 Flowers

- Determine flower species by features such as petal length



SETOSA

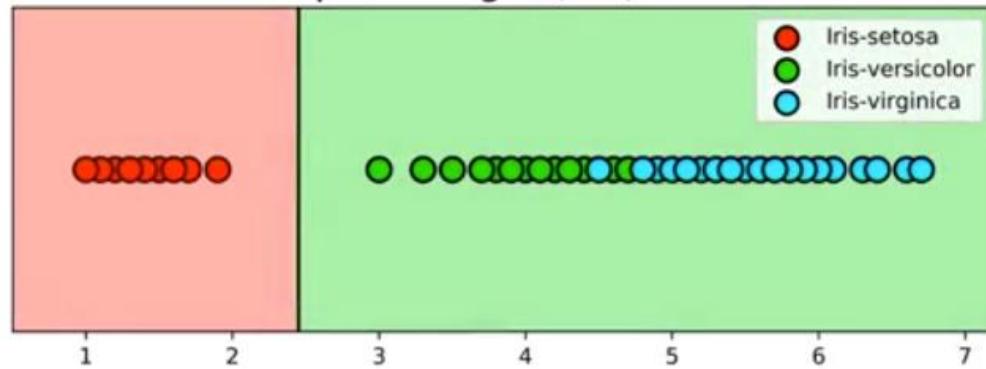


VERSICOLOR



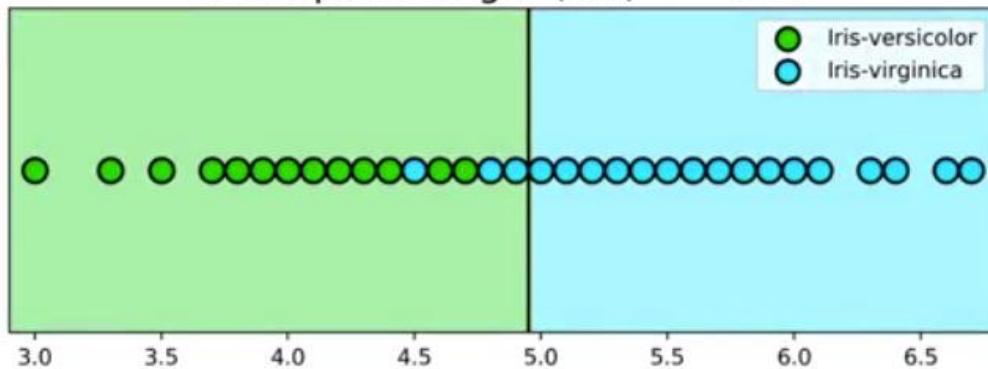
VIRGINICA

Is the petal length (cm) ≤ 2.45



Separate Setosa from other Species

Is the petal length (cm) ≤ 4.95



Separate Versicolor from other Virginica

You could keep writing rules until you completely separate out the flower species

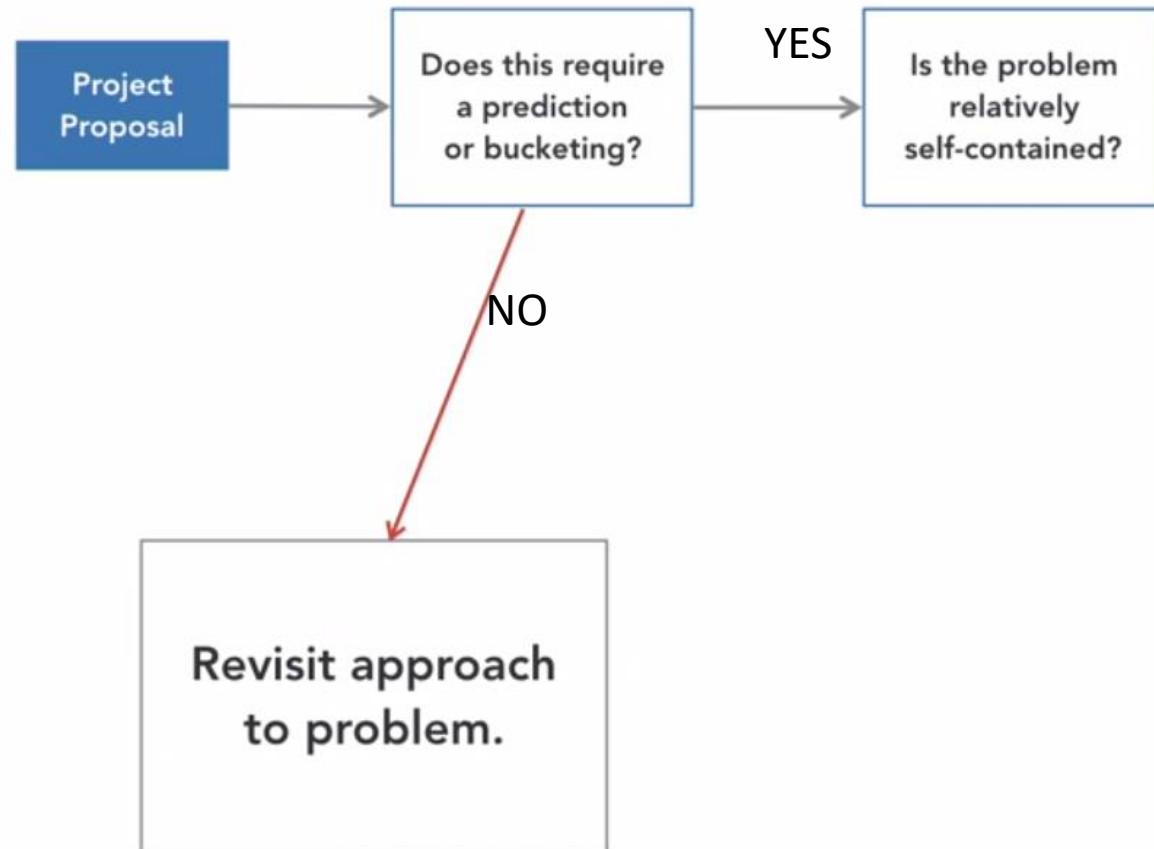
Assess the Problem You're Trying to Solve

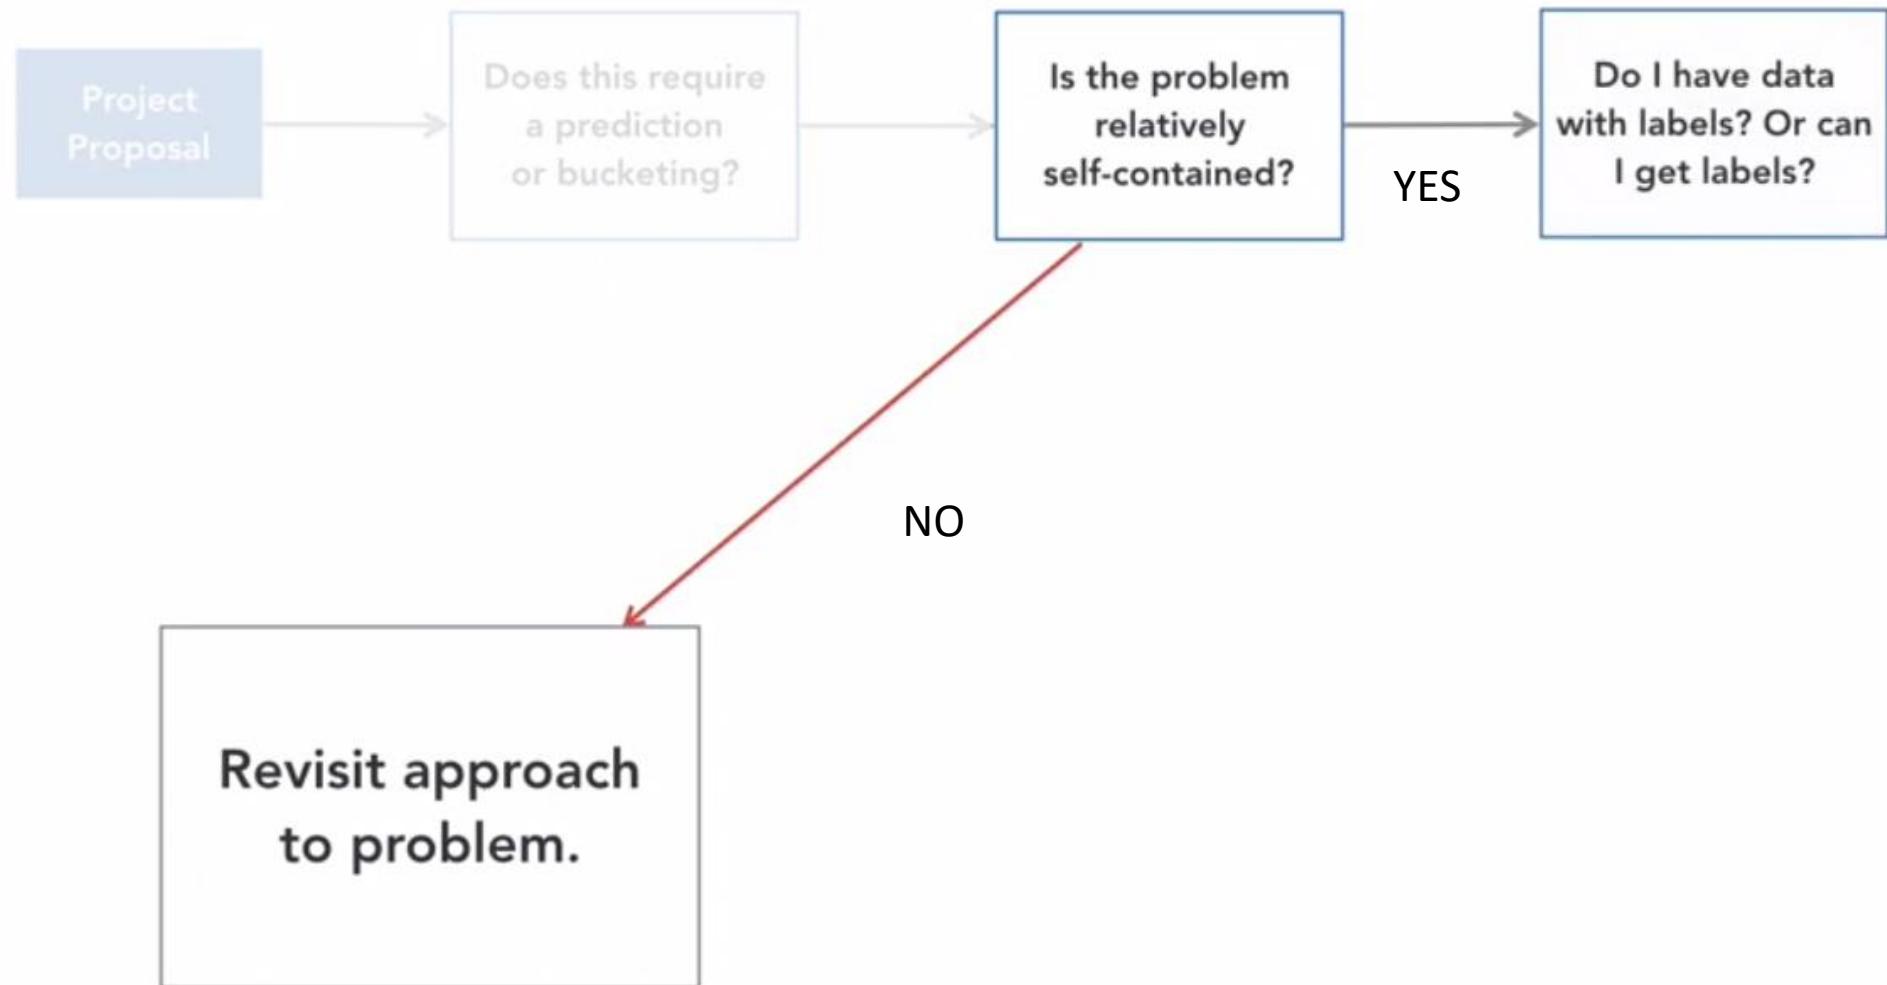
Is this a type of problem that can be solved using machine learning?

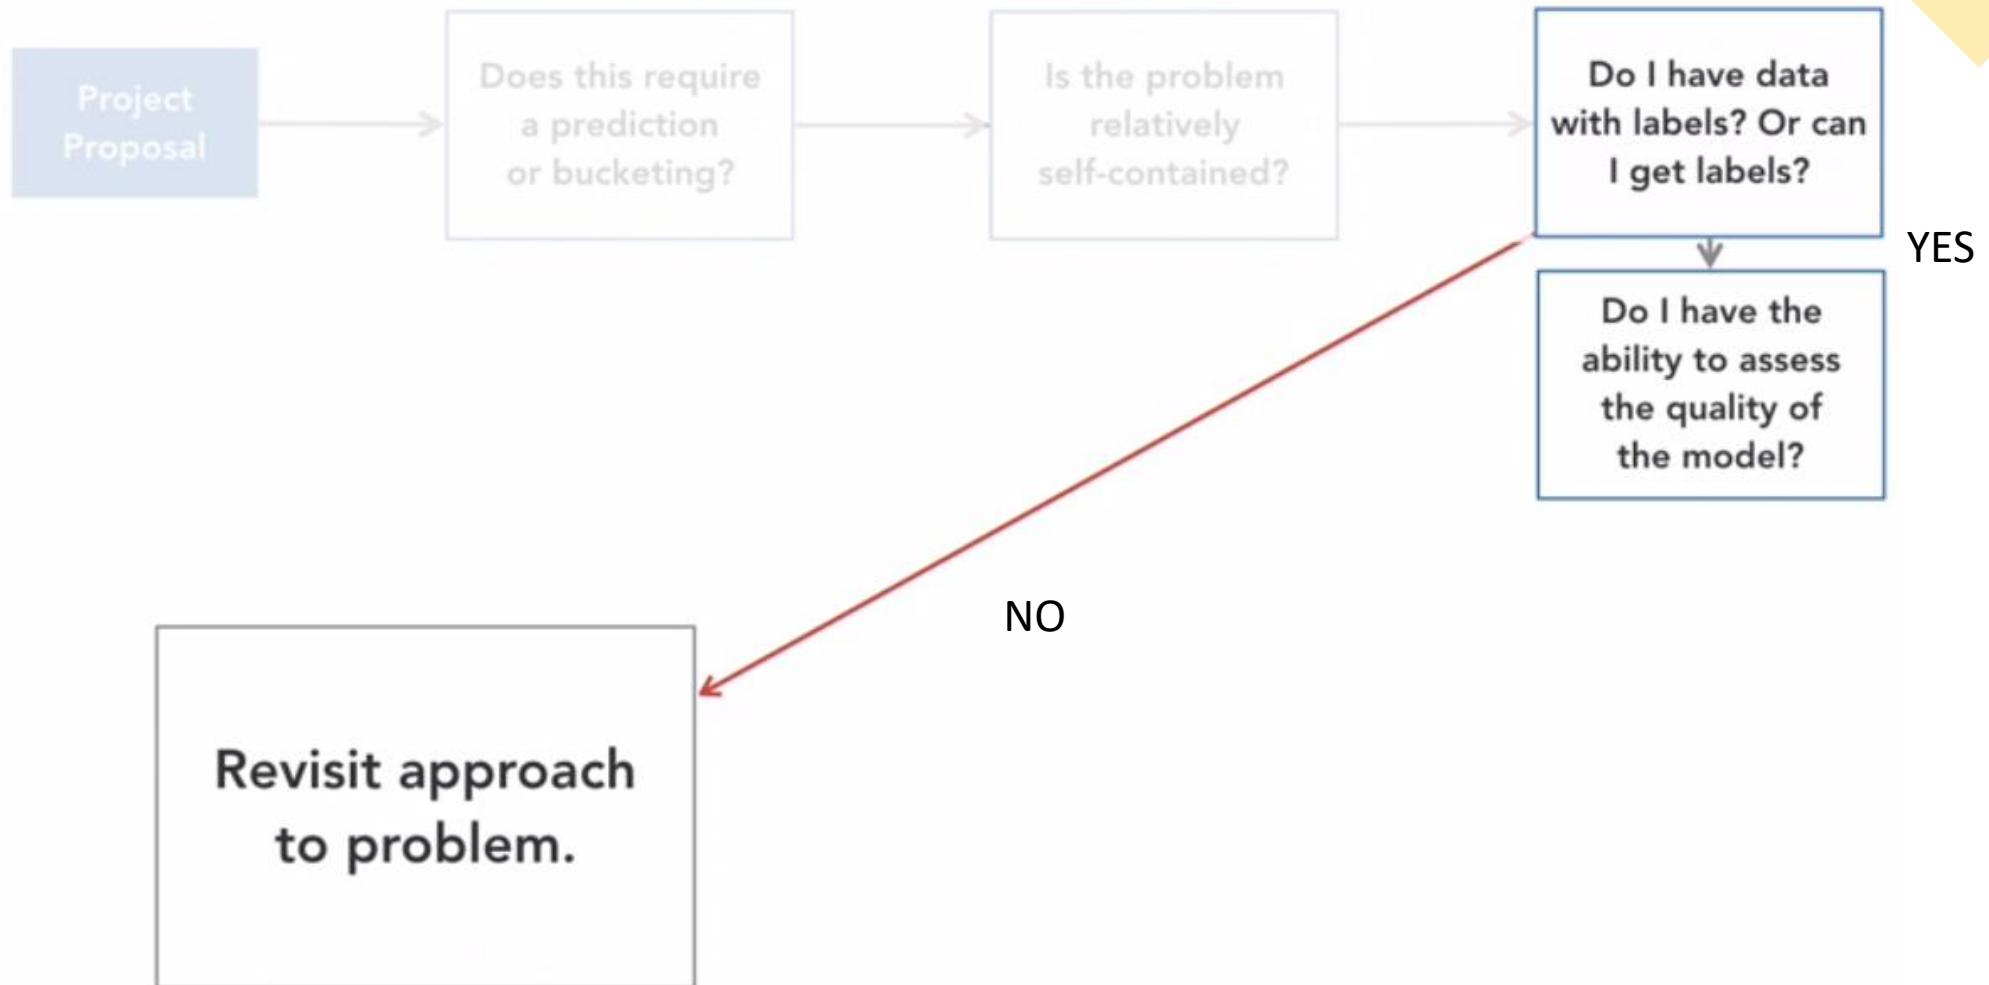
- Does it require a prediction or bucketing something into categories?
- Is the problem relatively self-contained?

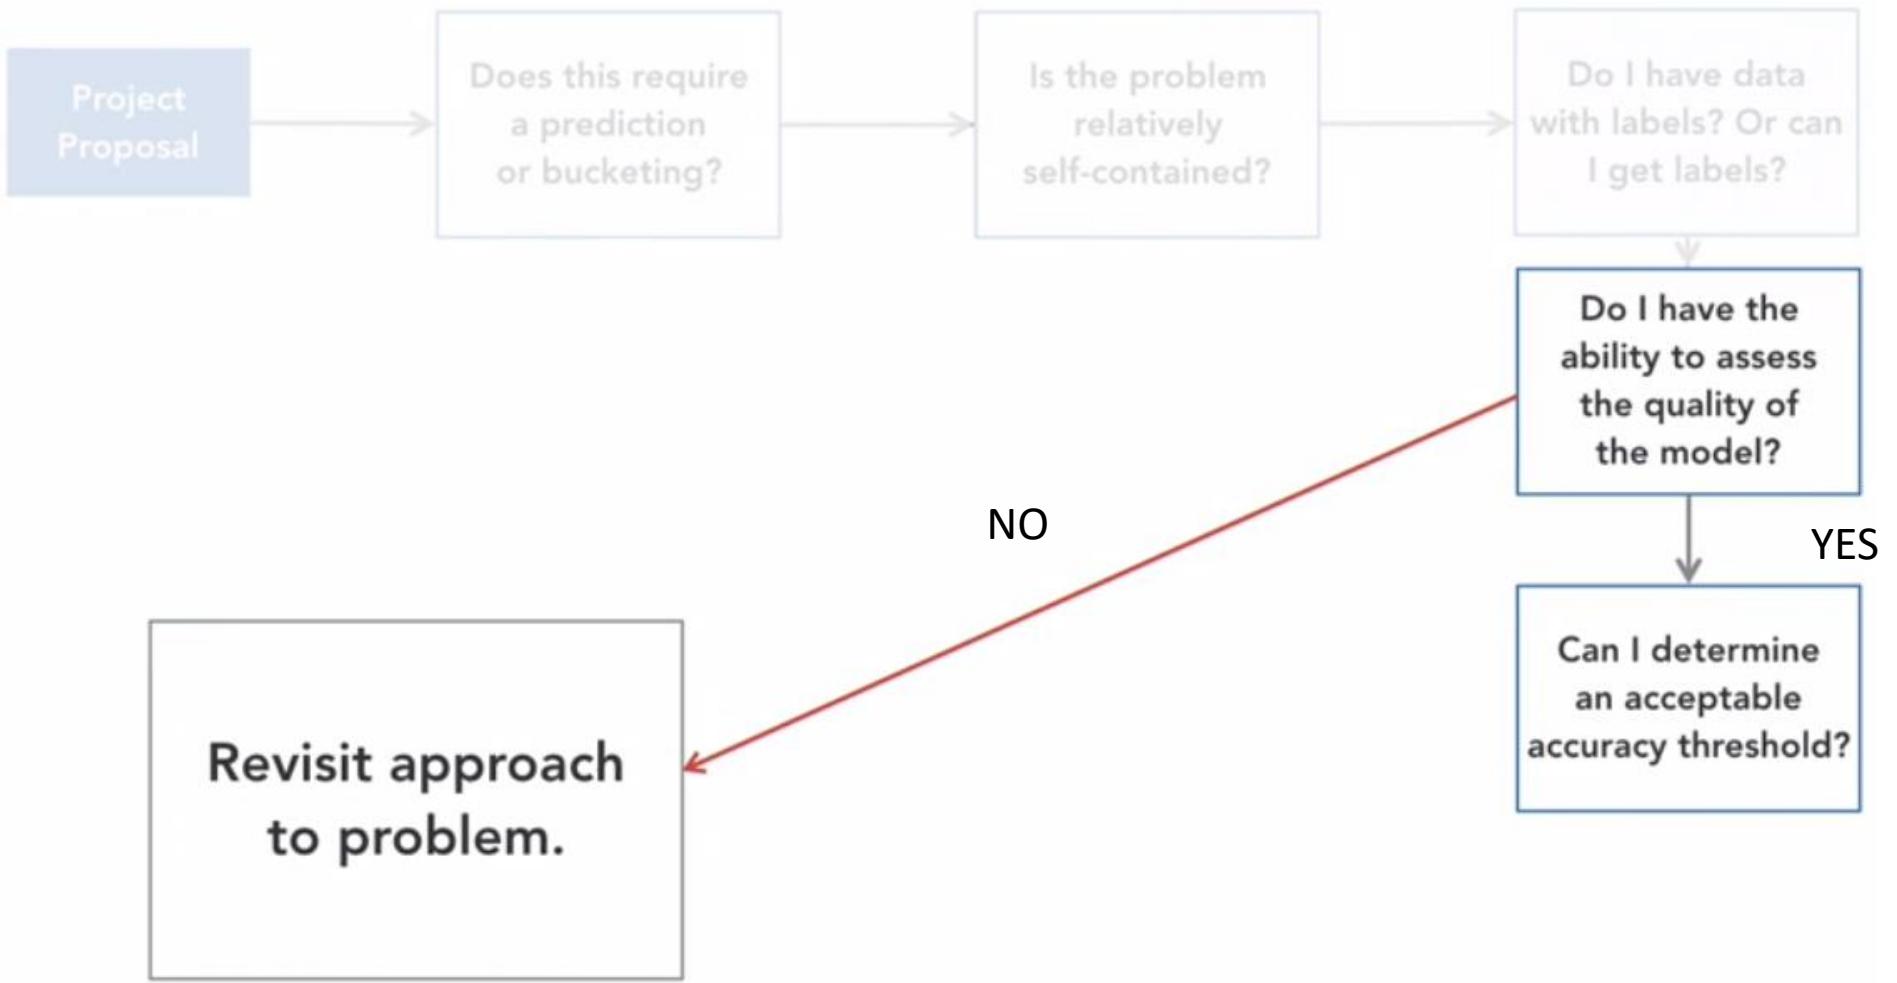
Is my particular problem one that can be solved with machine learning?

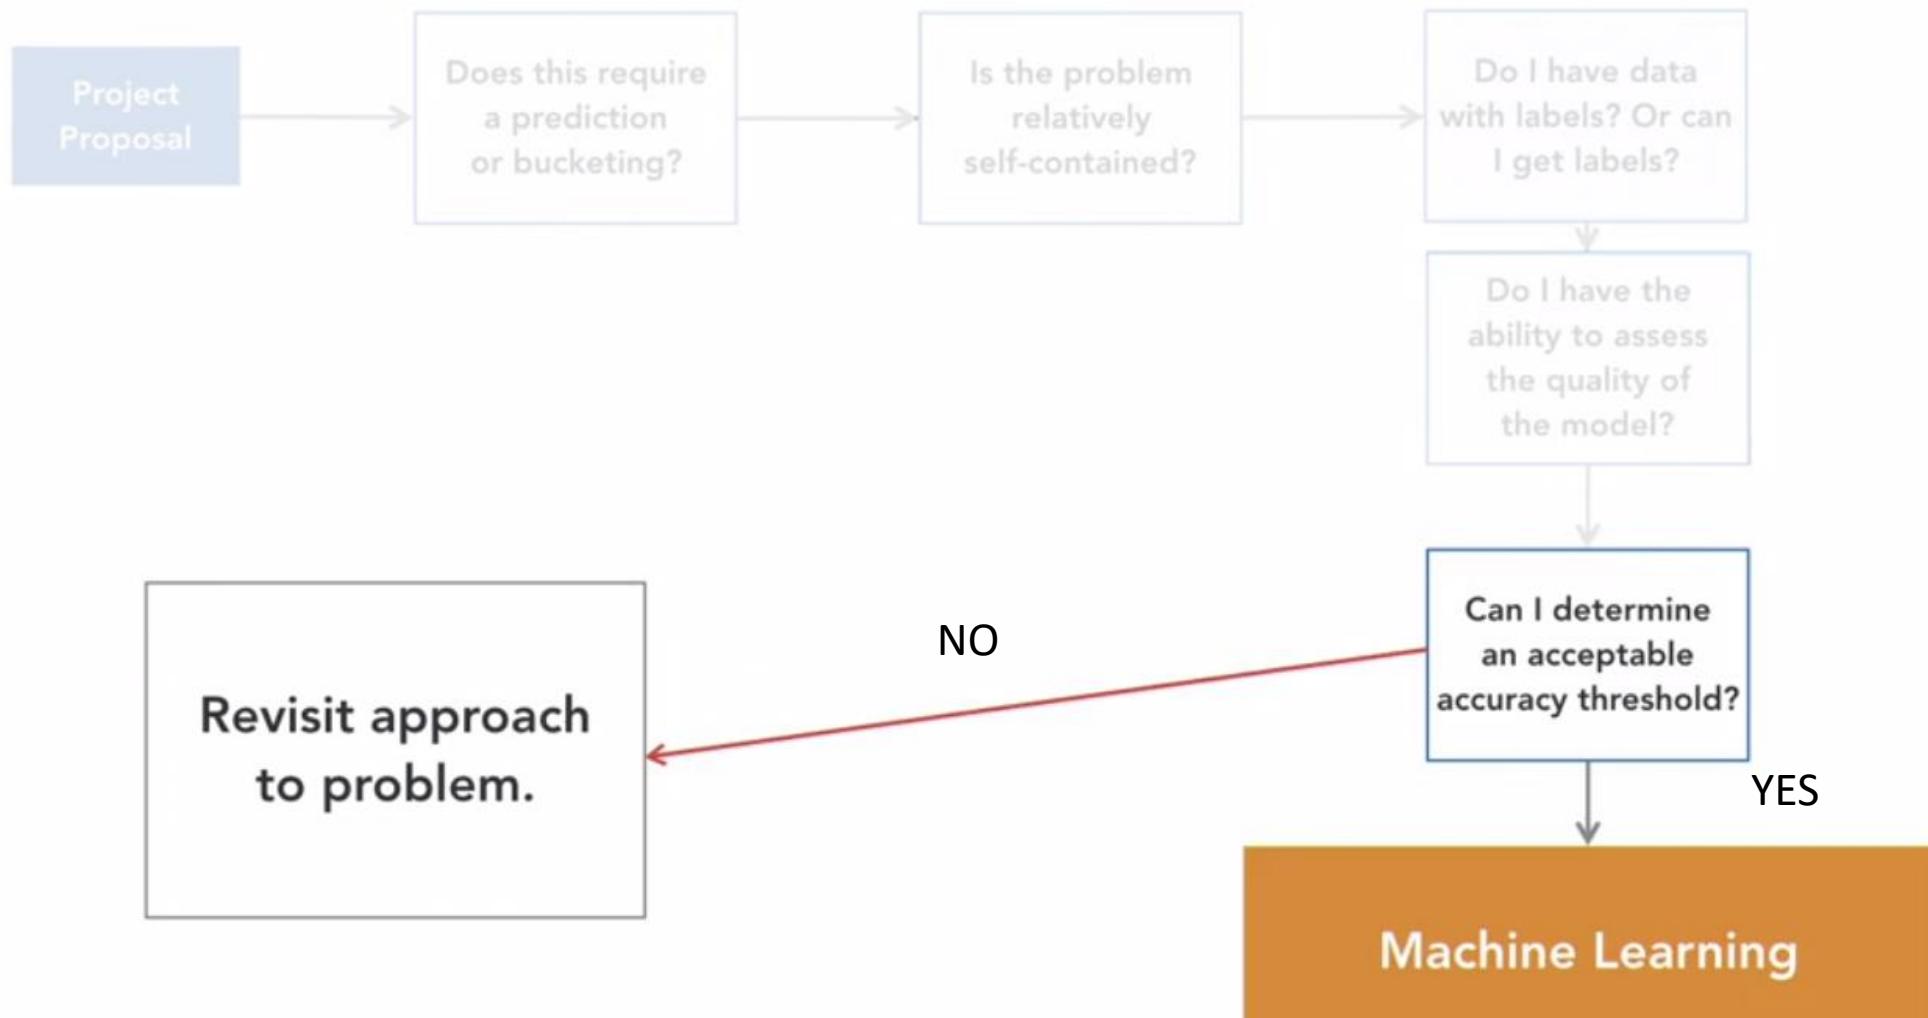
- Do I have data with labels?
- Do I have the ability to assess the quality of the model?
- Can I determine an acceptable accuracy threshold?





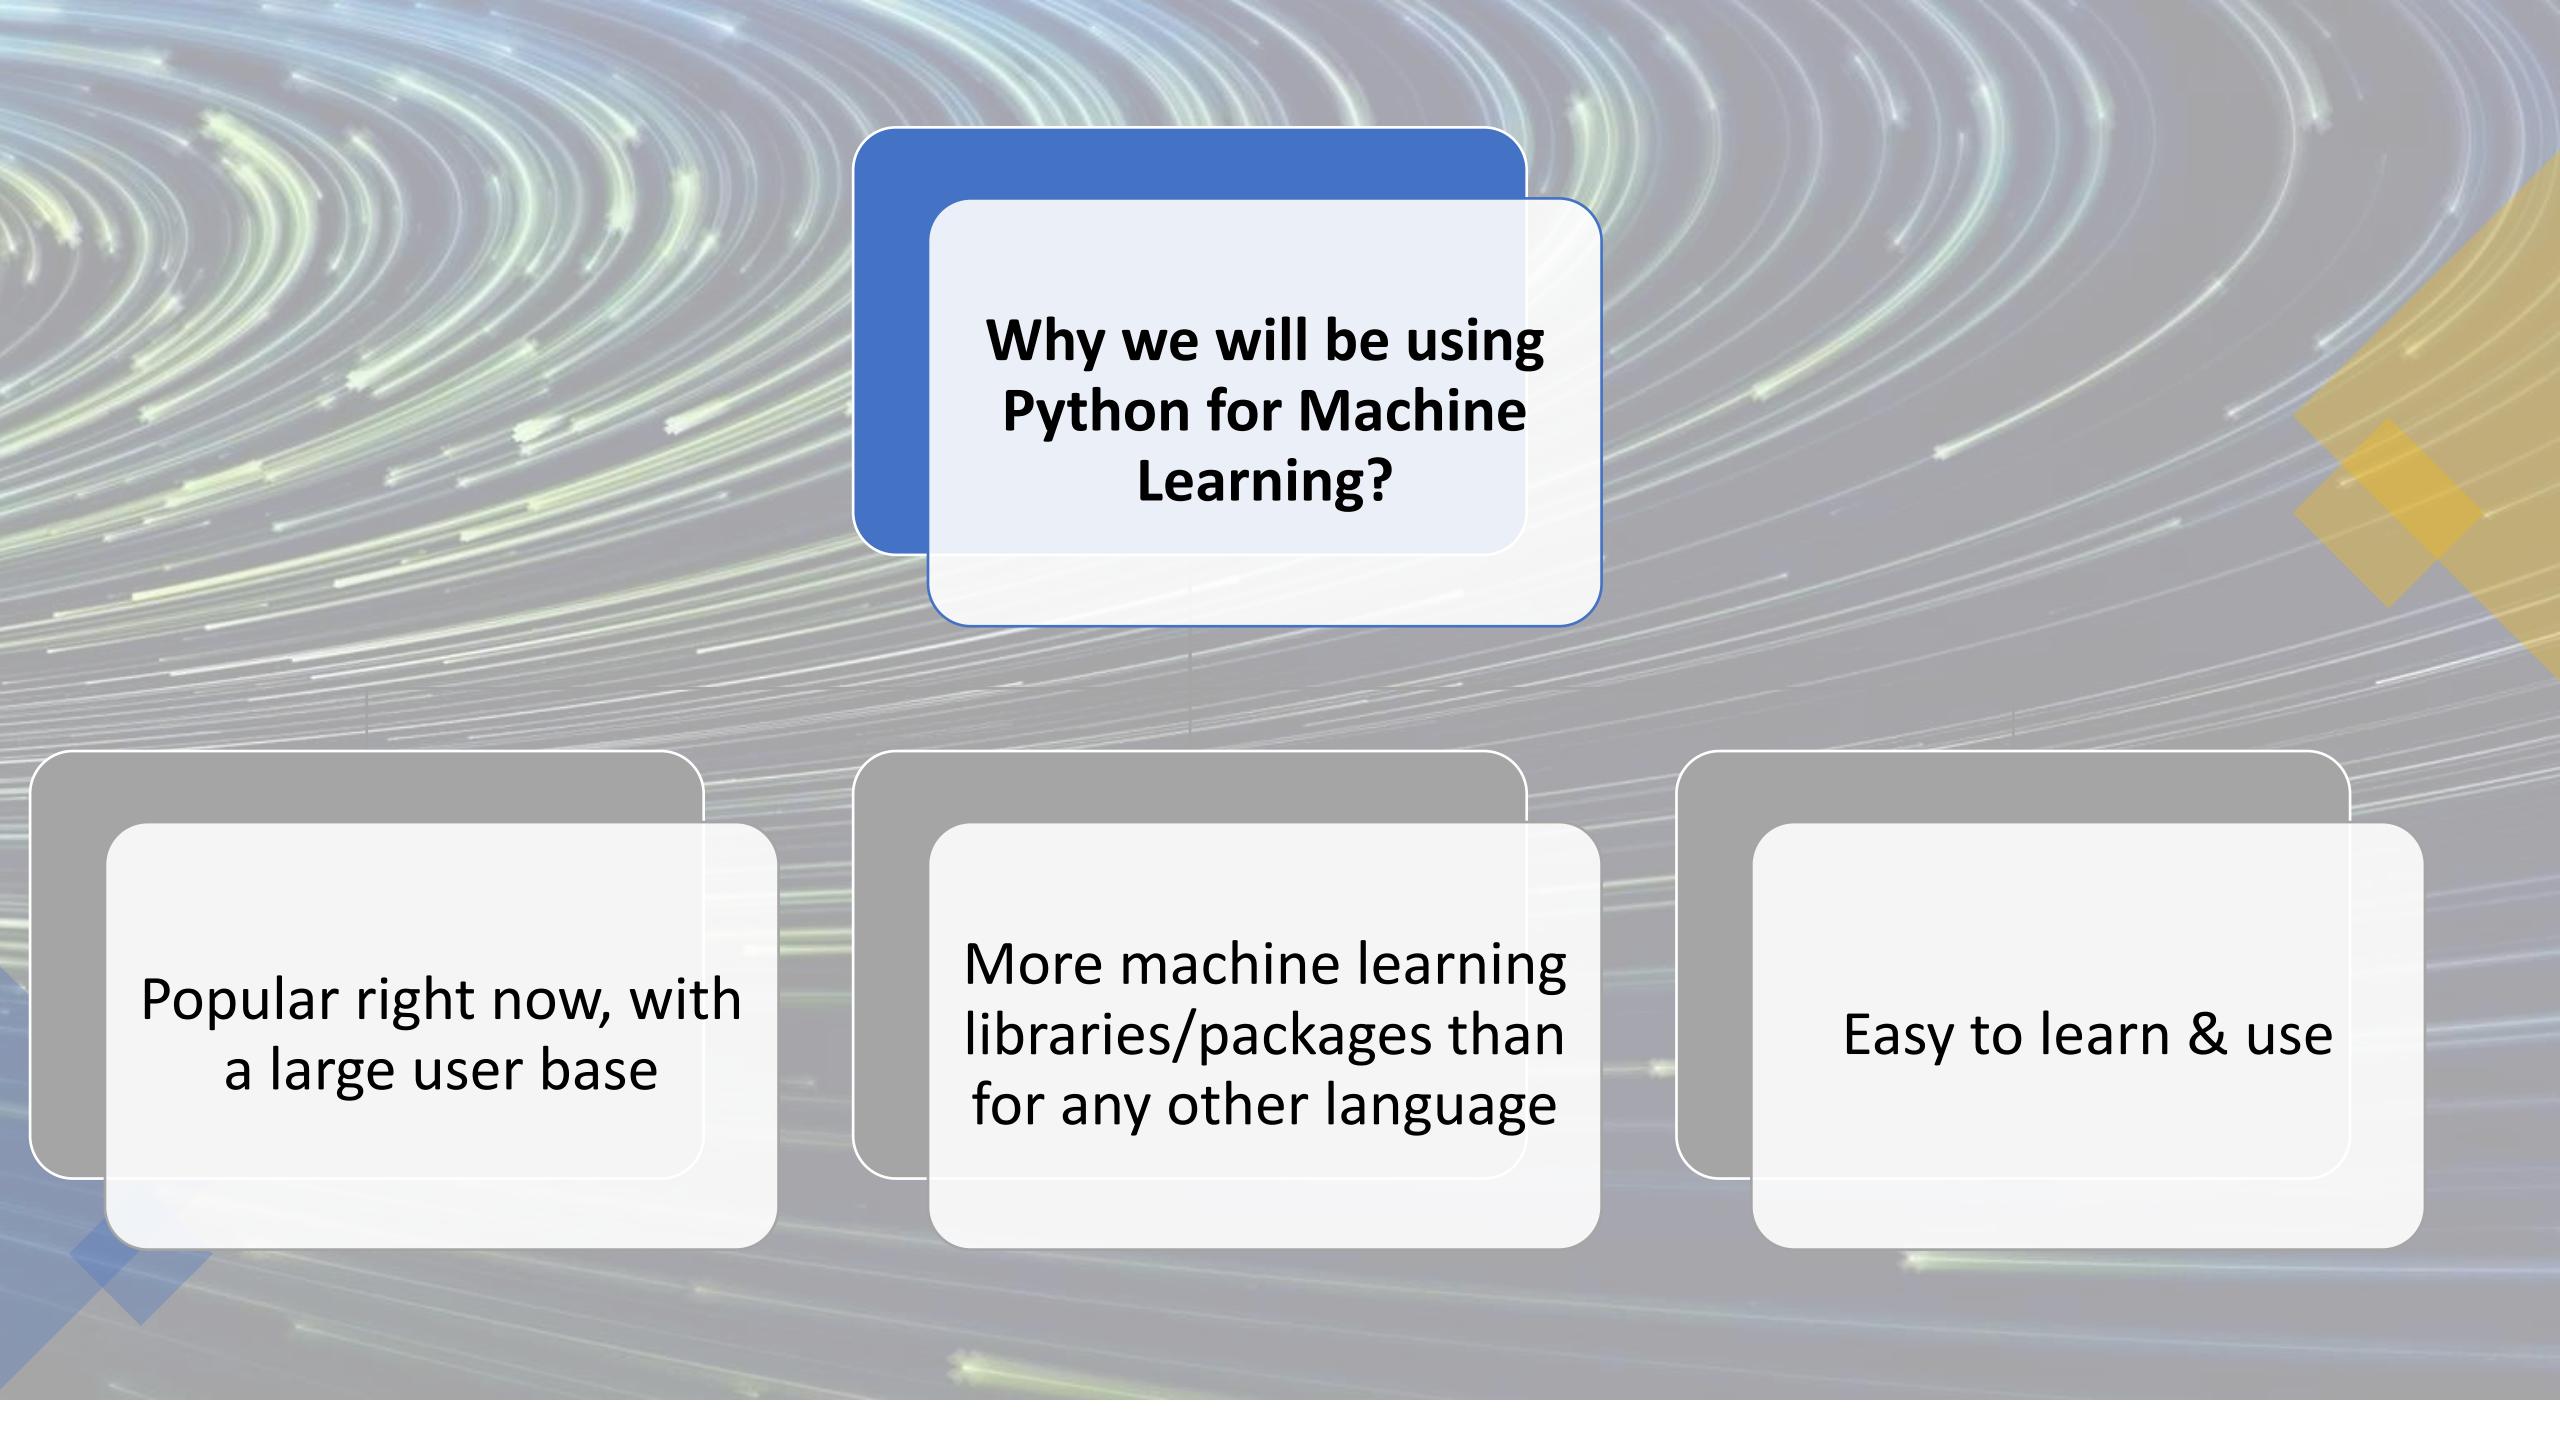






Examples of machine learning in everyday life

- spam detection
- speech recognition
- recommendations
- fraud detection.



Why we will be using Python for Machine Learning?

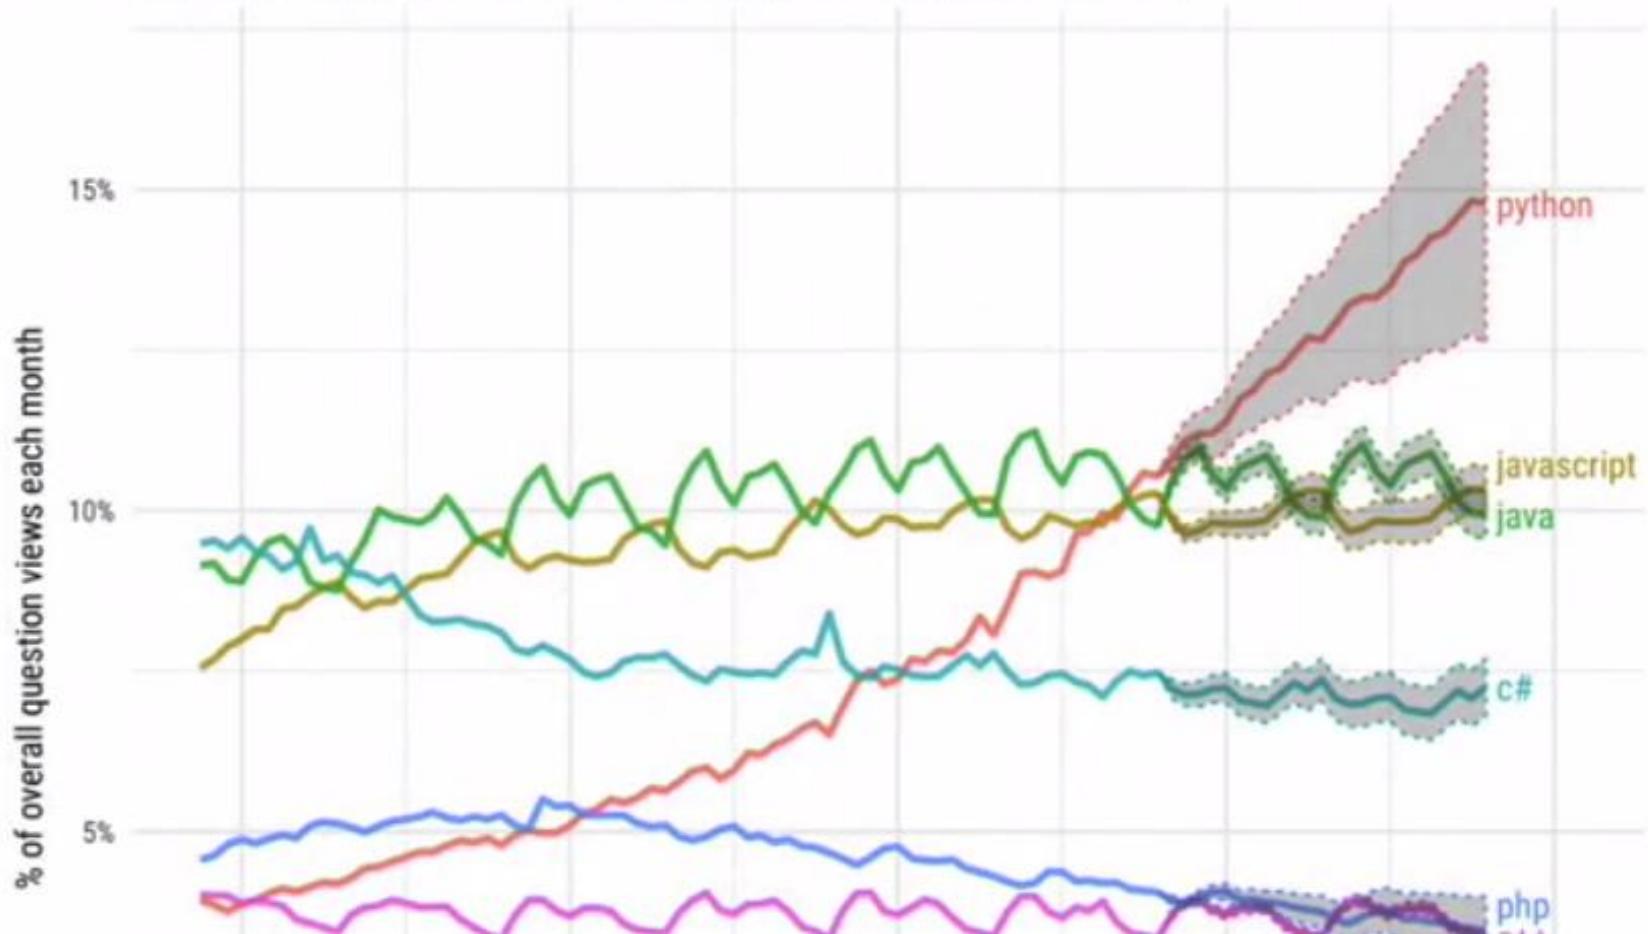
Popular right now, with
a large user base

More machine learning
libraries/packages than
for any other language

Easy to learn & use

Projections of future traffic for major programming languages

Future traffic is predicted with an S&L model, along with an 80% prediction interval.



Google Trends

Compare



Interest over time



100
75
50
25
0

Mar 23, 2013 Dec 13, 2015 Sep 3, 2017

Average

Note

Python Is a Swiss Army Knife

- Machine learning – scikit-learn, LightGBM, SciPy
- Deep learning – TensorFlow, PyTorch, keras
- Text – NLTK, genism



Python Is a Swiss Army Knife

- Images –
OpenCV, scikit-learn
- Visualization –
Matplotlib, Seaborn
- Data analysis –
pandas, NumPy



```
In [2]: import warnings  
warnings.filterwarnings("ignore", category=FutureWarning)  
  
from sklearn.datasets import load_iris  
from sklearn.linear_model import LogisticRegression
```

```
In [3]: iris = load_iris()
```

```
In [4]: model = LogisticRegression().fit(iris['data'], iris['target'])
```

```
In [6]: model.predict(iris['data'])
```

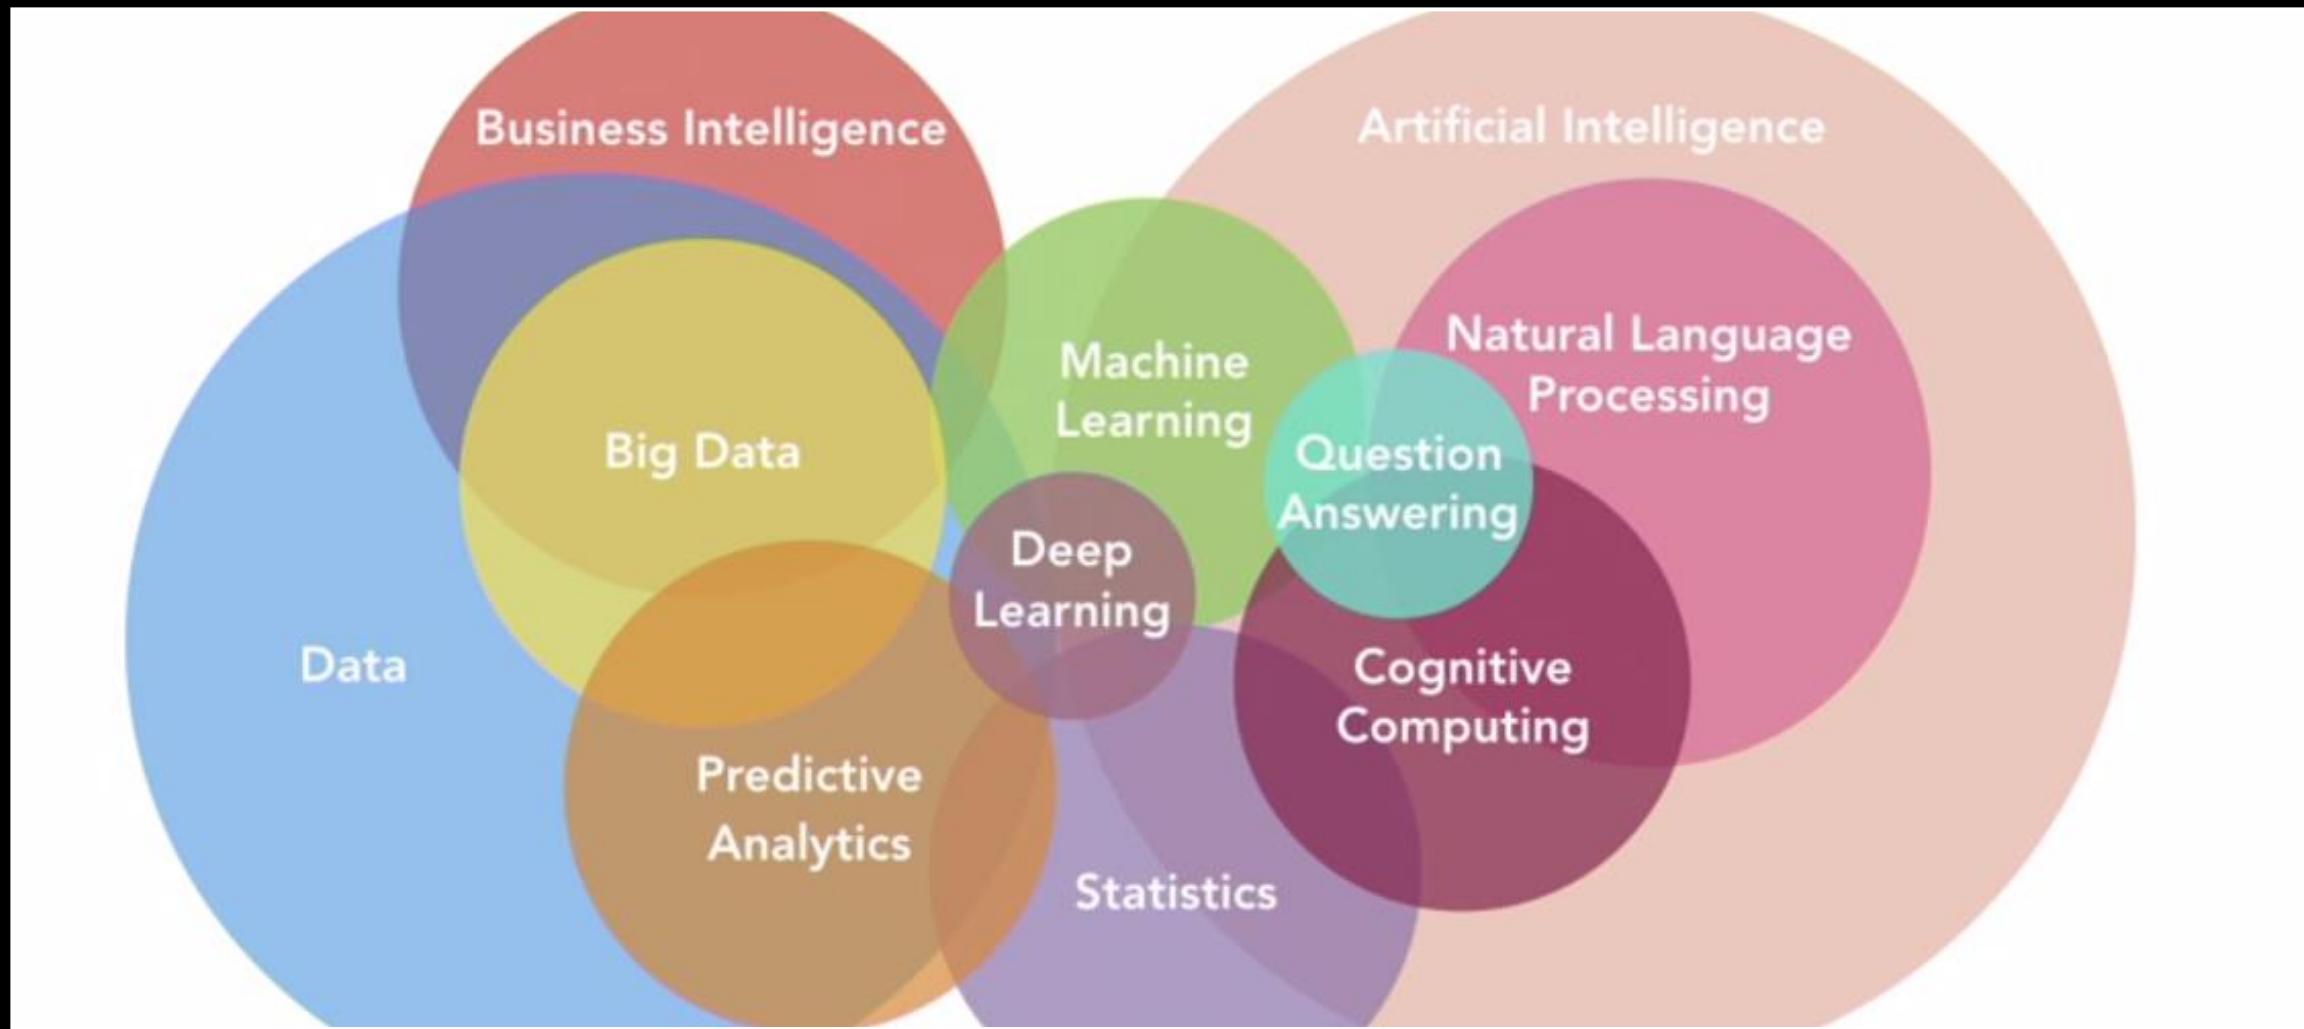
```
Out[6]: array([0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,  
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,  
0, 0, 0, 0, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,  
2, 1, 1, 1, 2, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 2, 2, 2, 2, 1, 1,  
1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,  
2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 1, 2, 2,  
2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2])
```

Scikit-learn is a very popular open-source machine learning project which is constantly being developed and improved.

Why should you use scikit-learn?

- Provides a large number of machine learning algorithms
- Has a clean, uniform, and streamlined API
- Works well with other python libraries
- Switching algorithms is very straightforward
- Widely used in industry and academia

Scikit-learn is a great library for machine learning.



Machine Learning Vs Deep Learning Vs Artificial Intelligence !!!

Fitting a function to examples and using
that function to generalize and make
predictions about new examples

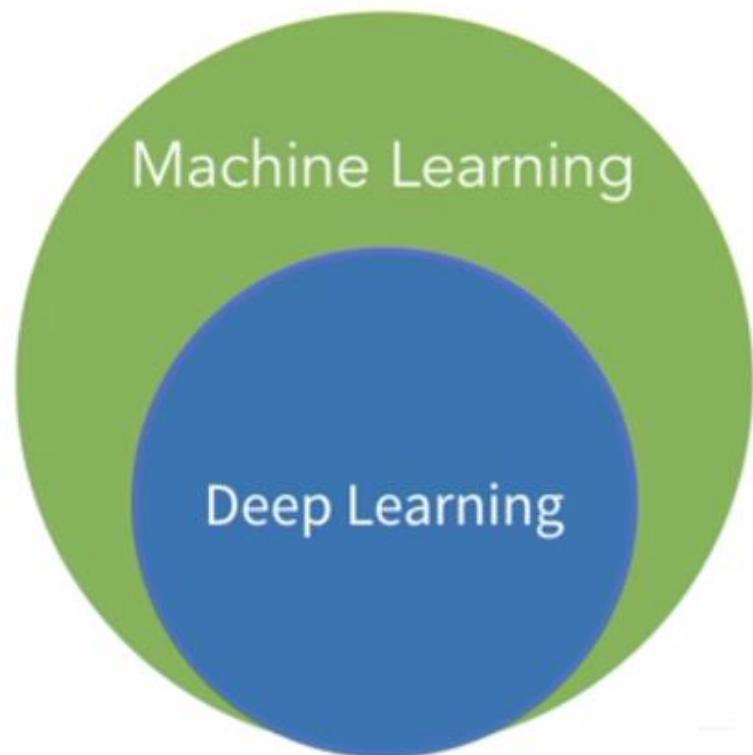


Machine Learning

Alternative: pattern matching

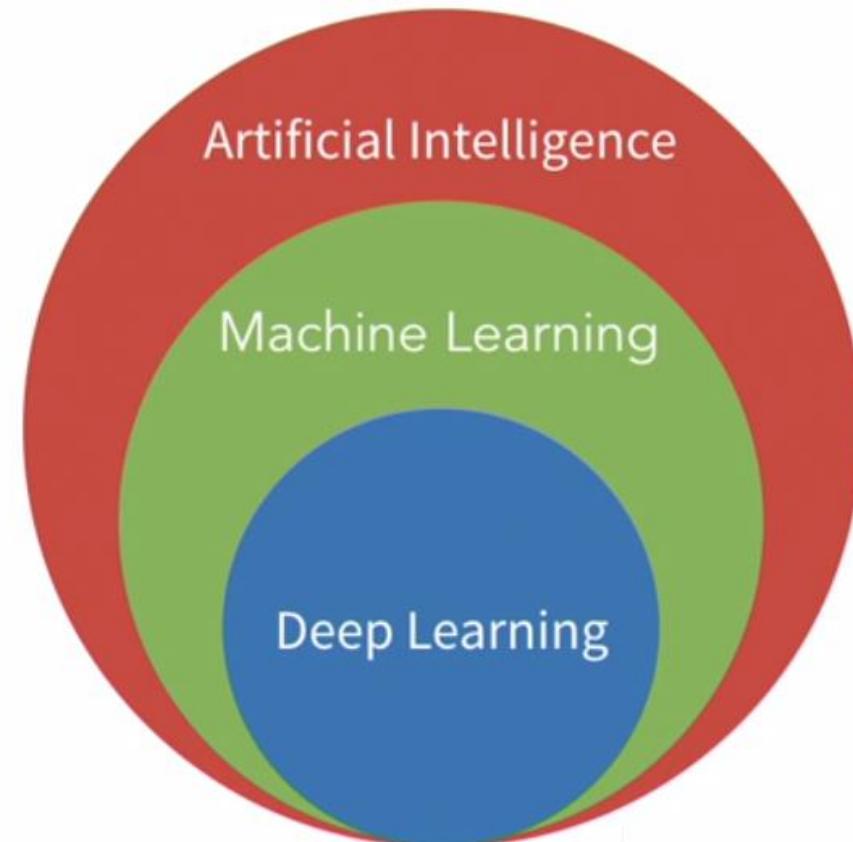
Fitting functions to examples where those functions are connected layers of nodes; with the intent to generalize and make predictions about new examples

Alternative: connected pattern matching



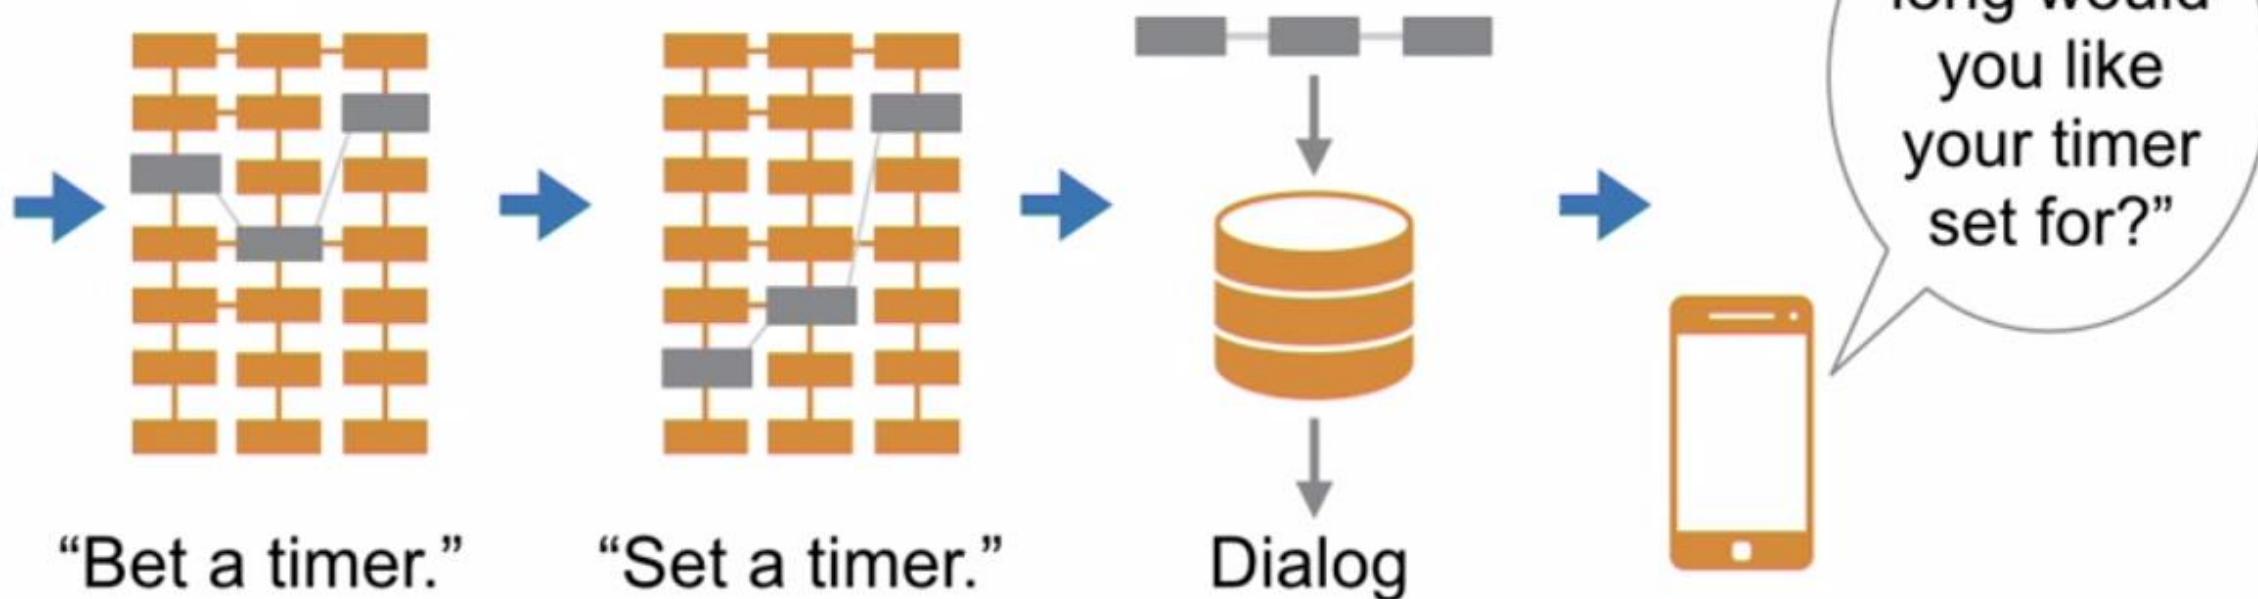
Weak AI – intelligence specifically designed to focus on a narrow task

Strong/general AI – a machine with consciousness, sentience, and a mind; general intelligence capable of any and all cognitive functions and reasoning that a human is capable of

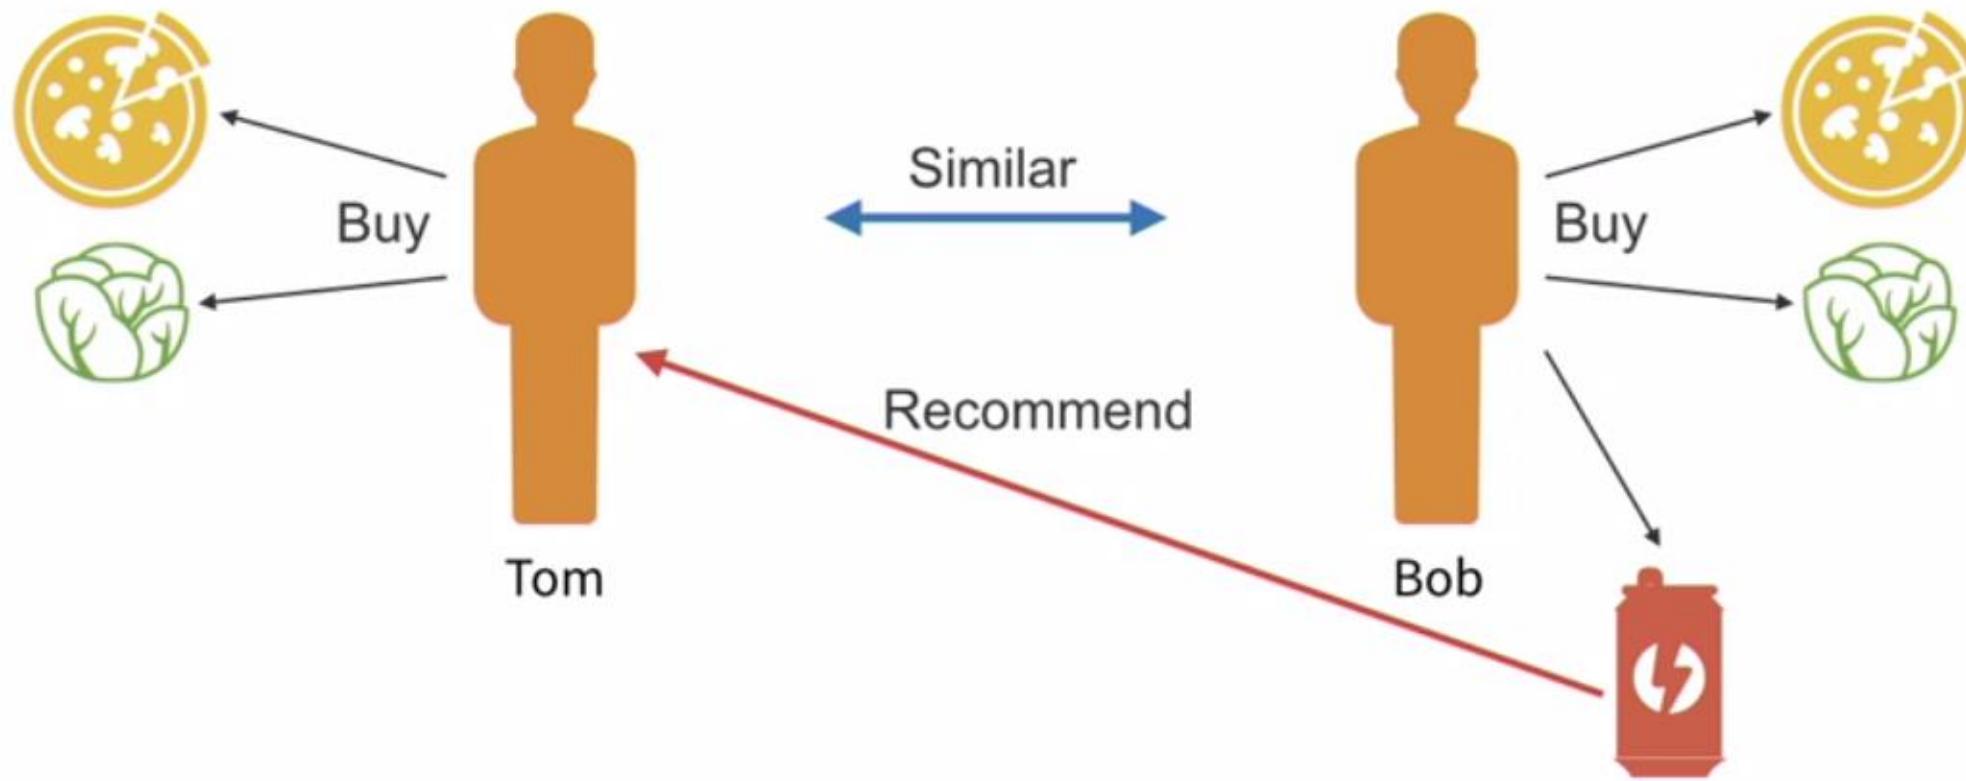


Personal assistants – Siri, Google Home, Amazon Echo

SIRI Teardown



Recommendation systems – Amazon, Netflix, Facebook, Twitter



Los Angeles



FILTERS

CHARTS

Zone Type

Census Tracts

Date-Time Range

1/1/2018 - 1/31/2018 (Every day, Daily Average)

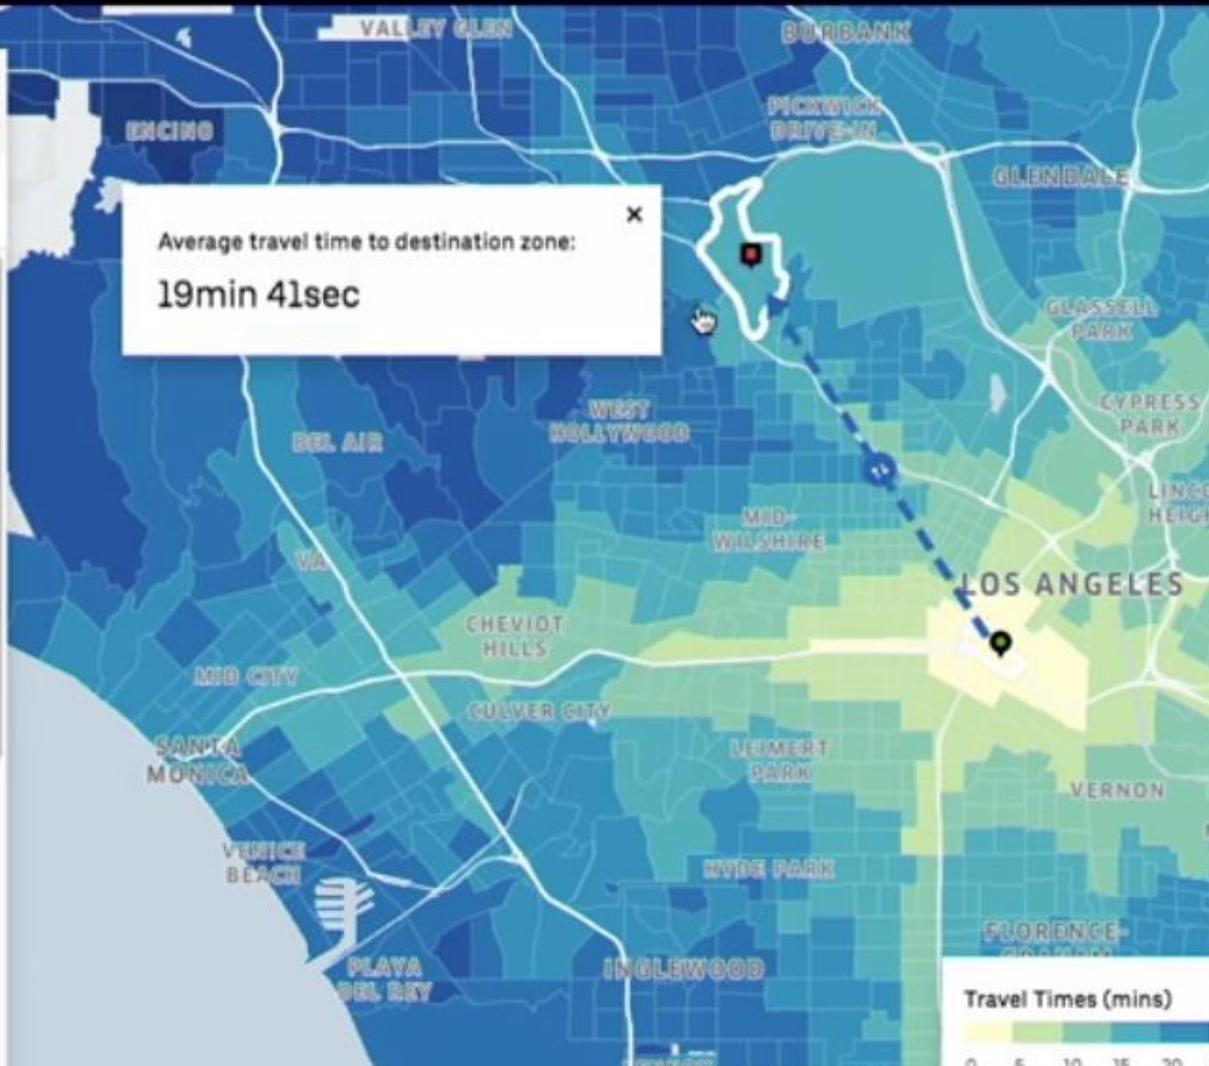
+ Add second date range to compare

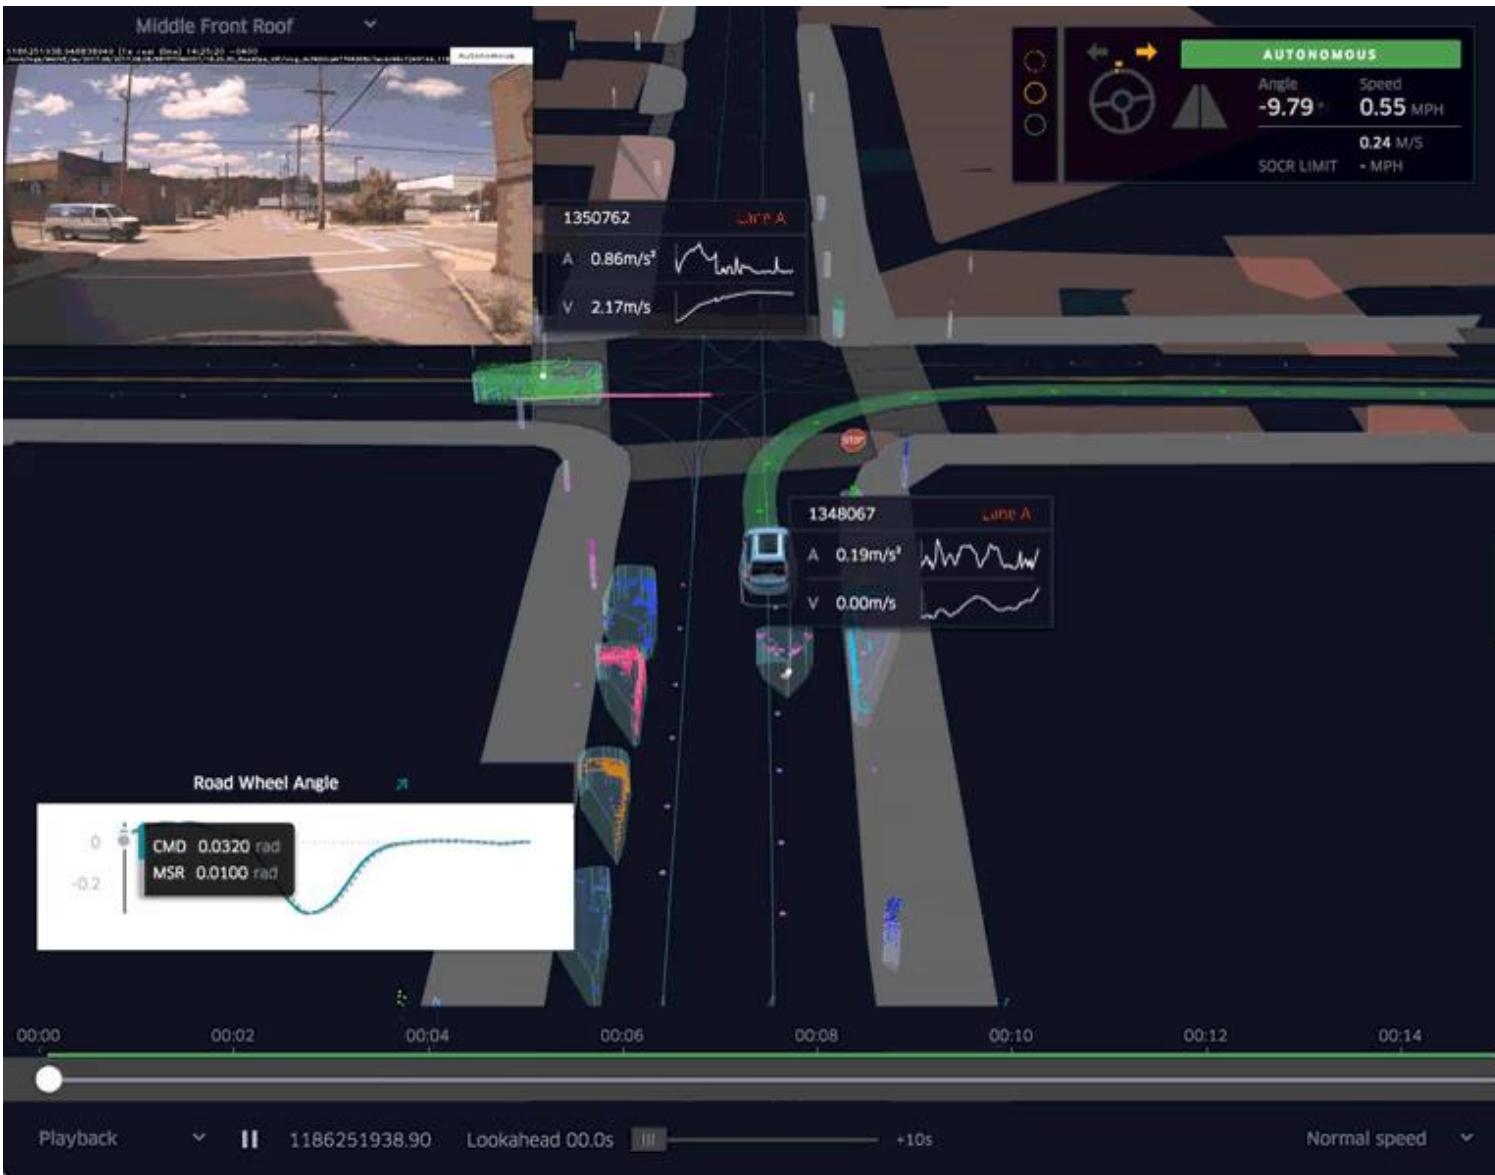
Origin and Destination Zones

Census Tract 224010



Census Tract 189701





Assess the Problem You're Trying to Solve

Is this a type of problem that can be solved using machine learning?

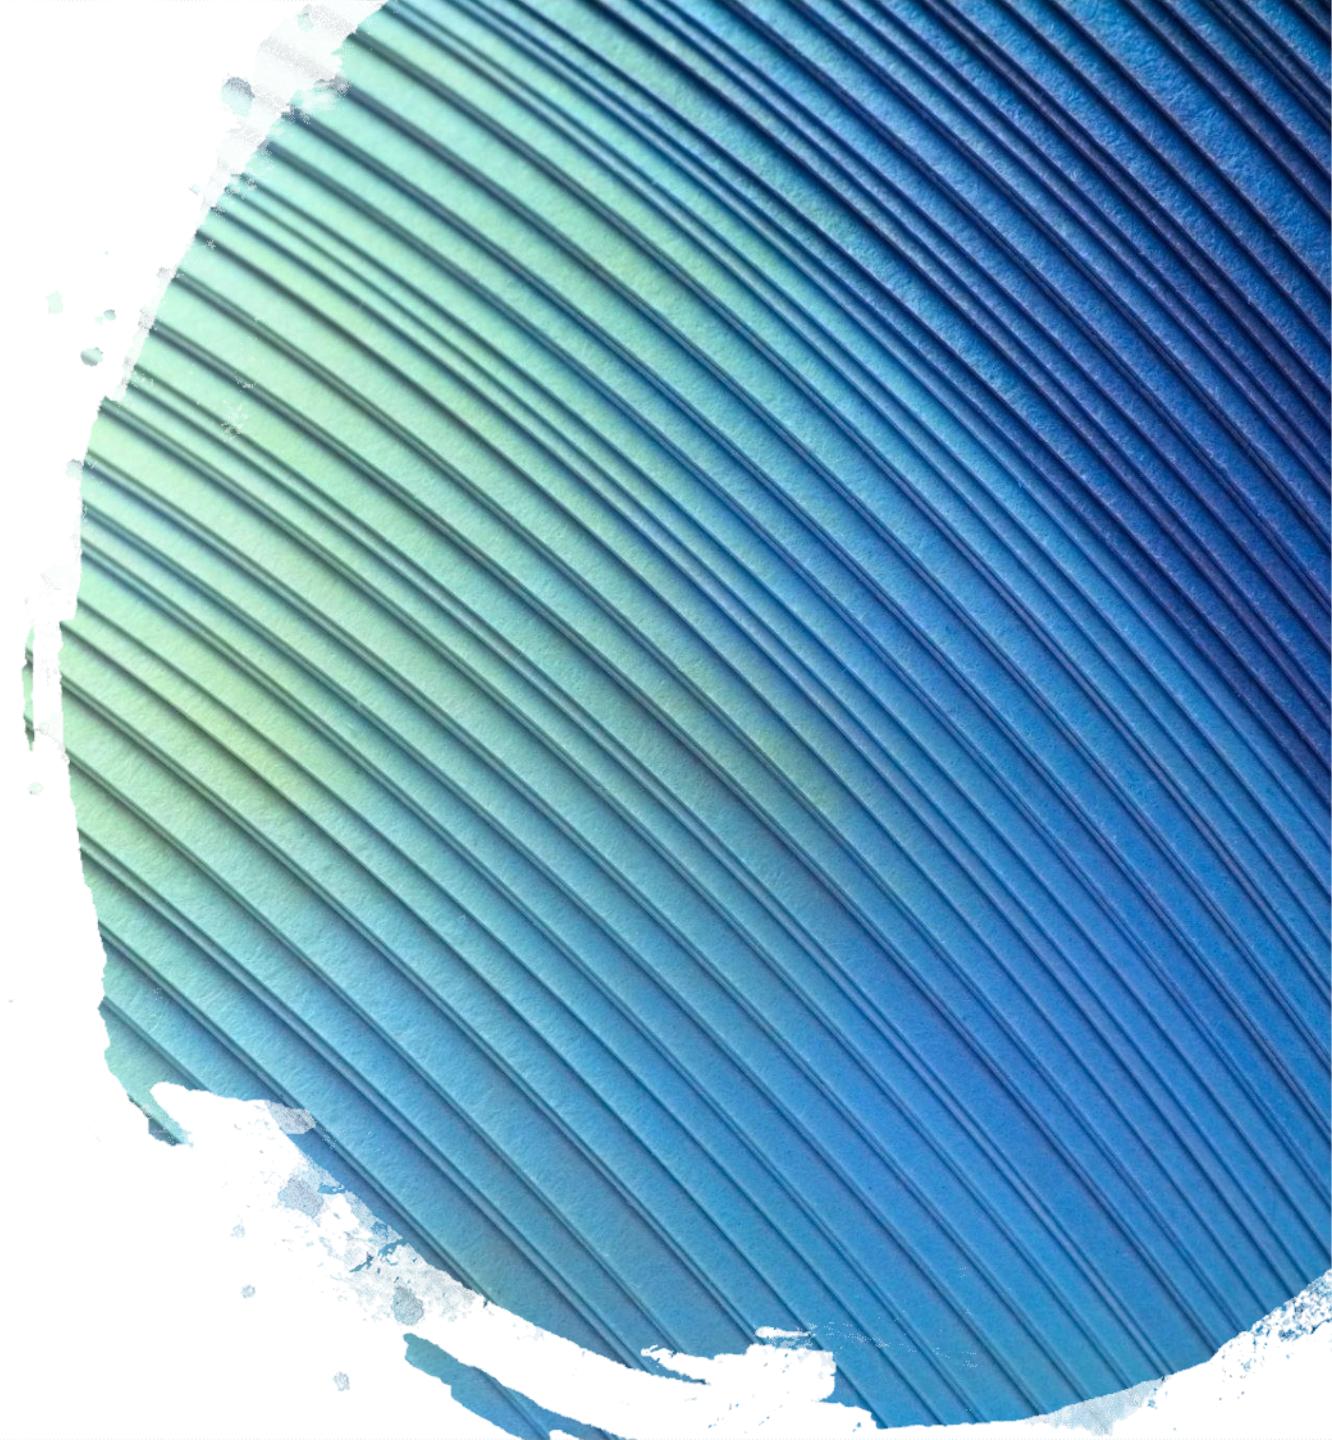
- Does it require a prediction or bucketing something into categories?
- Is the problem relatively self-contained?

Is my particular problem one that can be solved with machine learning?

- Do I have data with labels?
- Do I have the ability to assess the quality of the model?
- Can I determine an acceptable accuracy threshold?

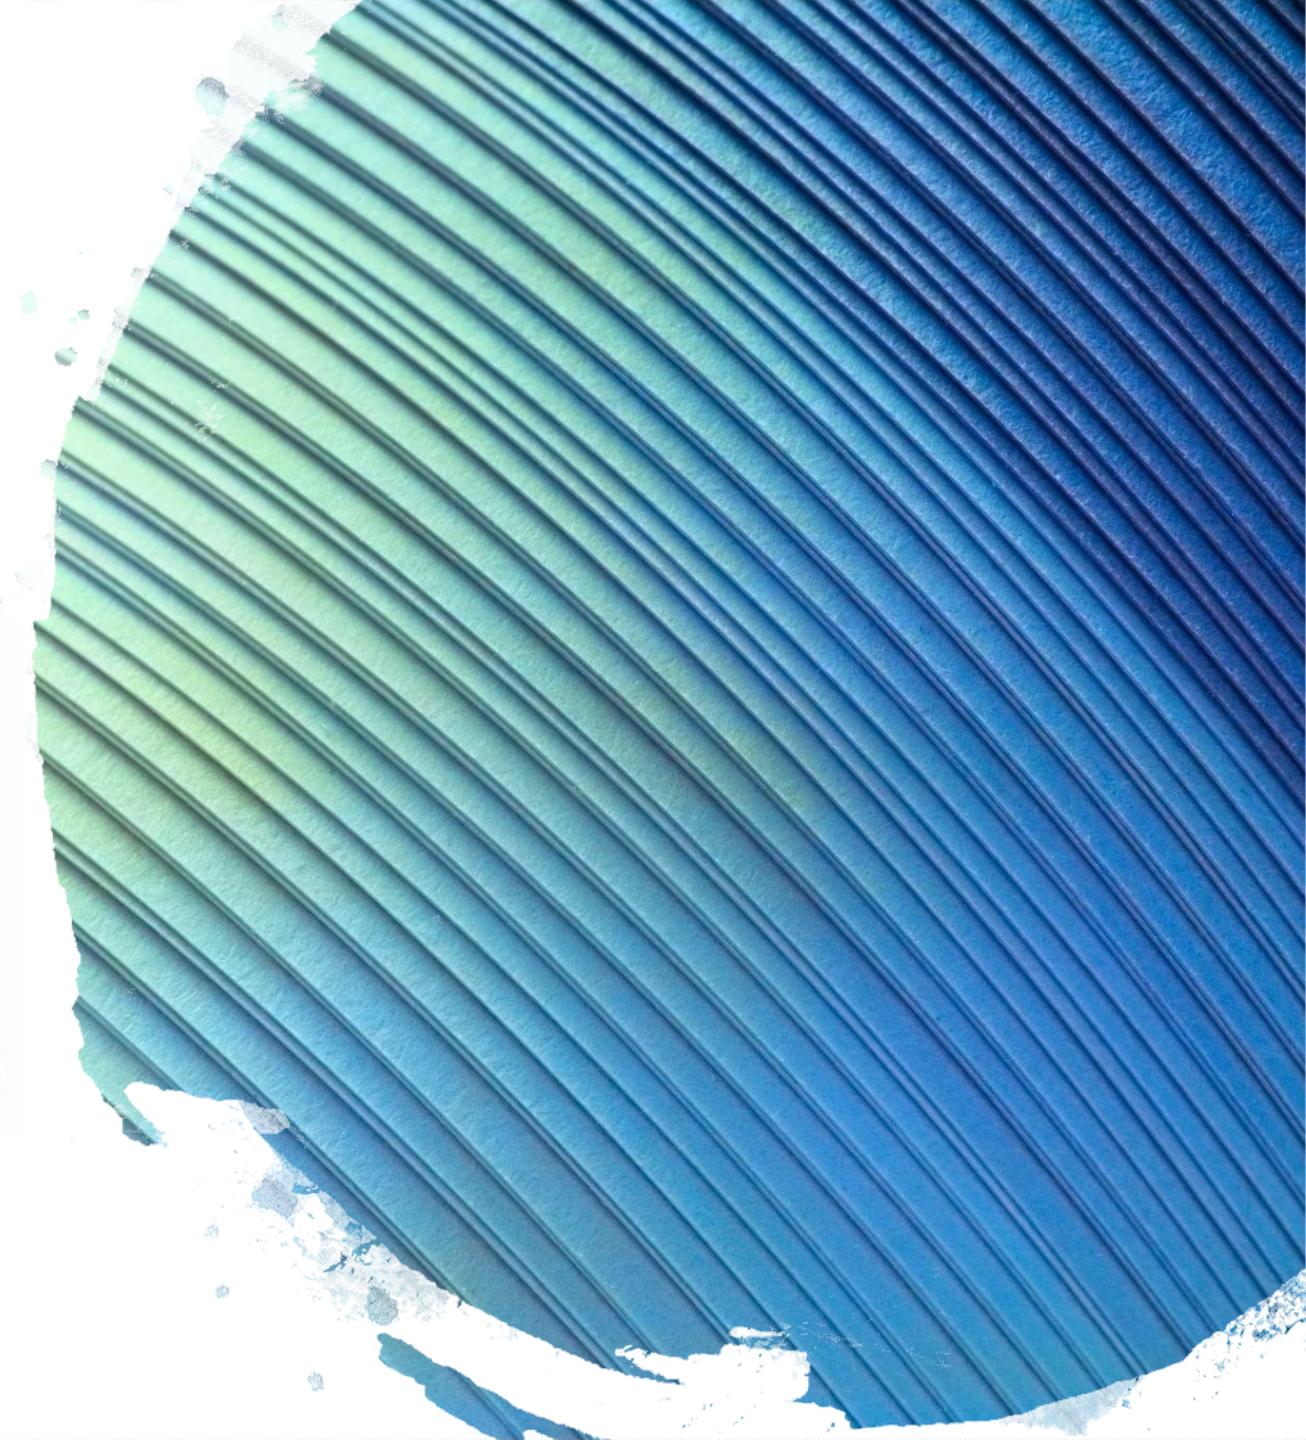
Common Challenges!!!

- Problem Scoping
- Data
- Infrastructure
- Latency



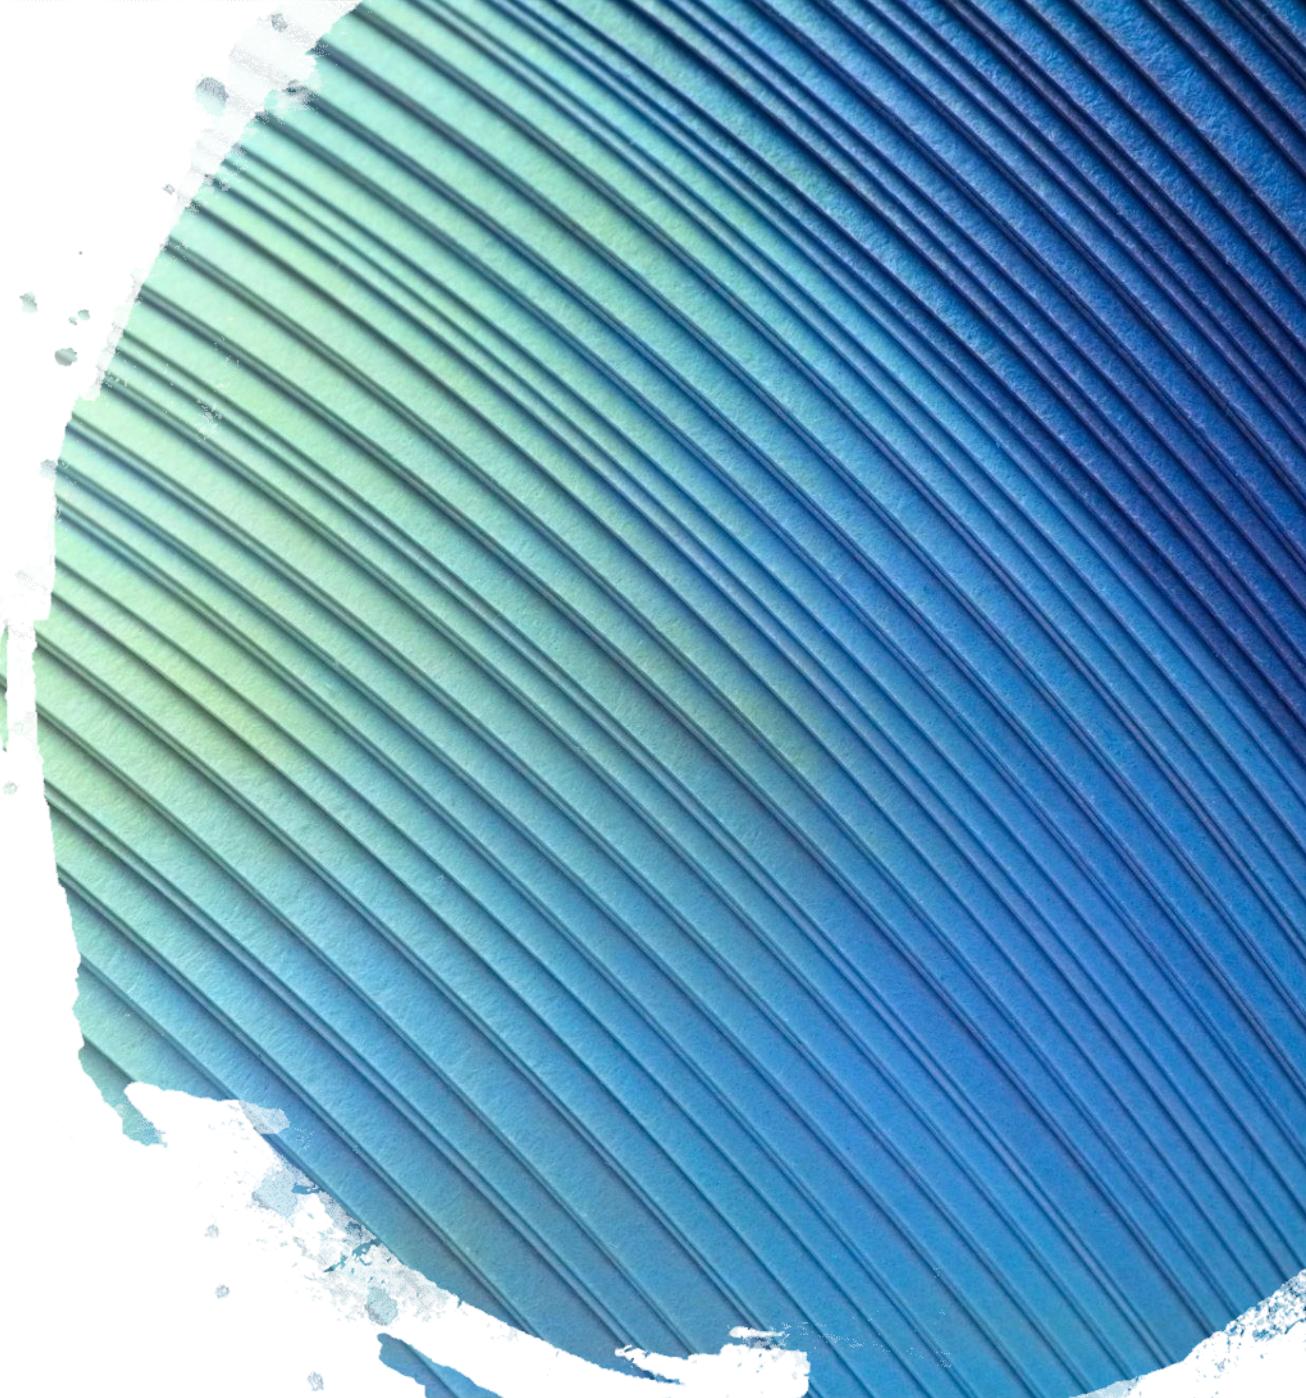
Problem scoping

- The right problem isn't being solved
- Tolerance threshold is undetermined
- Impact in real environment can't be measured



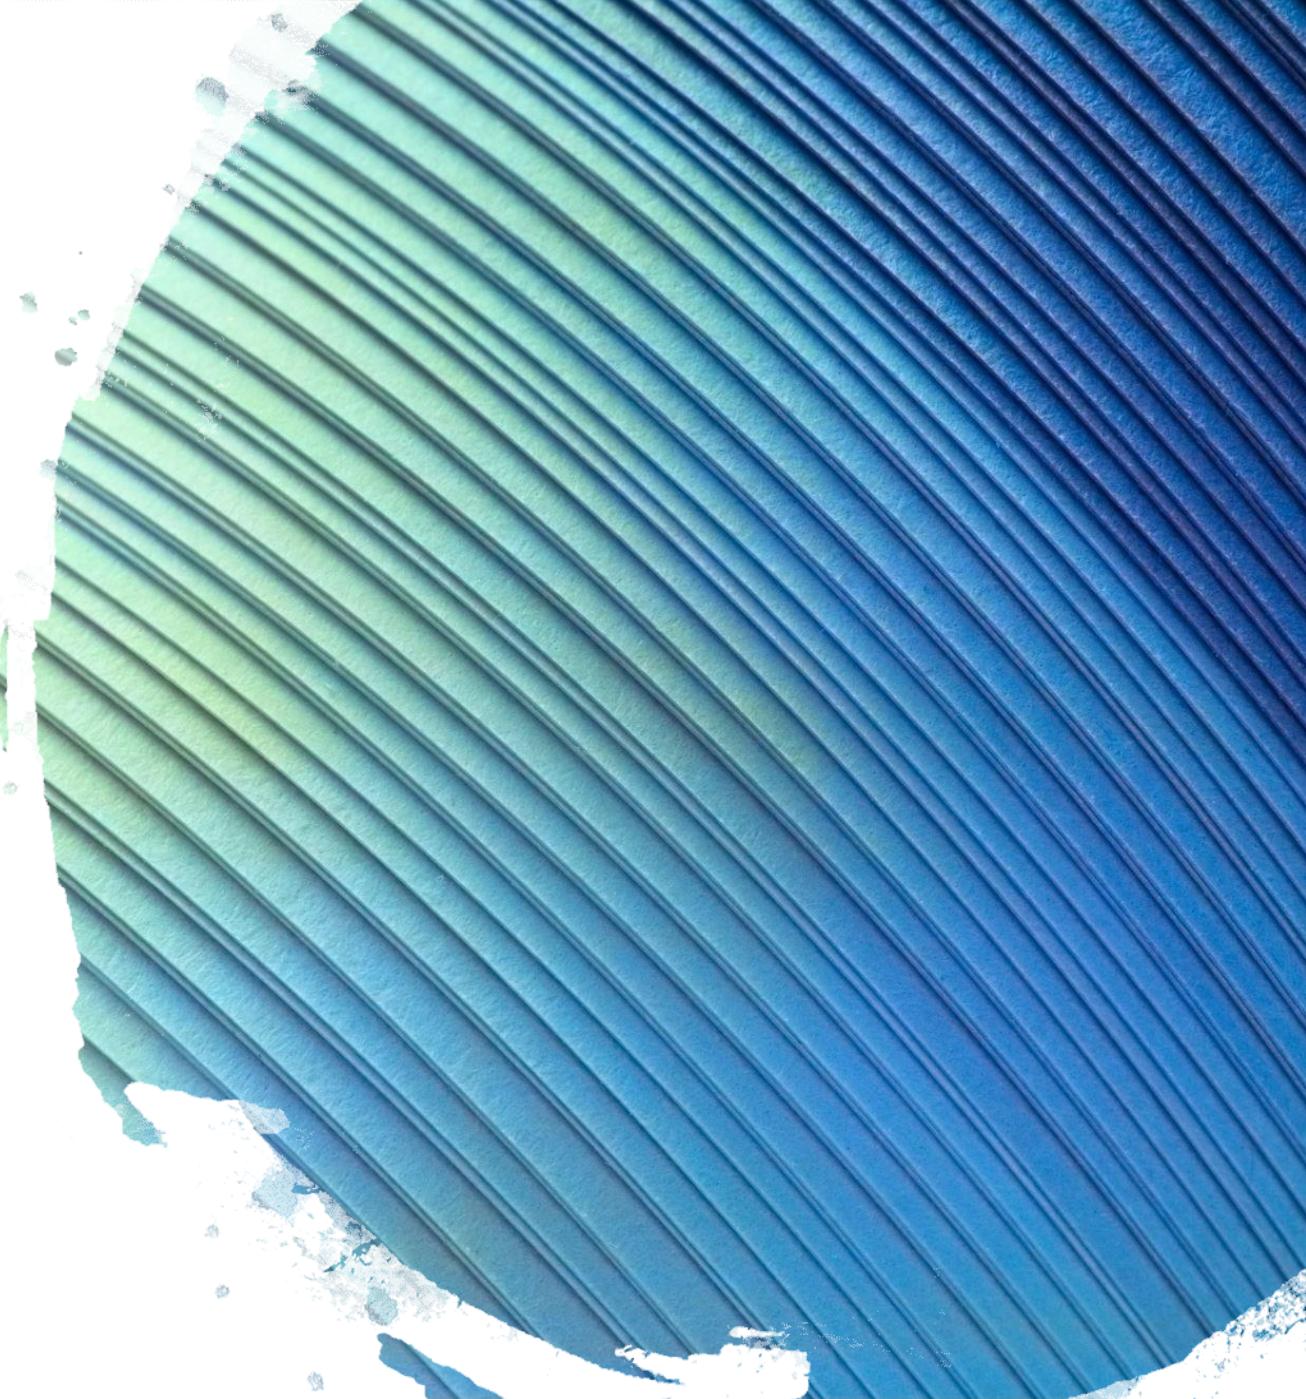
Data

- Lack of data/data sensitivity
- Too much data
- Data doesn't have labels



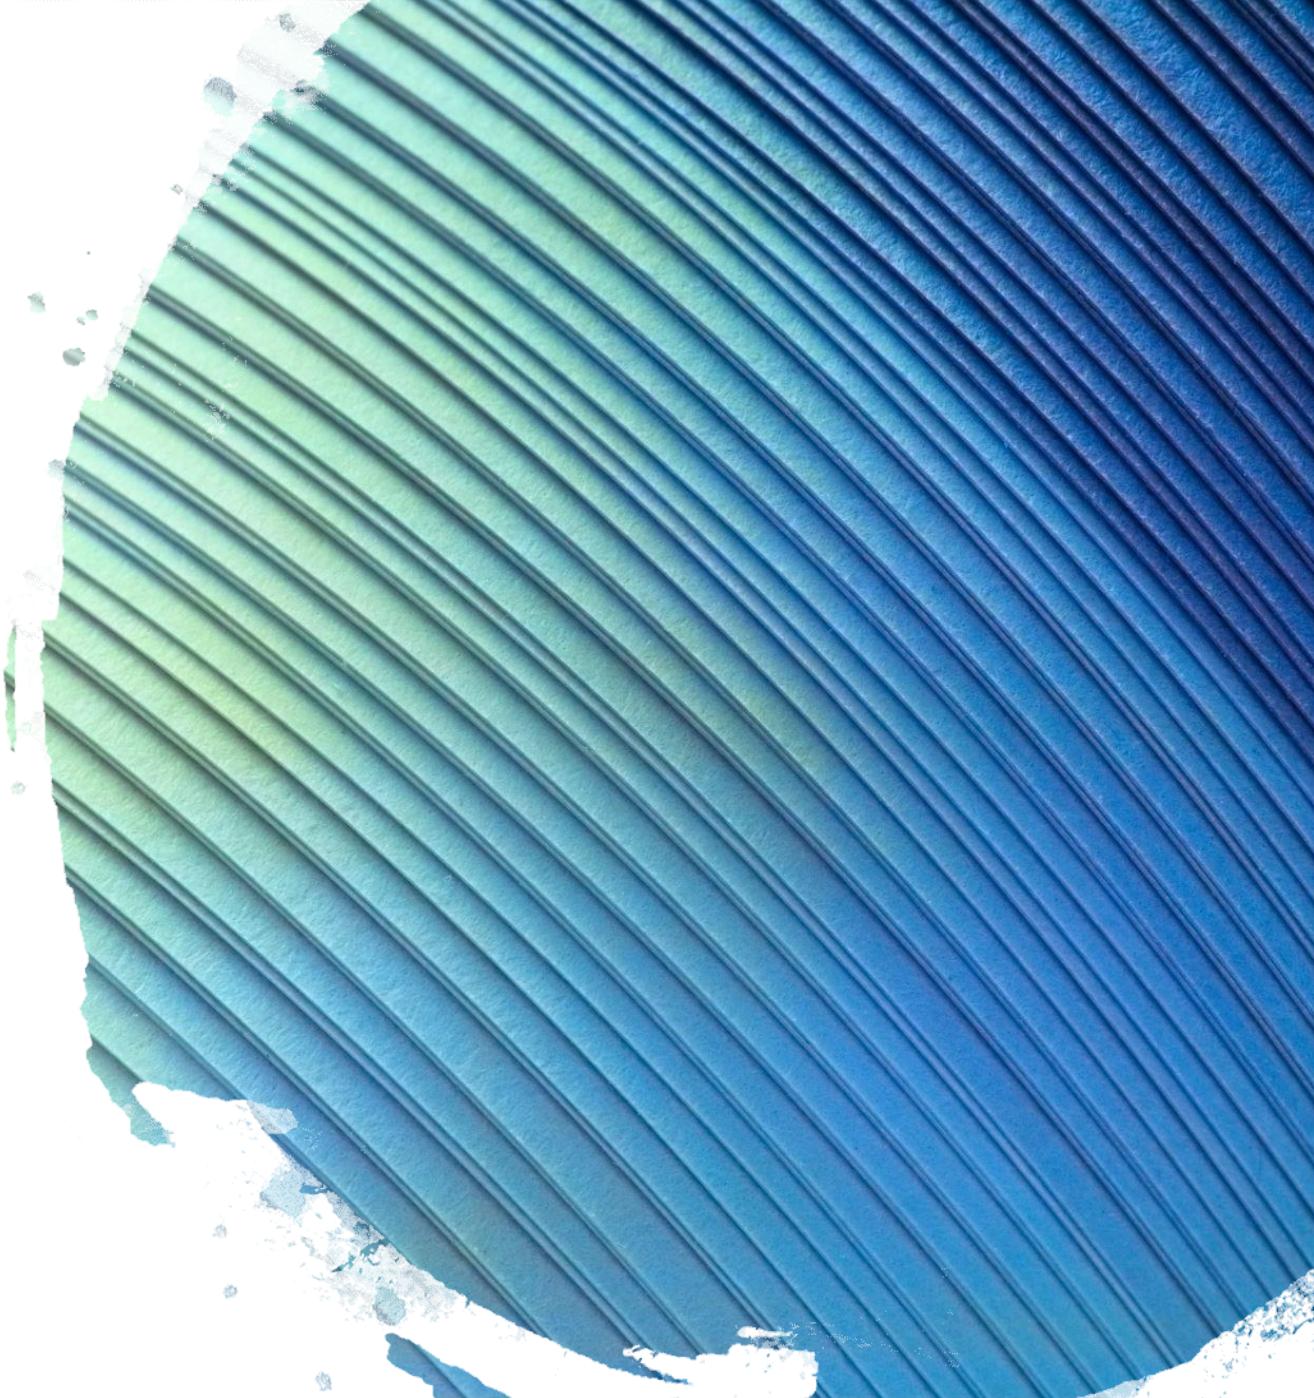
Data

- Lack of data/data sensitivity
- Too much data
- Data doesn't have labels
- Data is too biased, dirty, noisy, incomplete, etc.



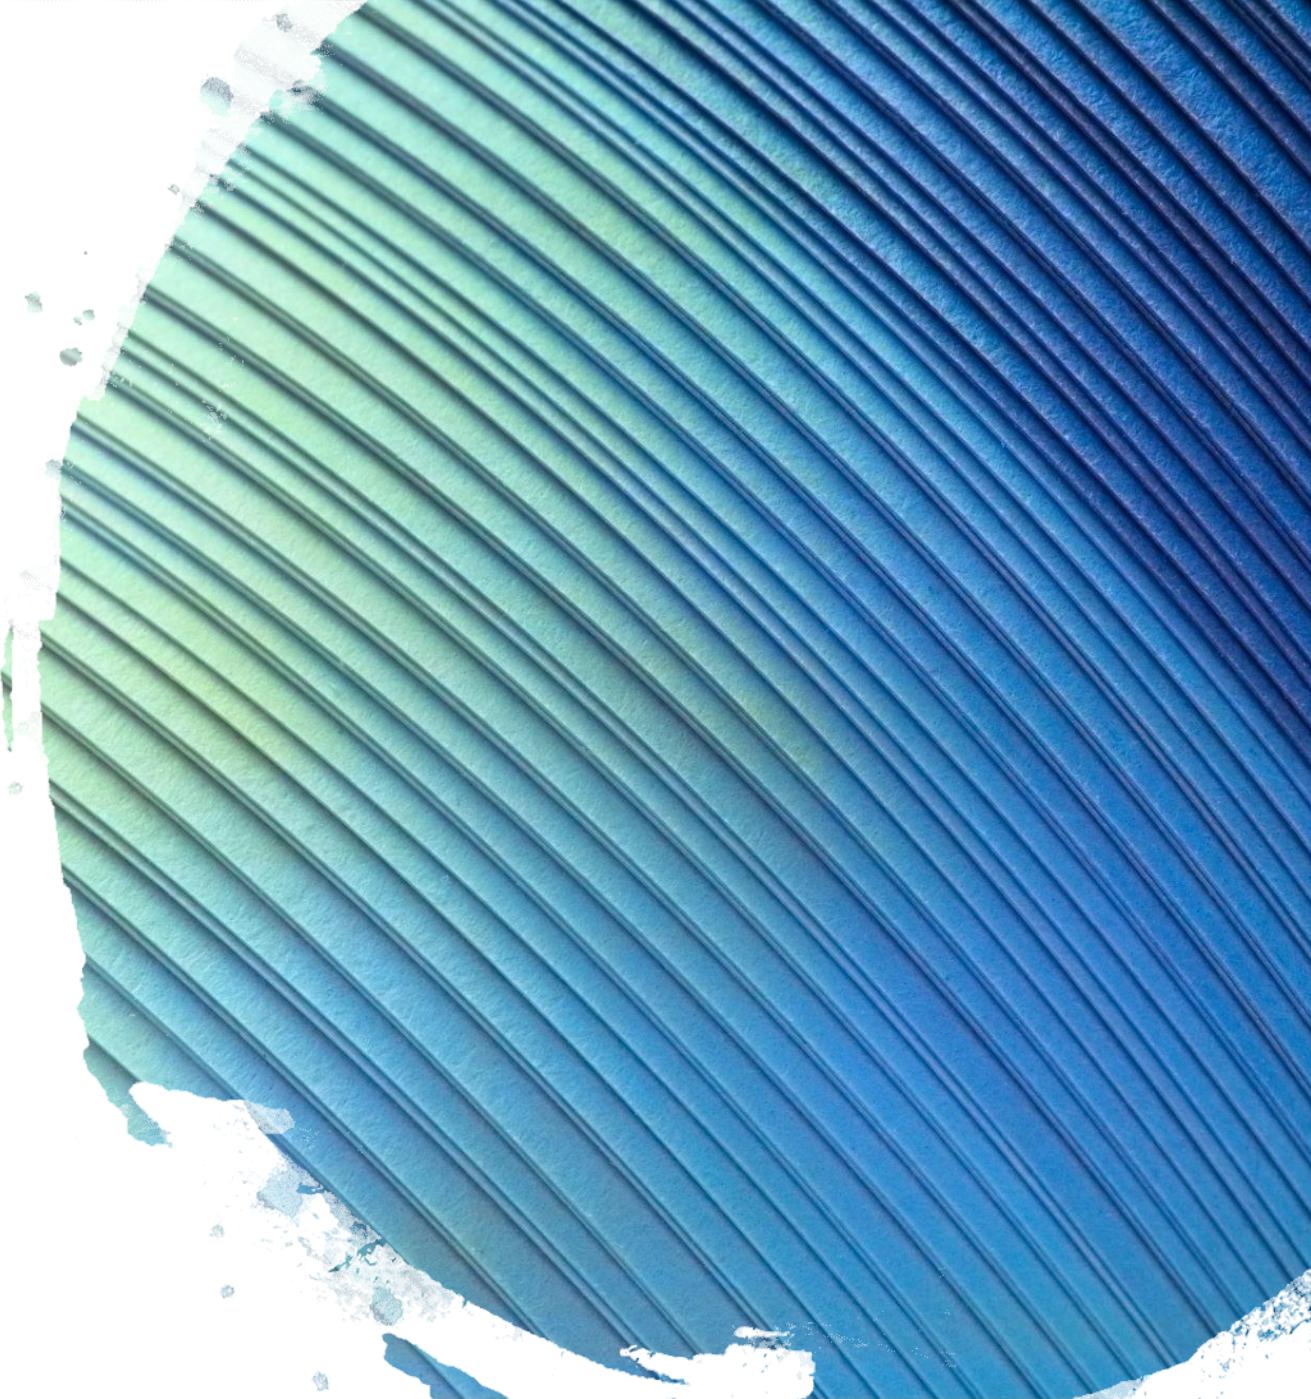
Infrastructure

- Lack skills to robustly automate
- Not enough compute power
- Inability to A/B test with existing solution
- Inability to continuously track model quality



Latency

- Model takes too long to train
- Model takes too long at inference time



Exploratory Data Analysis

Why?

- Understand the shape of the data
- Learn which features might be useful
- Inform the cleaning that will come next

What?

- Counts or distributions of all variables
- Data type for each feature
- Missing data
- Correlations
- Duplicates

Data Cleaning



Why?

- Shape data in a way a model can best pick up on the signal
- Remove irrelevant data
- Adjust features to be acceptable for a model

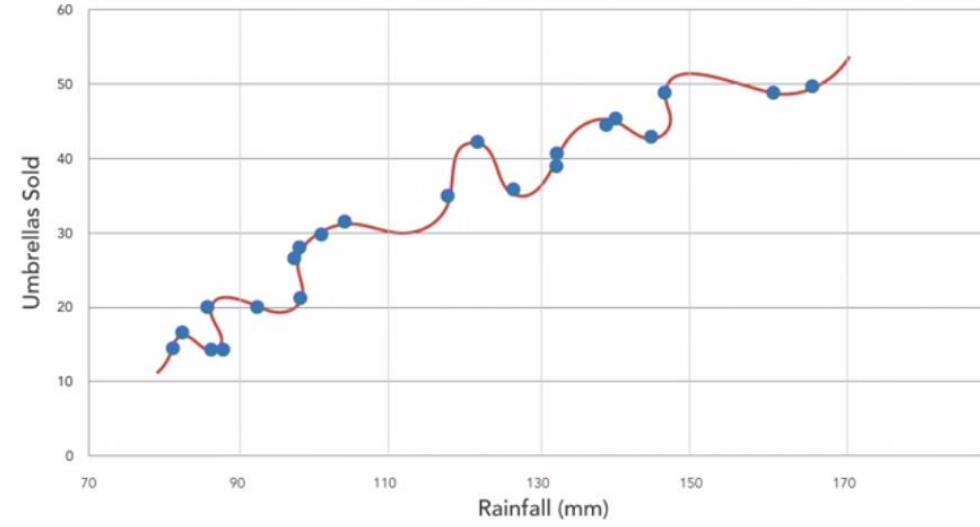
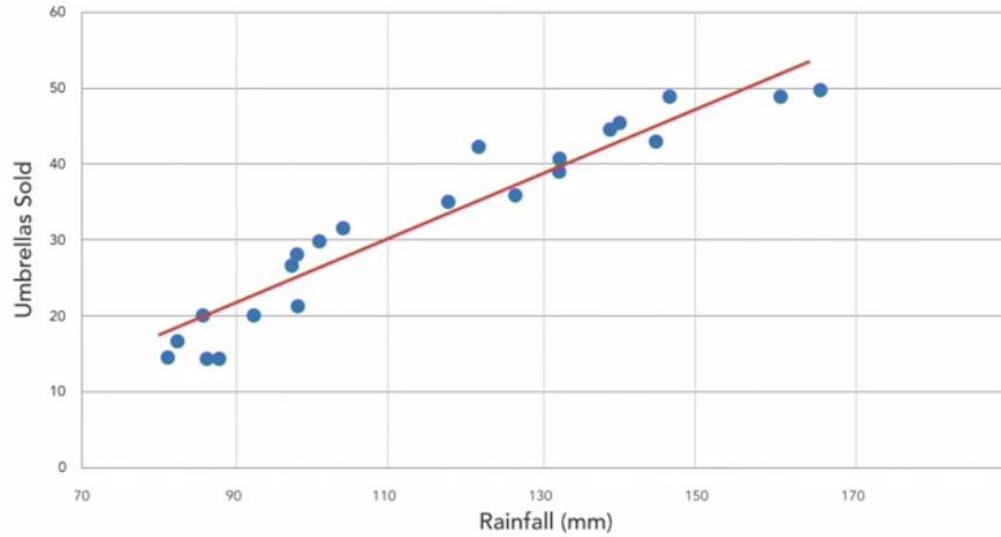
What?

- Anonymize
- Encode categorical variables
- Fill missing data (if necessary)



Fitting a function to examples and
using that function to generalize and
make predictions about new examples

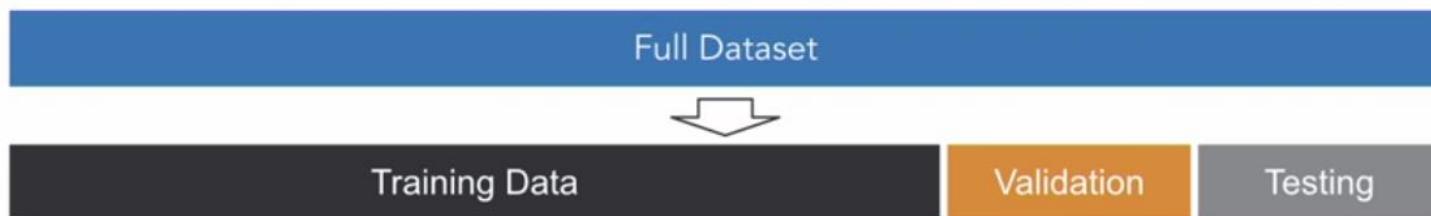


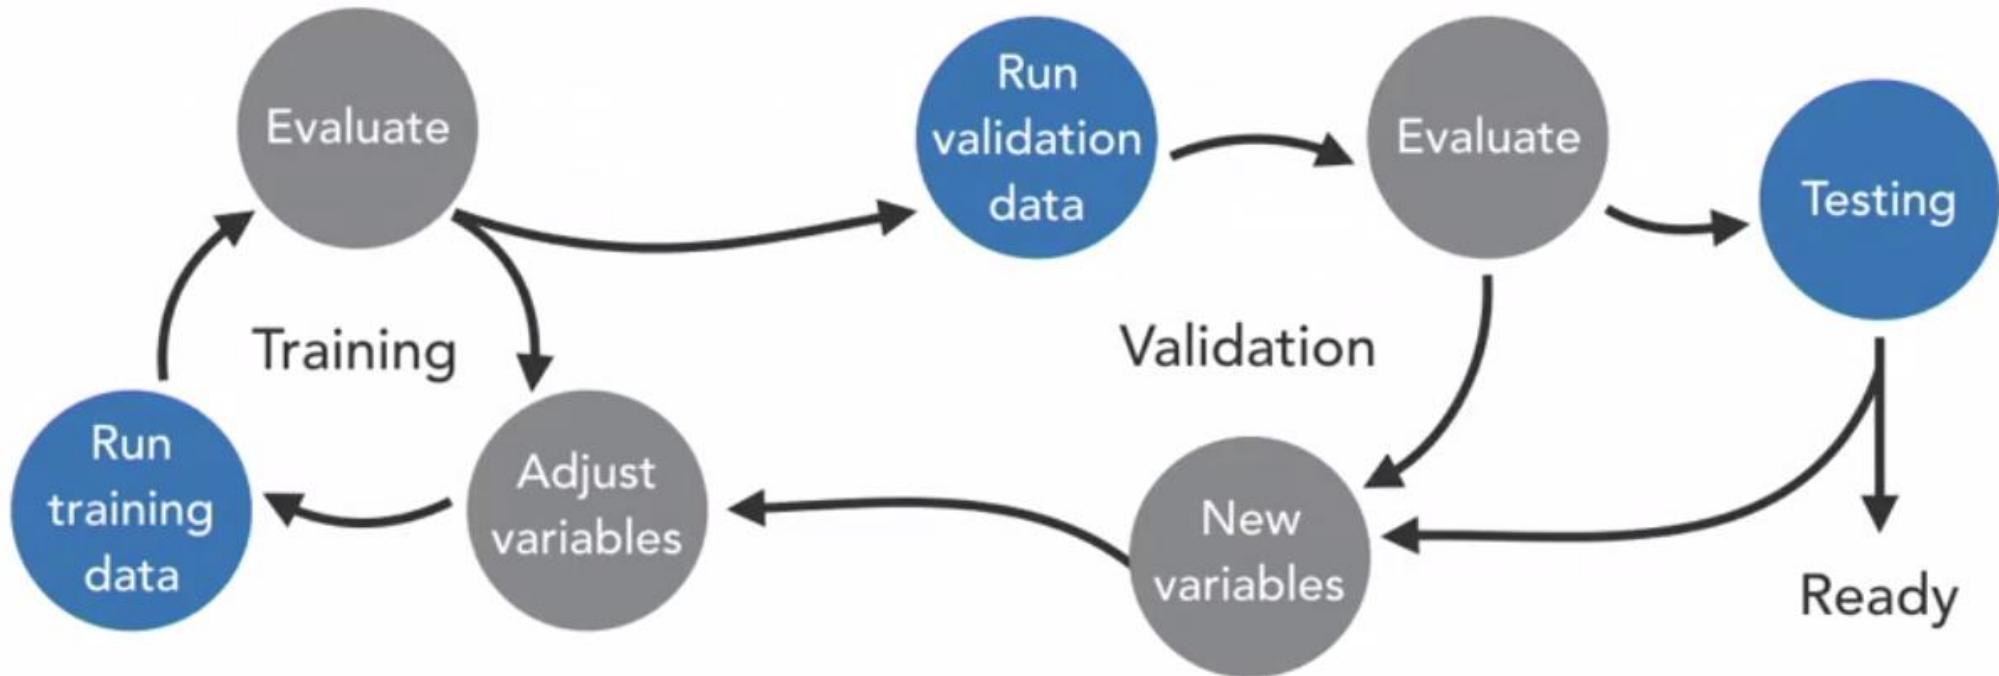


Will This Generalize?

How Do We Split Up Our Data?

- Training dataset – data used to train the model (allow the algorithm to learn from this data)
- Validation dataset – data used to select the best model (optimal algorithm and hyperparameter settings)
- Test dataset – data used to provide an unbiased evaluation of what the model will look like in its real environment





What if we don't split the dataset?

- Overfitting or underfitting to the data
- Inaccurate representation of how the model will generalize

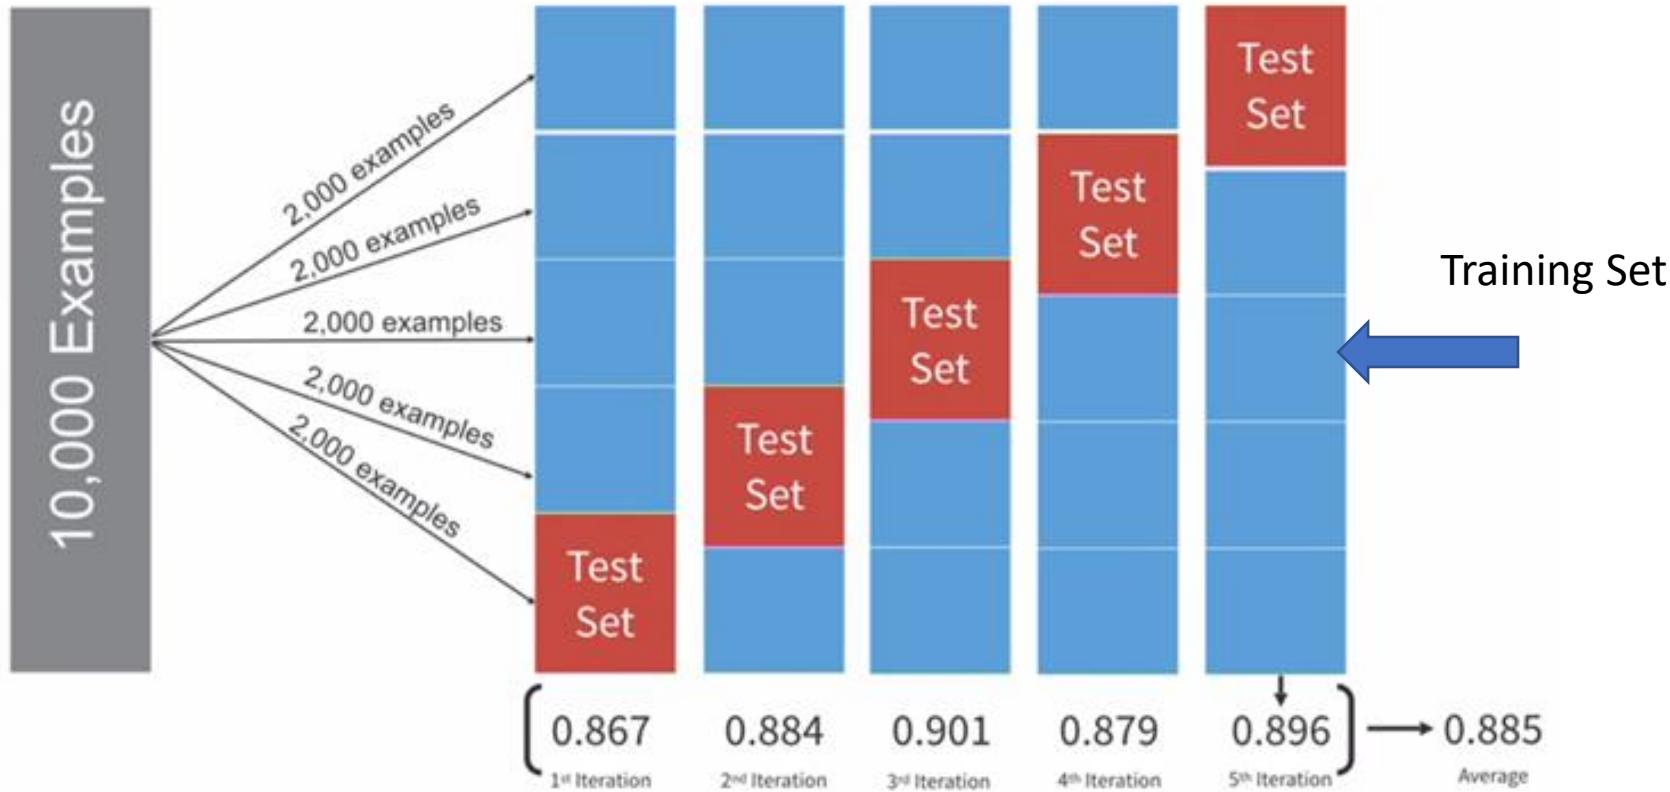
Holdout Test Set

Sample of data not used in fitting a model; used to evaluate the model's ability to generalize to unseen data

K-Fold Cross Validation

Data is divided into k subsets and the holdout method is repeated k times. Each time, one of the k subsets is used as the test set and the other $k-1$ subsets are combined to be used to train the model.

K-Fold Cross Validation Process

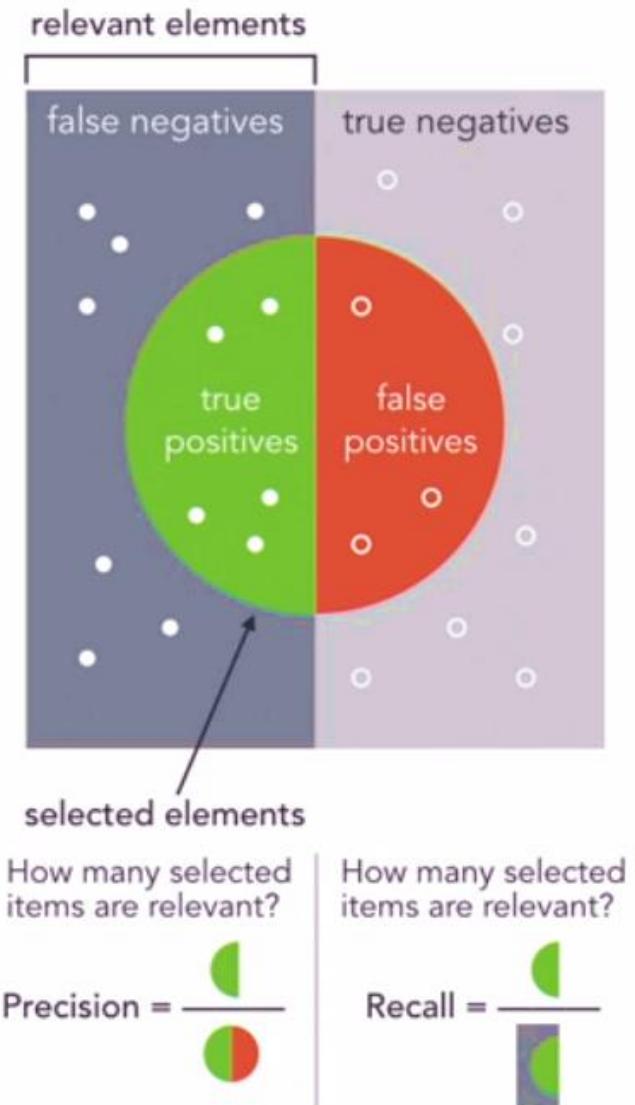


Evaluation Framework

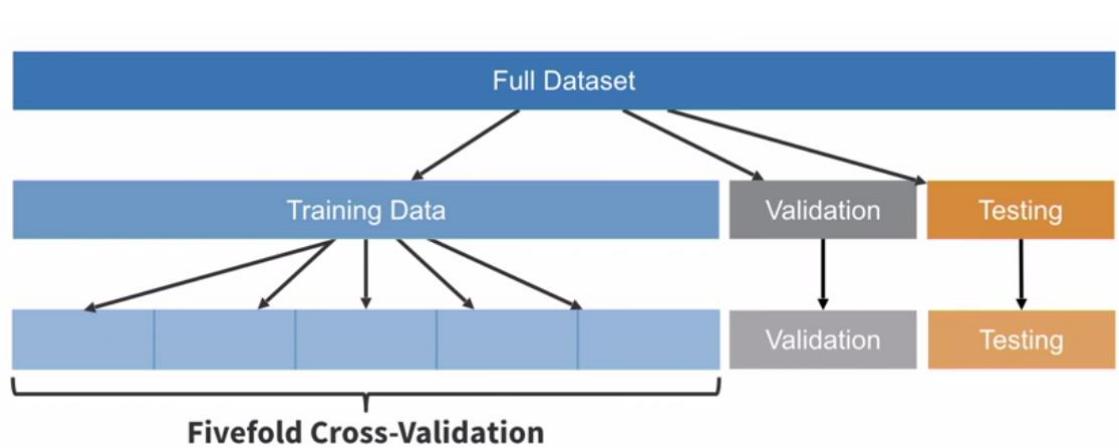
1. Evaluation metrics – how are we gauging the accuracy of the model?
2. Process (how to split the data) – how do we leverage our full dataset to mitigate likelihood of overfitting or underfitting?

Evaluation Metrics

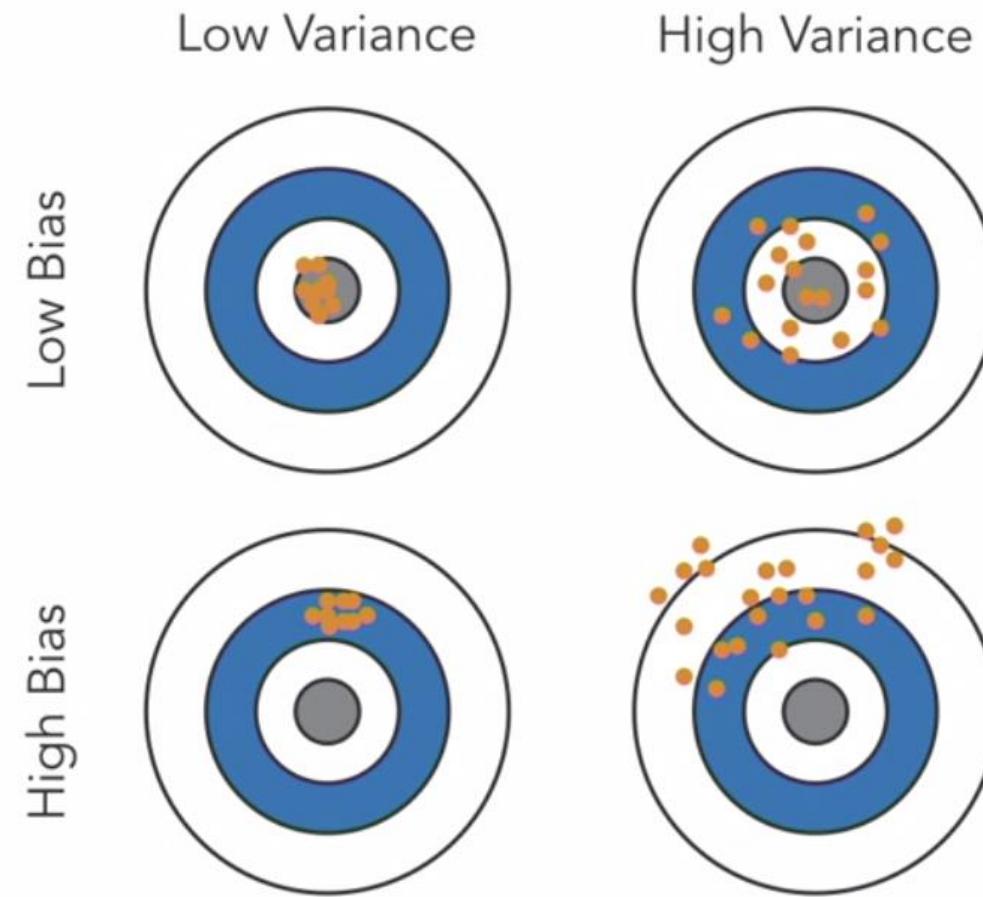
- Accuracy =
$$\frac{\text{\# predicted correctly}}{\text{total \# of examples}}$$
- Precision =
$$\frac{\text{\# predicted as surviving that actually survived}}{\text{total \# predicted to survive}}$$
- Recall =
$$\frac{\text{\# predicted as surviving that actually survived}}{\text{total \# that actually survived}}$$

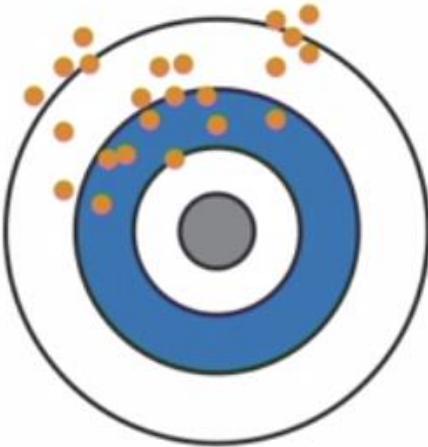
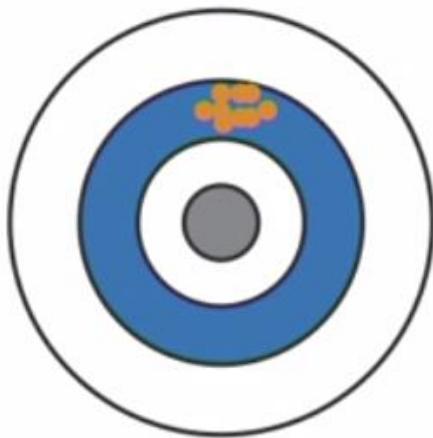


Process



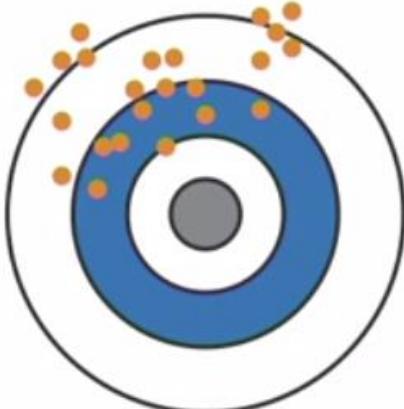
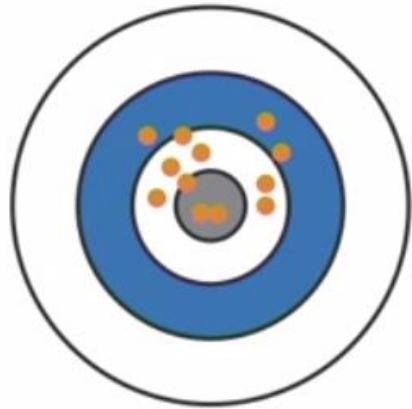
1. Run fivefold cross-validation and select the best models.
2. Re-fit models on full training set, evaluate those models on the validation set and pick the best one.
3. Evaluate that best model on the test set to gauge its ability to generalize to unseen data.





Bias is the algorithm's tendency to consistently learn the wrong thing by not taking into account all the information in the data.

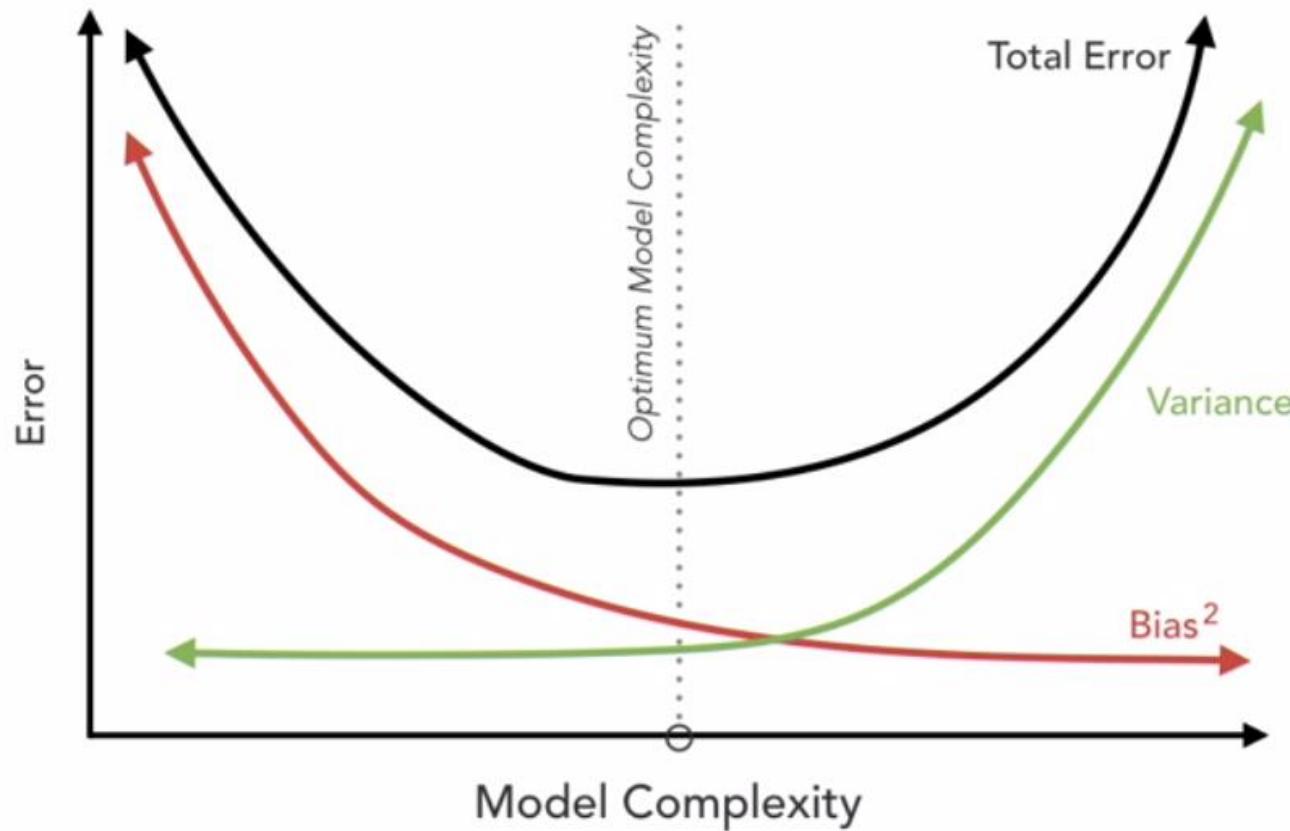
High bias is a result of the algorithm missing the relevant relations between features and target outputs.



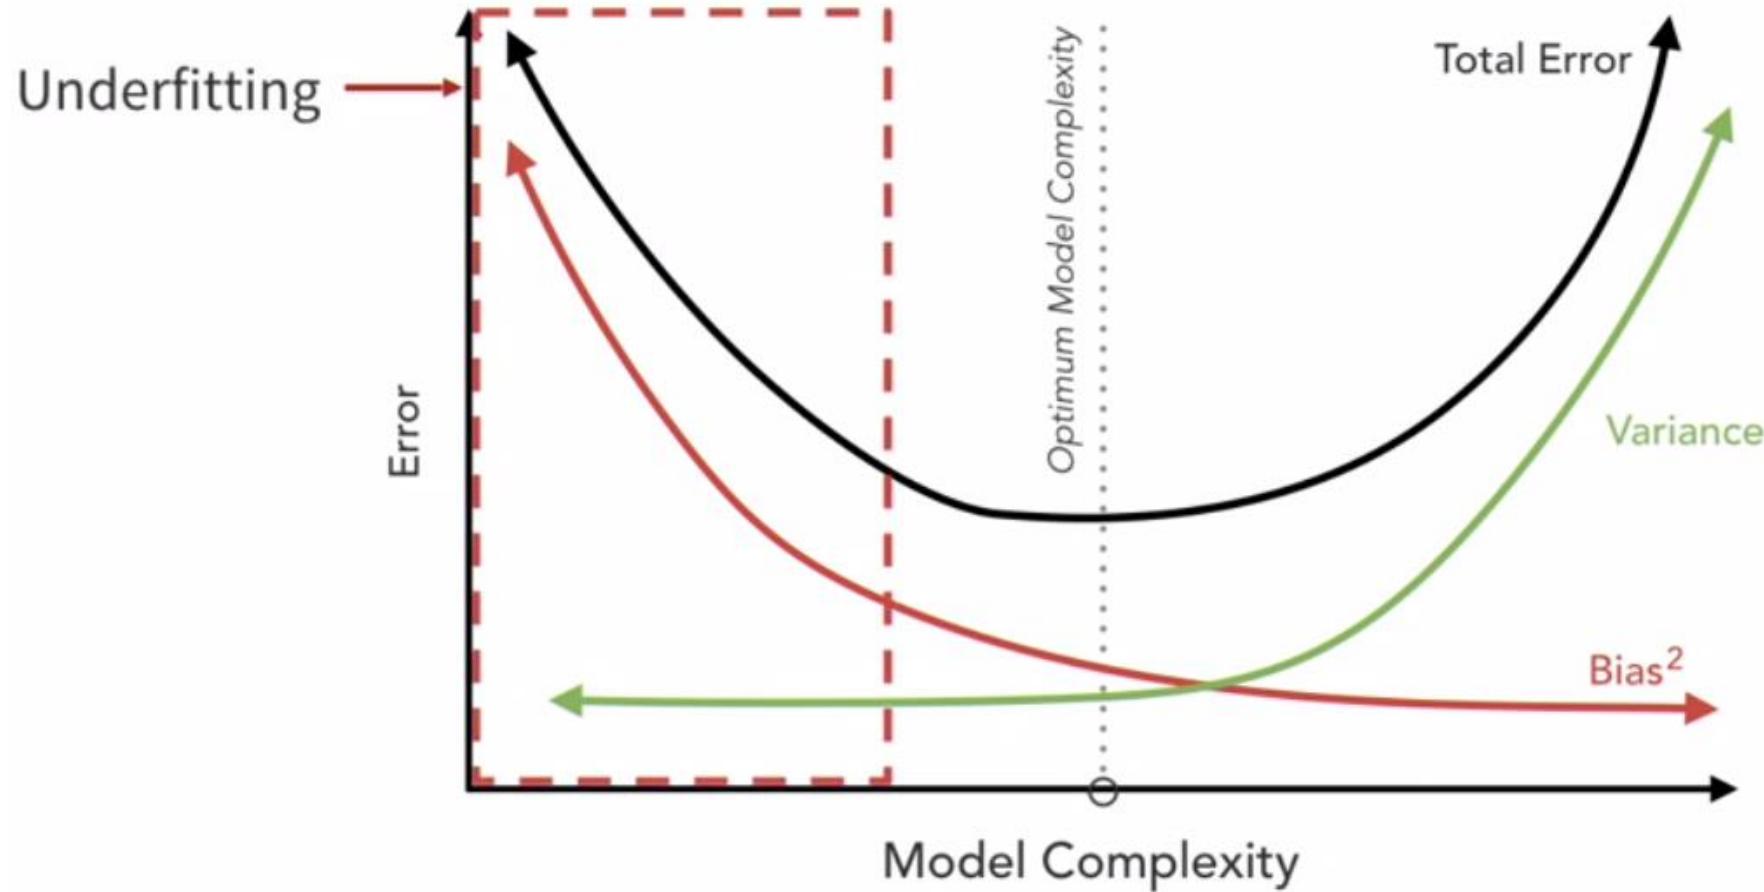
Variance refers to an algorithm's sensitivity to small fluctuations in the training set.

High variance is a result of the algorithm fitting to random noise in the training data.

Total Error = (Bias + Variance) + Irreducible Error



Total Error = (Bias + Variance) + Irreducible Error

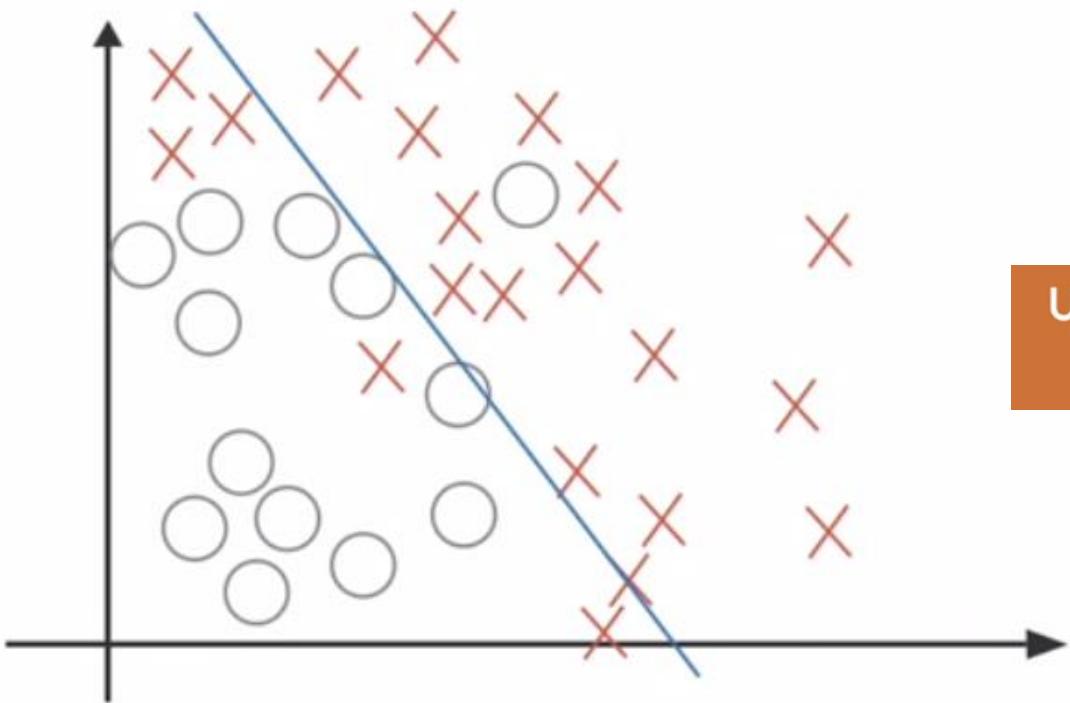




Bias is the algorithm's tendency to consistently learn the wrong thing by not taking into account all the information in the data.

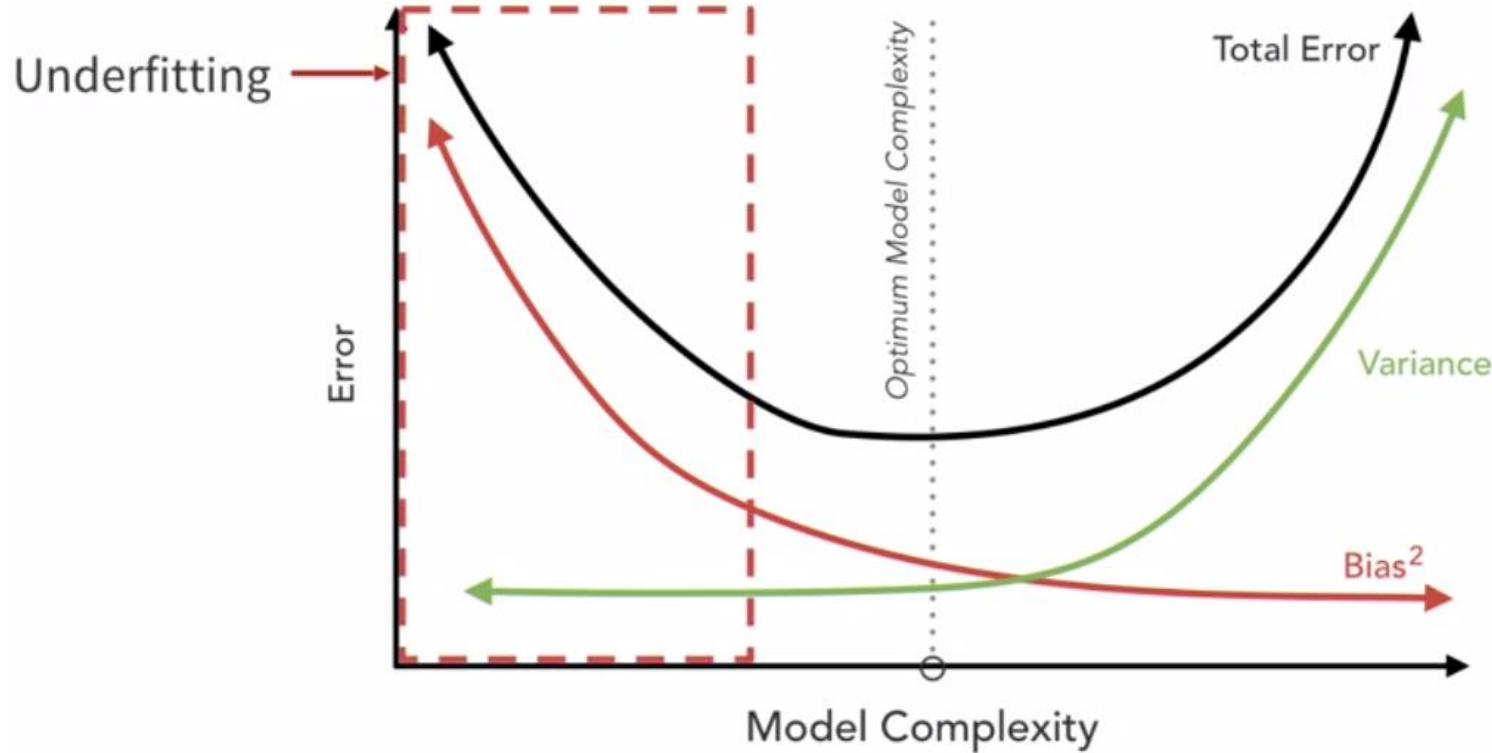
High bias is a result of the algorithm missing the relevant relations between features and target outputs.

Underfitting occurs when an algorithm cannot capture the underlying trend of the data.



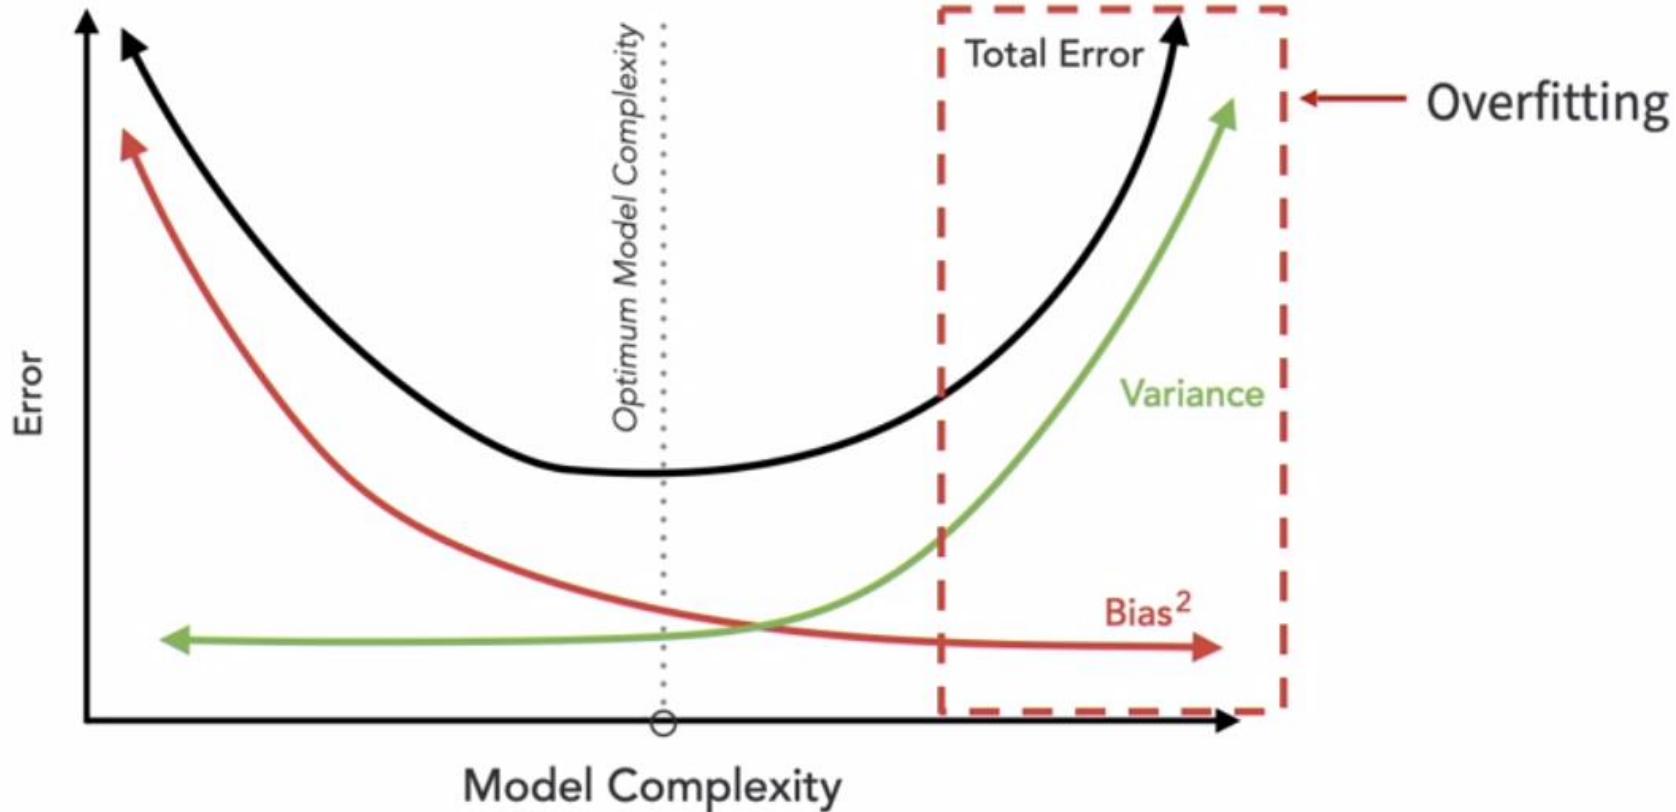
Underfitting occurs when an algorithm cannot capture the underlying trend of the data.

Total Error = (Bias + Variance) + Irreducible Error



Underfitting occurs when an algorithm cannot capture the underlying trend of the data.

Total Error = (Bias + Variance) + Irreducible Error

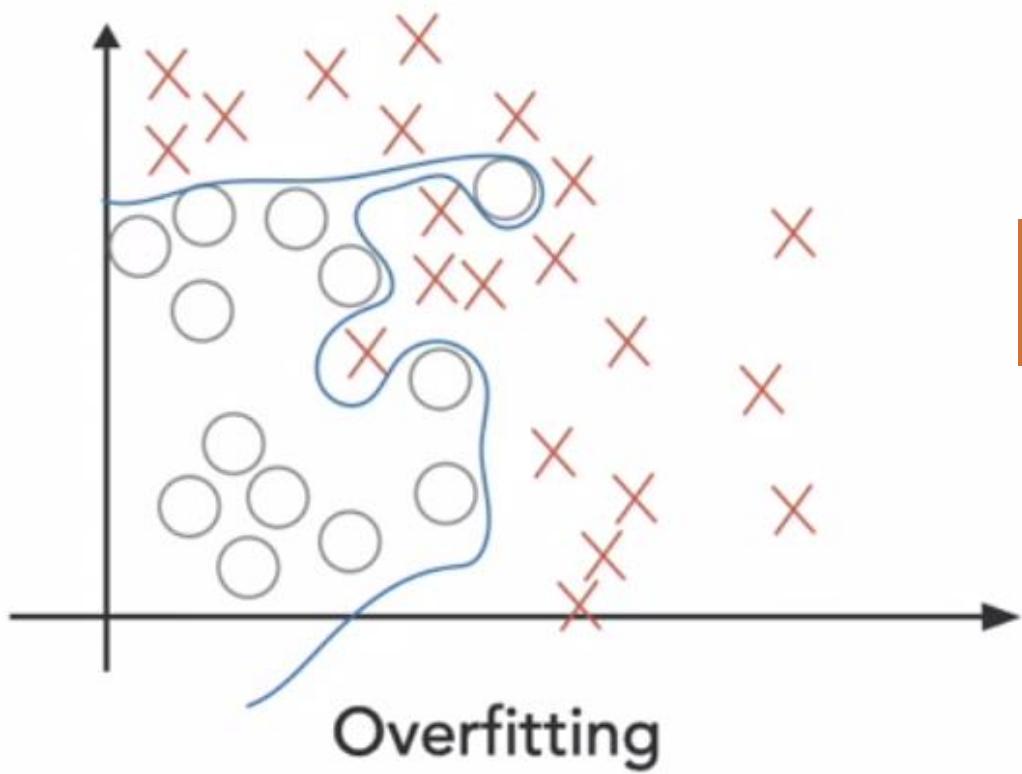




Variance refers to an algorithm's sensitivity to small fluctuations in the training set.

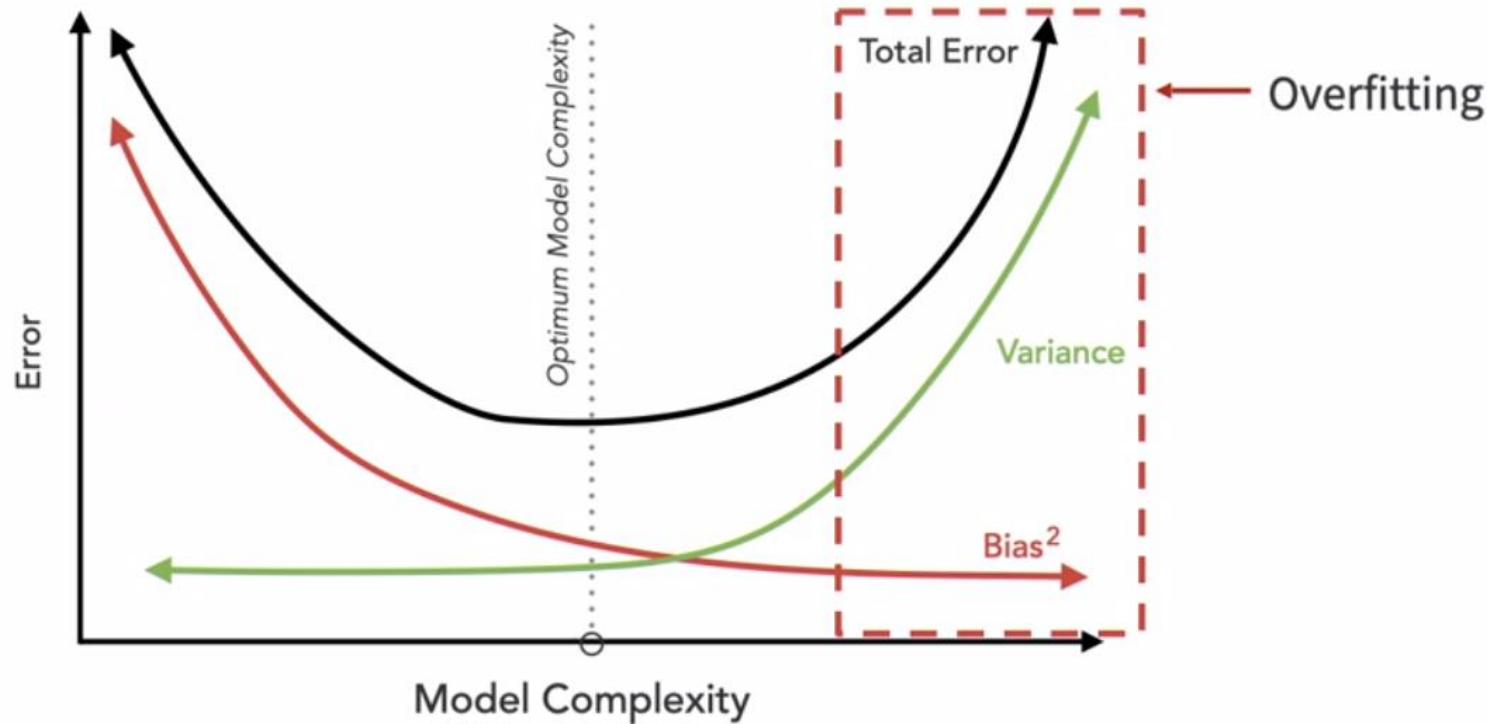
High variance is a result of the algorithm fitting to random noise in the training data.

Overfitting occurs when an algorithm fits too closely to a limited set of data.



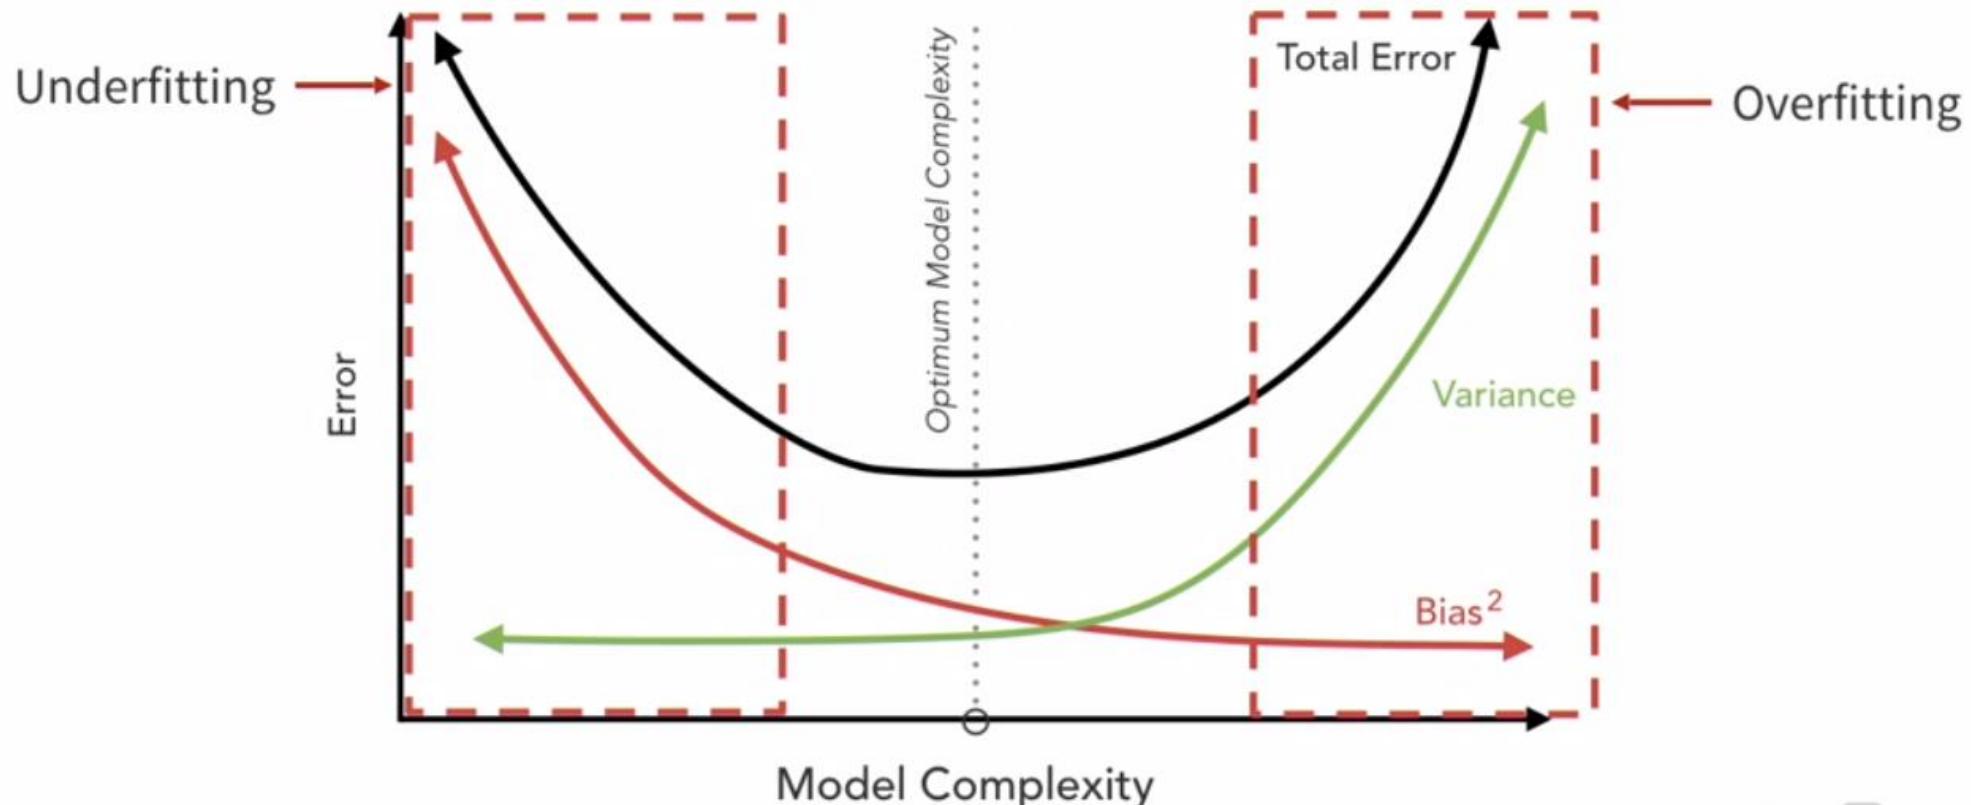
Overfitting occurs when an algorithm fits too closely to a limited set of data.

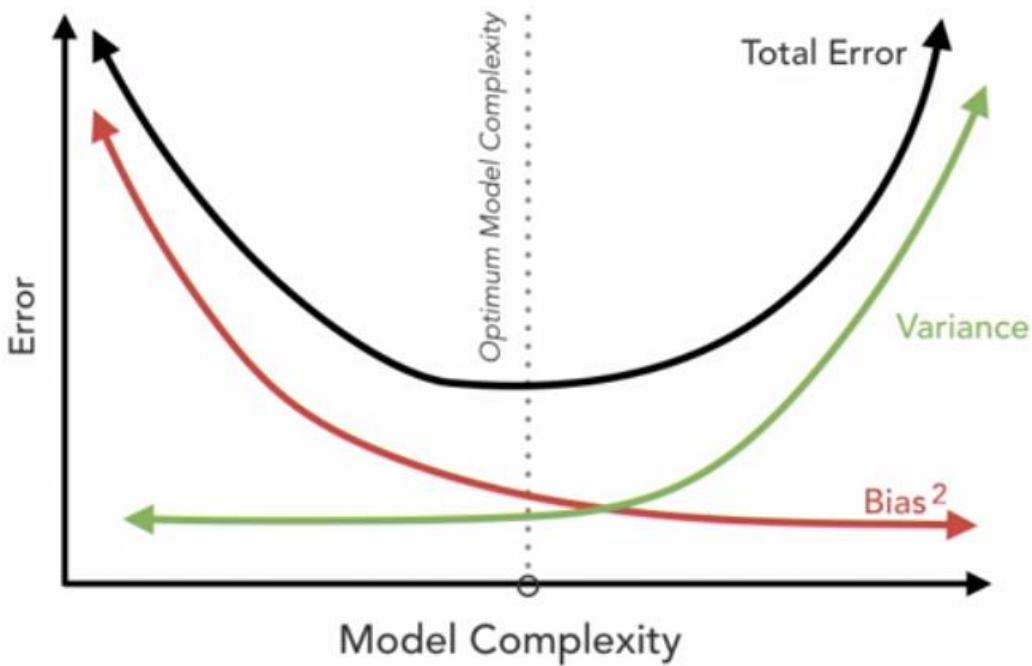
Total Error = (Bias + Variance) + Irreducible Error

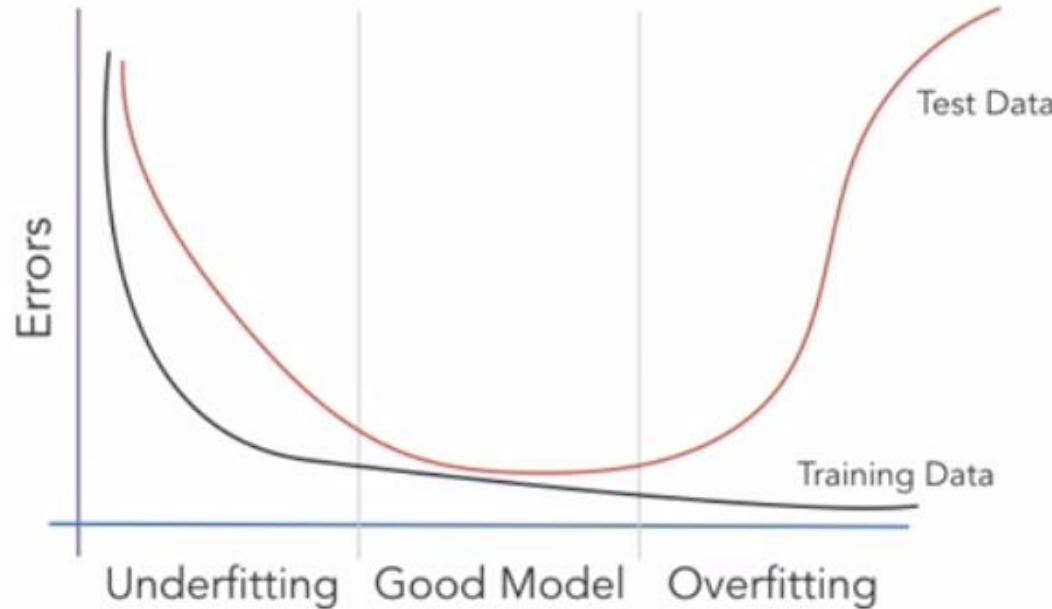


Overfitting occurs when an algorithm fits too closely to a limited set of data.

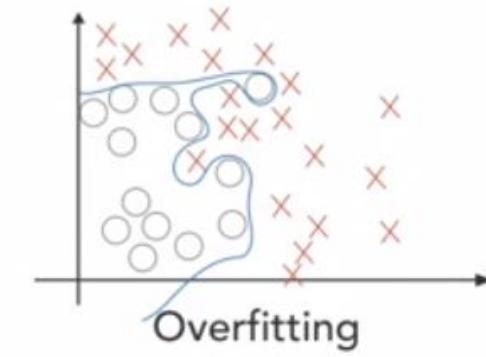
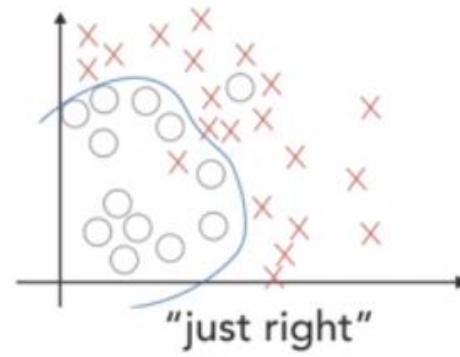
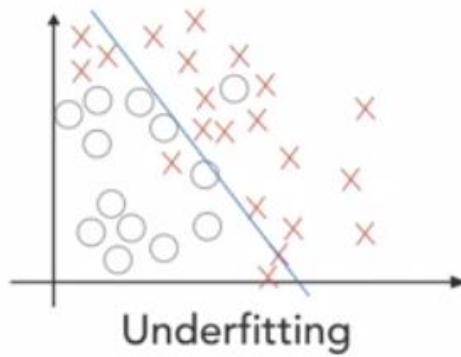
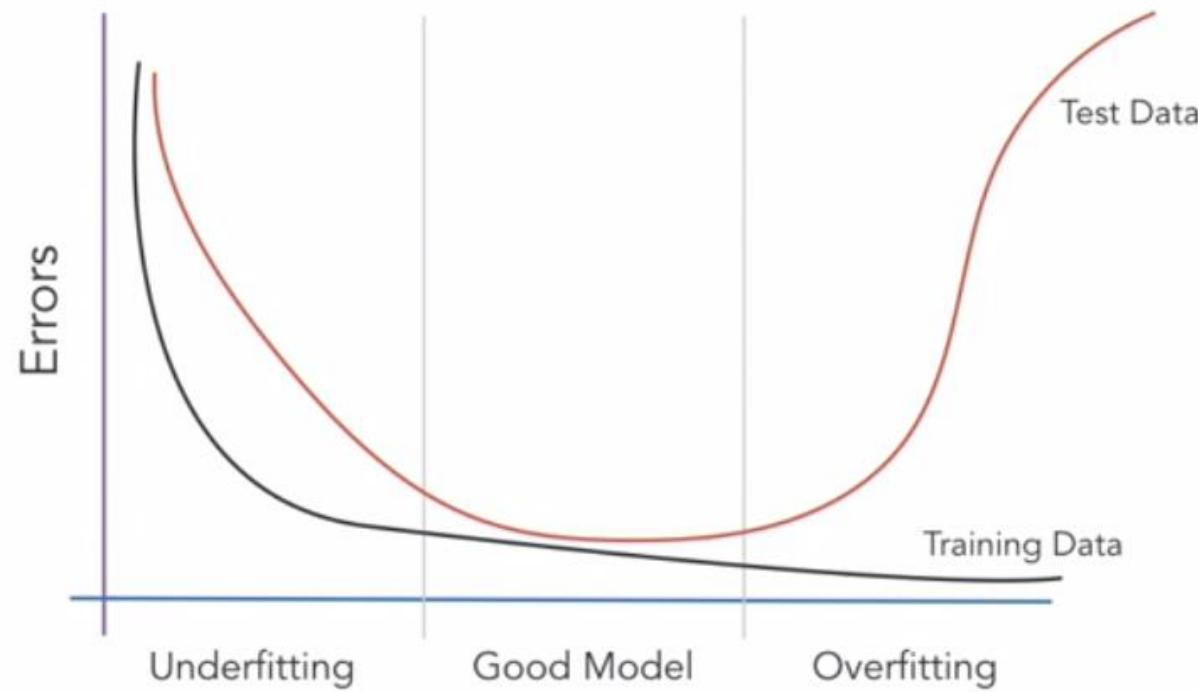
$$\text{Total Error} = (\text{Bias} + \text{Variance}) + \text{Irreducible Error}$$







	Underfitting	Good Model	Overfitting
Complexity	Low	Medium	High
Bias	High	Low	Low
Variance	Low	Low	High
Train Error	High	Low	Low
Test Error	High	Low	High



Two Primary ways to Tune the Model for Optimal Complexity

1. Hyperparameter tuning – choosing a set of optimal hyperparameters for fitting an algorithm
2. Regularization – technique used to reduce overfitting by discouraging overly complex models in some way

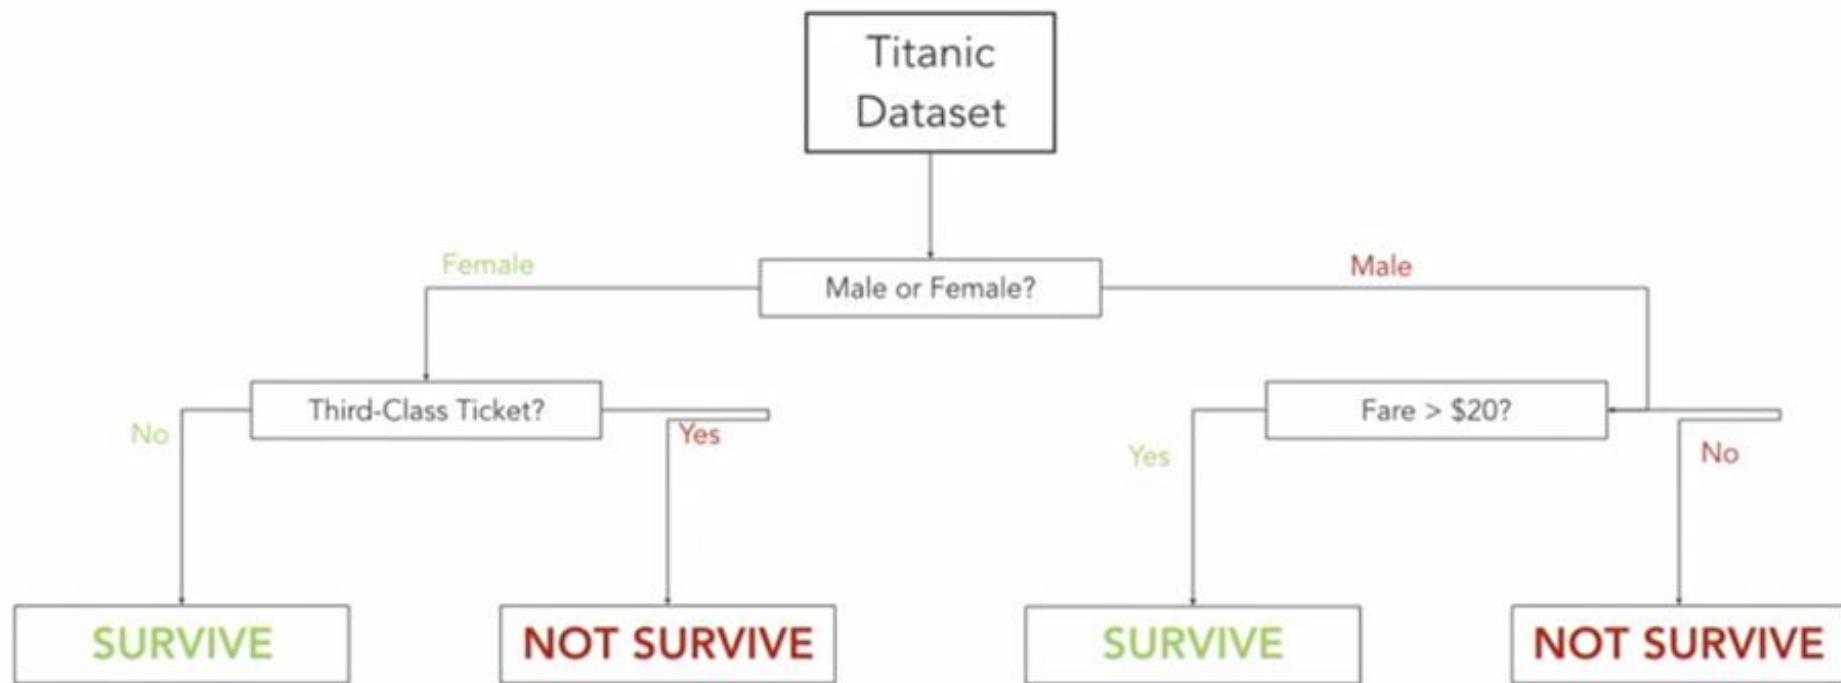


A model **parameter** is a configuration variable that is internal to the model and whose value can be estimated from data.

A model **hyperparameter** is a configuration that is external to the model, whose value cannot be estimated from data, and whose value guides how the algorithm learns parameter values from the data.

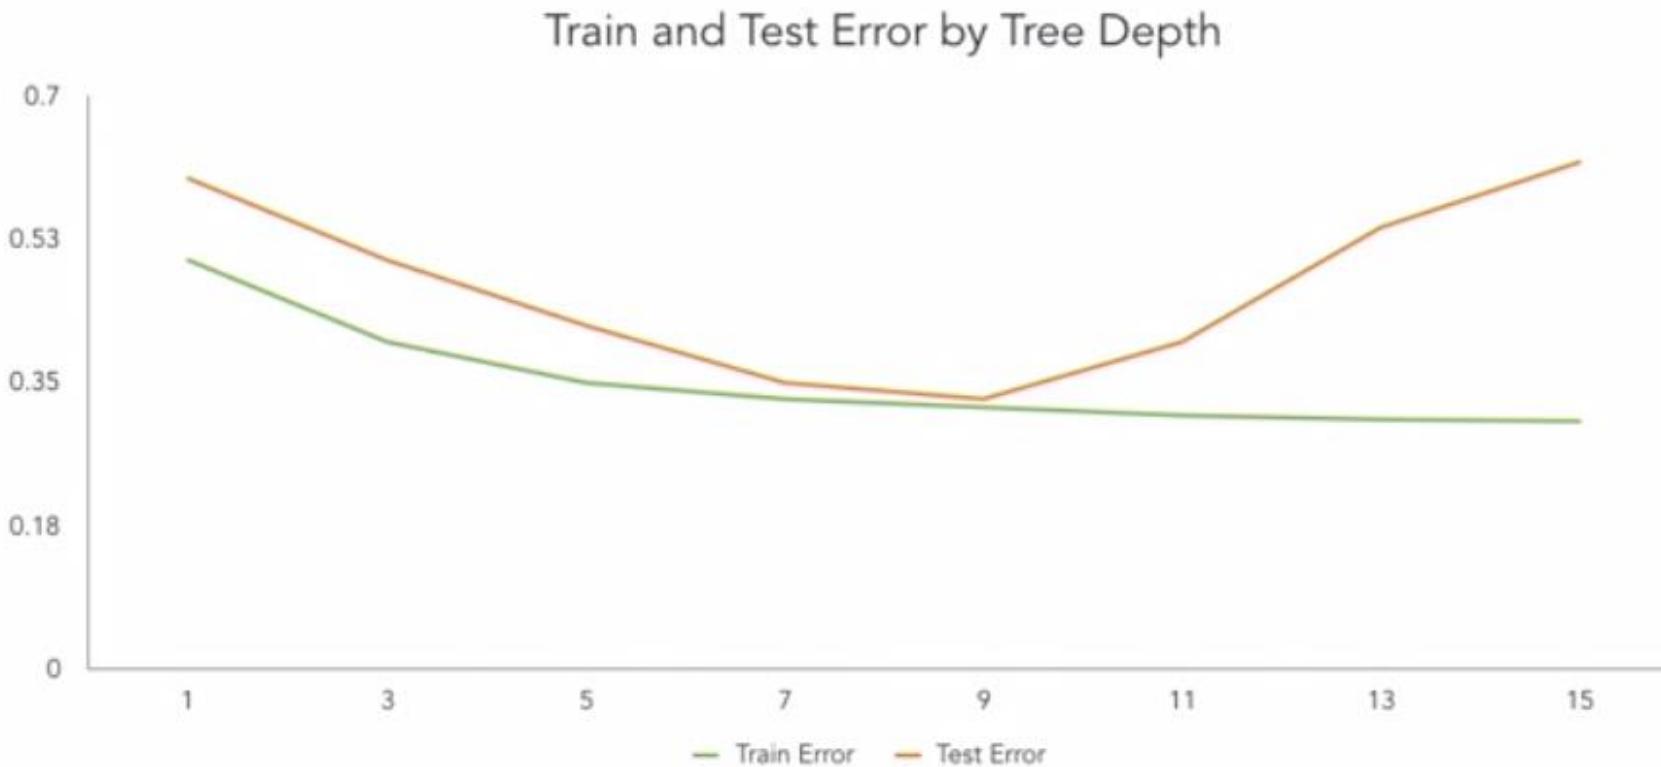
Parameters – fare = \$20, ticket class = third

Hyperparameters – max depth of tree, features to consider, etc.



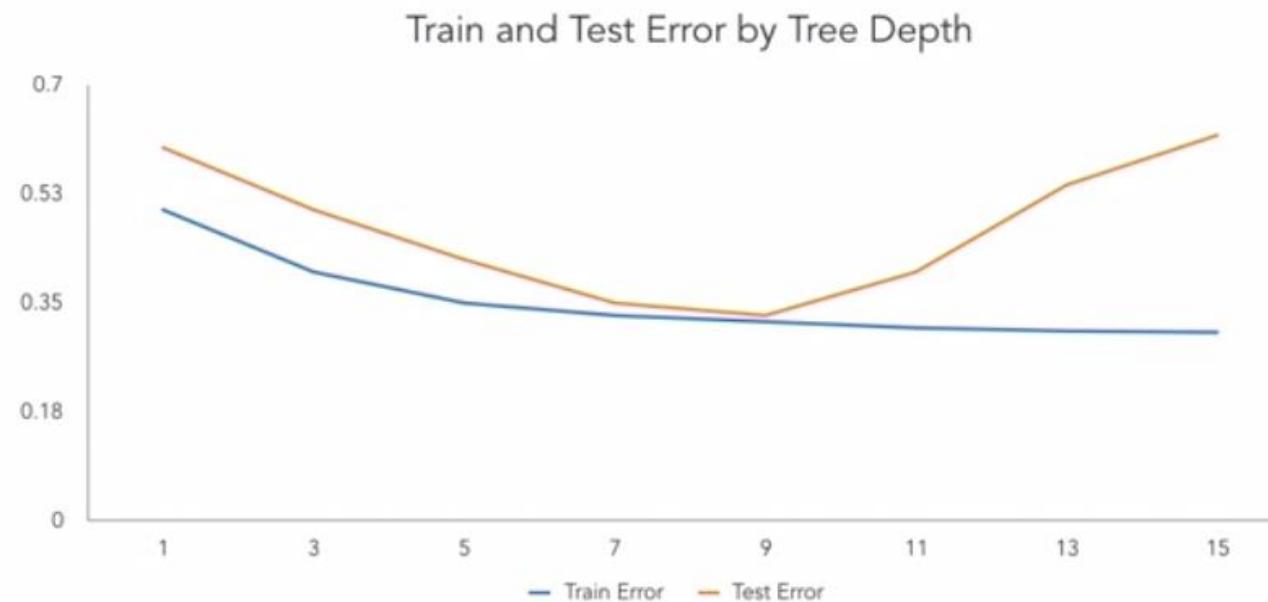
Parameters – fare = \$20, ticket class = third

Hyperparameters – max depth of tree, features to consider, etc.



Two Primary ways to Tune the Model for Optimal Complexity

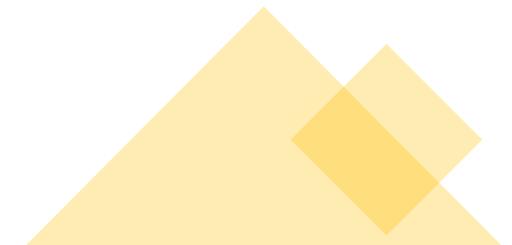
1. Hyperparameter tuning – choosing a set of optimal hyperparameters for fitting an algorithm



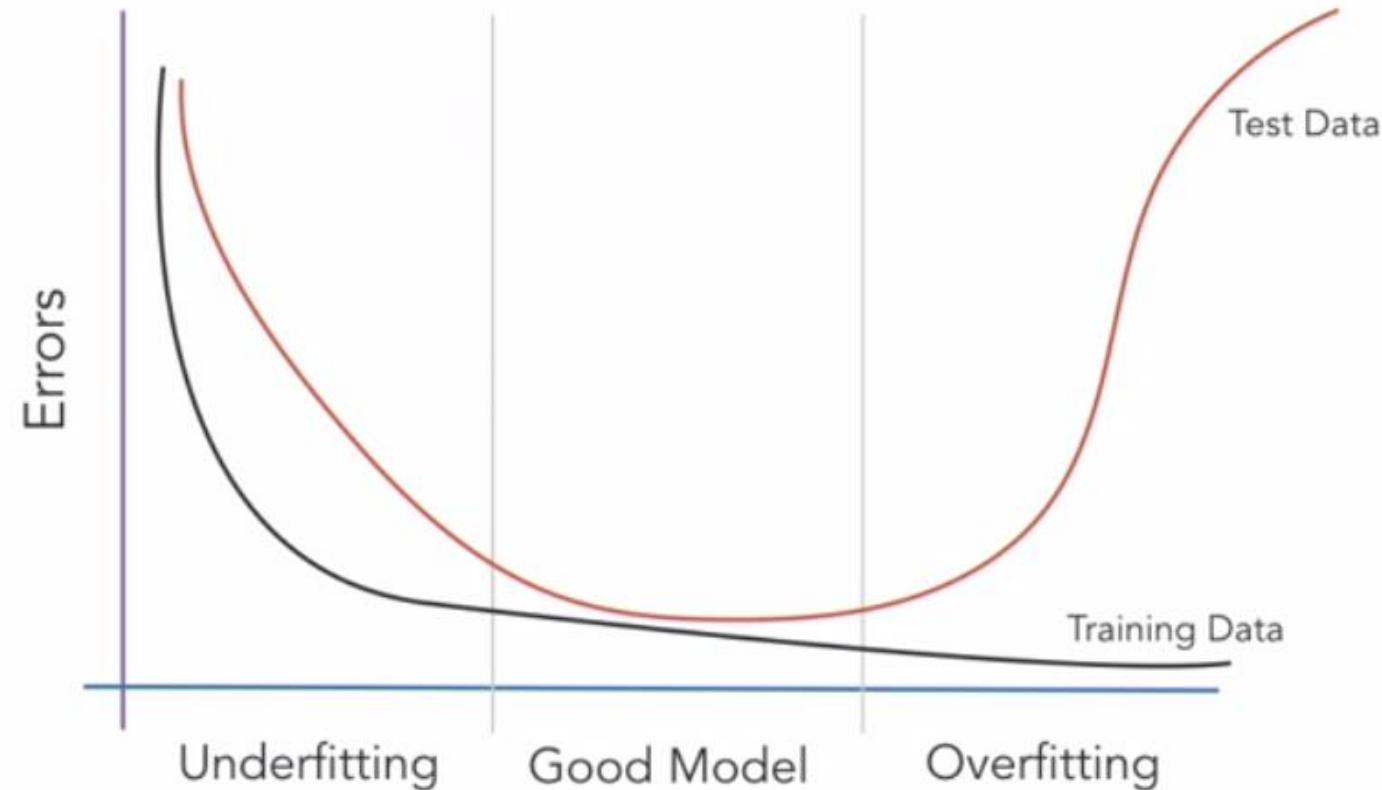
Two Primary ways to Tune the Model for Optimal Complexity

- 
1. Hyperparameter tuning – choosing a set of optimal hyperparameters for fitting an algorithm
 2. Regularization – technique used to reduce overfitting by discouraging overly complex models in some way

Goal – allow enough flexibility for the algorithm to learn the underlying patterns in the data but provide guardrails so it doesn't overfit



Occam's Razor – Whenever possible, choose the simplest answer to a problem



Examples of Regularization

1. Ridge regression and lasso regression – adding a penalty to the loss function to constrain coefficients
2. Dropout – some nodes are ignored during training which forces the other nodes to take on more or less responsibility for the input/output