

# Vision Systems

---

## Lecture 5

# PART 1

---

## Image Sampling and Interpolation

# Cost Improvement using Convolution Theorem?

## Convolution Theorem

Fourier transform of convolution of two functions is a product of their Fourier transforms:

$$F[g * h] = F[g]F[h]$$

**Convolution** in the spatial domain can be obtained through **multiplication** in the frequency domain!

$$g * h = F^{-1}[F[g]F[h]]$$

# Cost Improvement using Convolution Theorem?

## Convolution Theorem

Fourier transform of convolution of two functions is a product of their Fourier transforms:

$$F[g * h] = F[g]F[h]$$


**Convolution** in the spatial domain can be obtained through **multiplication** in the frequency domain!

$$g * h = F^{-1}[F[g]F[h]]$$


- Image convolution needs  $O(N^2 \cdot k^2)$  time, where  $N \times N$  is image size, and  $k \times k$  is kernel size
- By performing convolution in Fourier domain, cost is:  $O(N^2)$  for a single pass over the image + cost of FFT:  $O(N^2 \log N^2)$  for the image and  $O(k^2 \log k^2)$  for the kernel  $\approx O(N^2 \log N^2 + k^2 \log k^2)$ , in total (other terms additive)

## Exercise: Match spatial domain image to Fourier magnitude image


1




2




3



4



5




A




B




C












D



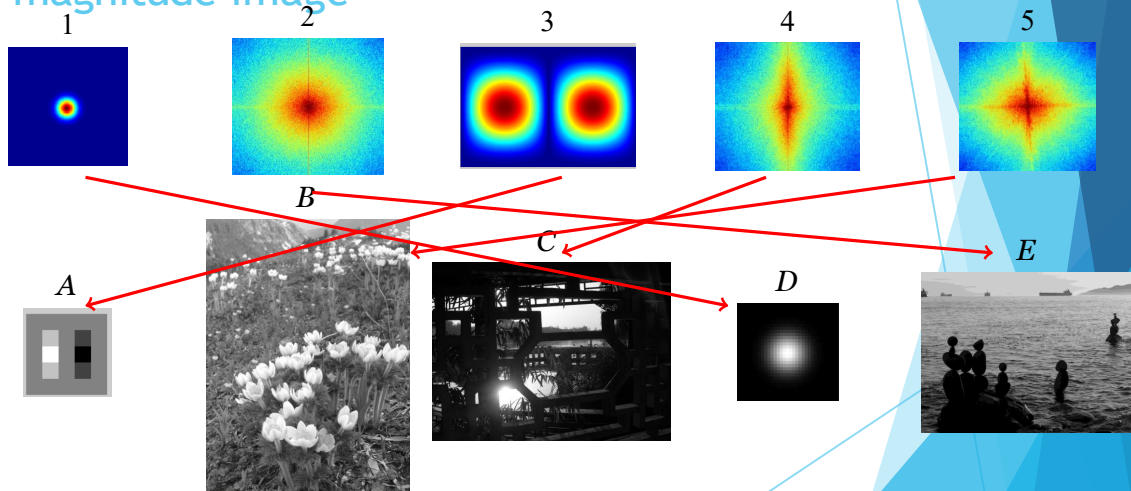
E



## Exercise: Match spatial domain image to Fourier magnitude image



# Exercise: Match spatial domain image to Fourier magnitude image



# What sense does a low-resolution image make to us?



Original



Subsampled & zoomed

Clues from human perception:

- Early processing in human's filters for various orientations and scales of frequency.
- Perceptual cues in mid-high frequencies dominate perception.
- When we see an image from far away, we are effectively **sub-sampling it**.

Credit: Ron Hansen (Unsplash)



## What sense does a low-resolution image make to us?



Original



Subsampled & zoomed


### Clues from human perception:

- Early processing in human's filters for various orientations and scales of frequency.
- Perceptual cues in mid-high frequencies dominate perception.
- When we see an image from far away, we are effectively **sub-sampling** it.

*Credit: Ron Hansen (Unsplash)*

## Sub-sampling

Throw away every other row and column to create a  $1/2$  size image.



$1/4$



$1/8$

# Sub-sampling

Why does this look so crufty?



Credit: S Seitz, R Urtasun


# Sub-sampling

What's happening?



*Credit: S Seitz, R Urtasun*

# Aliasing




- Occurs when your sampling rate is not high enough to capture the amount of detail in your image.
- To do sampling right, need to understand the structure of your signal/image.
- The minimum sampling rate is called the **Nyquist rate**.

Aliased




# Aliasing




- Occurs when your sampling rate is not high enough to capture the amount of detail in your image.
- To do sampling right, need to understand the structure of your signal/image.
- The minimum sampling rate is called the **Nyquist rate**.

Aliased



# Aliasing



- Occurs when your sampling rate is not high enough to capture the amount of detail in your image.
- To do sampling right, need to understand the structure of your signal/image.
- The minimum sampling rate is called the **Nyquist rate**.


Aliased



# Aliasing: Problems

**Shannon's Sampling Theorem** shows that the minimum sampling is:

$$f_s \geq 2f_{max}$$



Example of a 1D signal

## Examples

### Image

- Striped shirt's pattern look weird on screen.

### Video

- **Wagon Wheel effect:** Wheels spins in the opposite direction at high speed.

### Graphics


- Checkerboards disintegrate in ray tracing.



# Aliasing: Problems

**Shannon's Sampling Theorem** shows that the minimum sampling is:

$$f_s \geq 2f_{\max}$$



Example of a 1D signal

## Examples

### Image

- Striped shirt's pattern look weird on screen.

### Video

- **Wagon Wheel effect:** Wheels spins in the opposite direction at high speed.


### Graphics

- Checkerboards disintegrate in ray tracing.

# Aliasing: Problems

Shannon's Sampling Theorem shows that the minimum sampling is:

$$f_s \geq 2f_{max}$$



Example of a 1D signal

## Examples

### Image

- Striped shirt's pattern look weird on screen.

### Video

- **Wagon Wheel effect:** Wheels spins in the opposite direction at high speed.


### Graphics

- Checkerboards disintegrate in ray tracing.

# Aliasing: Problems

Shannon's Sampling Theorem shows that the minimum sampling is:

$$f_s \geq 2f_{max}$$



Example of a 1D signal

## Examples

- **Image**

- Striped shirt's pattern look weird on screen.

- **Video**

- **Wagon Wheel effect:** Wheels spins in the opposite direction at high speed.

- **Graphics**


- Checkerboards disintegrate in ray tracing.

## Aliasing: Image




Striped shirt's pattern look weird on screen.

## Aliasing : Video



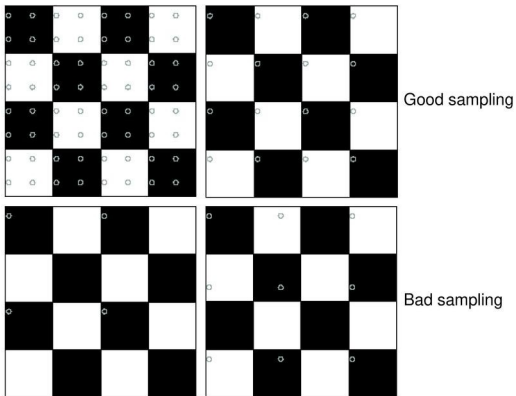
**Wagon Wheel effect:** Wheels spins in the opposite direction at high speed.

## Aliasing: Graphics



Checkerboards disintegrate in ray tracing.

## Aliasing: Nyquist Limit 2D example



Credit: S Seitz, R Urtasun


# Anti-aliasing





# Anti-aliasing

Example: Gaussian Pre-filtering



Credit: N Snavely, R Urtasun




$\frac{1}{2}$

$\frac{1}{4}$  (2x Zoom)


$\frac{1}{8}$  (4x Zoom)

# Subsampling with Gaussian Pre-filtering




Credit: N Snavely, R Urtasun

# Upsampling



How to go from left to right? **Interpolation**. Simple method: Repeat each row and column 10 times (Nearest Neighbour Interpolation).


# Interpolation



Recall how a digital image is formed,

$$F[x, y] = \text{quantize}\{f(xd, yd)\}$$

- It is a discrete point-sampling of a continuous function.
- If we could somehow reconstruct the original function, any new image could be generated, at any resolution and scale.




What if we don't know  $f$ ?

- Guess an approximation: Can be done in a principled way via filtering.
- Convert  $F$  to a continuous function:

$$f_F(x) = \begin{cases} F(\frac{x}{d}) & \text{if } \frac{x}{d} \text{ is an integer} \\ 0 & \text{otherwise} \end{cases}$$

- Reconstruct:  $\hat{f} = h * f_F$


# Interpolation



Recall how a digital image is formed,

$$F[x, y] = \text{quantize}\{f(xd, yd)\}$$

- It is a discrete point-sampling of a continuous function.
- If we could somehow reconstruct the original function, any new image could be generated, at any resolution and scale.




- What if we don't know  $f$ ?
  - Guess an approximation: Can be done in a principled way via filtering.
  - Convert  $F$  to a continuous function:

$$f_F(x) = \begin{cases} F\left(\frac{x}{d}\right) & \text{if } \frac{x}{d} \text{ is an integer} \\ 0 & \text{otherwise} \end{cases}$$

- Reconstruct:  $\hat{f} = h * f_F$


# Interpolation



Recall how a digital image is formed,

$$F[x, y] = \text{quantize}\{f(xd, yd)\}$$

- It is a discrete point-sampling of a continuous function.
- If we could somehow reconstruct the original function, any new image could be generated, at any resolution and scale.




- What if we don't know  $f$ ?
- Guess an approximation: Can be done in a principled way via filtering.
- Convert  $F$  to a continuous function:

$$f_F(x) = \begin{cases} F(\frac{x}{d}) & \text{if } \frac{x}{d} \text{ is an integer} \\ 0 & \text{otherwise} \end{cases}$$

- Reconstruct:  $\hat{f} = h * f_F$


# Interpolation



Recall how a digital image is formed,

$$F[x, y] = \text{quantize}\{f(xd, yd)\}$$

- It is a discrete point-sampling of a continuous function.
- If we could somehow reconstruct the original function, any new image could be generated, at any resolution and scale.




- What if we don't know  $f$ ?
  - Guess an approximation: Can be done in a principled way via filtering.
  - Convert  $F$  to a continuous function:

$$f_F(x) = \begin{cases} F(\frac{x}{d}) & \text{if } \frac{x}{d} \text{ is an integer} \\ 0 & \text{otherwise} \end{cases}$$

- Reconstruct:  $\hat{f} = h * f_F$


# Interpolation



Recall how a digital image is formed,

$$F[x, y] = \text{quantize}\{f(xd, yd)\}$$

- It is a discrete point-sampling of a continuous function.
- If we could somehow reconstruct the original function, any new image could be generated, at any resolution and scale.




- What if we don't know  $f$ ?
  - Guess an approximation: Can be done in a principled way via filtering.
  - Convert  $F$  to a continuous function:

$$f_F(x) = \begin{cases} F(\frac{x}{d}) & \text{if } \frac{x}{d} \text{ is an integer} \\ 0 & \text{otherwise} \end{cases}$$

- Reconstruct:  $\hat{f} = h * f_F$




# Interpolation



Recall how a digital image is formed,

$$F[x, y] = \text{quantize}\{f(xd, yd)\}$$

- It is a discrete point-sampling of a continuous function.
- If we could somehow reconstruct the original function, any new image could be generated, at any resolution and scale.




- What if we don't know  $f$ ?
  - Guess an approximation: Can be done in a principled way via filtering.
  - Convert  $F$  to a continuous function:

$$f_F(x) = \begin{cases} F\left(\frac{x}{d}\right) & \text{if } \frac{x}{d} \text{ is an integer} \\ 0 & \text{otherwise} \end{cases}$$

- Reconstruct:  $\hat{f} = h * f_F$


# Interpolation



Recall how a digital image is formed,

$$F[x, y] = \text{quantize}\{f(xd, yd)\}$$

- It is a discrete point-sampling of a continuous function.
- If we could somehow reconstruct the original function, any new image could be generated, at any resolution and scale.



- What if we don't know  $f$ ?
    - Guess an approximation: Can be done in a principled way via filtering.
    - Convert  $F$  to a continuous function:
- $$f_F(x) = \begin{cases} F\left(\frac{x}{d}\right) & \text{if } \frac{x}{d} \text{ is an integer} \\ 0 & \text{otherwise} \end{cases}$$
- Reconstruct:  $\hat{f} = h * f_F$


# Interpolation as Convolution

- To **interpolate** (or upsample) an image to a higher resolution, we need an **interpolation kernel** with which to **convolve** the image:

$$g(i,j) = \sum_{k,L} f(k,l)h(i - rk, j - rl)$$

Above formula similar to discrete convolution<sup>a</sup>, except that we replace  $k$  and  $l$  in  $h(\cdot)$  with  $rk$  and  $rl$ , where  $r$  is the upsampling rate.

- Linear interpolator (corresponding to *tent kernel*) produces interpolating piecewise linear curves.
- More complex kernels e.g., B-splines.



$$^a g = f * h \Rightarrow g(i,j) = \sum_{k,l} f(k,l)h(i - k, j - l)$$

# Interpolation as Convolution

- To **interpolate** (or upsample) an image to a higher resolution, we need an **interpolation kernel** with which to **convolve** the image:

$$g(i,j) = \sum_{k,l} f(k,l)h(i - rk, j - rl)$$

Above formula similar to discrete convolution<sup>a</sup>, except that we replace  $k$  and  $l$  in  $h(\cdot)$  with  $rk$  and  $rl$ , where  $r$  is the upsampling rate.

- Linear interpolator (corresponding to *tent kernel*) produces interpolating piecewise linear curves.
- More complex kernels e.g., B-splines.

---


$$^a g = f * h \Rightarrow g(i,j) = \sum_{k,l} f(k,l)h(i - k, j - l)$$

## Examples

Original Image:



Upsampled Images:



Left to right: Nearest Neighbour Interpolation, Bilinear Interpolation, Bicubic Interpolation.

# Interpolation and Decimation

## Interpolation

To **interpolate** (or upsample) an image to a higher resolution, we need an **interpolation kernel** with which to convolve the image ( $r$  is upsampling rate):

$$g(i,j) = \sum_{k,l} f(k,l)h(i - rk, j - rl)$$

## Decimation (Sub-sampling)

To **decimate** (or sub-sample) an image to a lower resolution, we need an **decimation kernel** with which to convolve the image ( $r$  is downsampling rate):

$$g(i,j) = \sum_{k,l} f(k,l)h(i - \frac{k}{r}, j - \frac{l}{r})$$

# Homework

## Readings

Chapter 3 (§3.5.1-3.5.2), Szeliski, *Computer Vision: Algorithms and Applications*

Chapter 7 (§7.4), Forsyth and Ponce, *Computer Vision: A Modern Approach*