

# Vision Systems

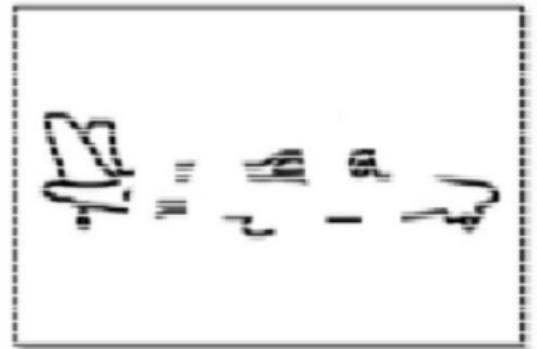
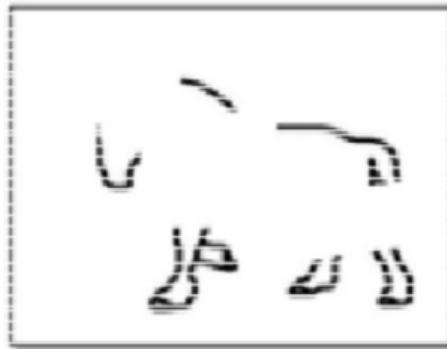
---

Lecture 6

# Part 1

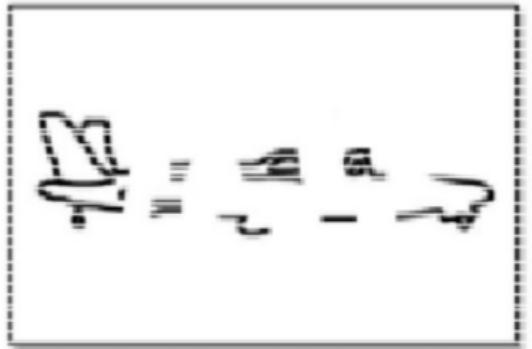
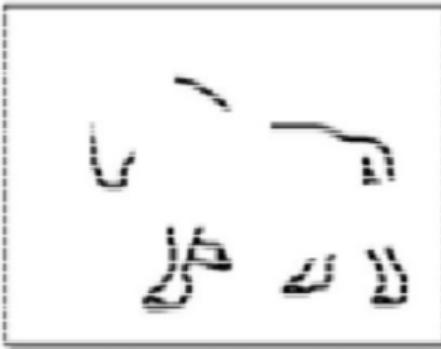
## Edge Detection

# Edge Detection



- Map image from 2D matrix of pixels to a set of curves or line segments or contours  $\Rightarrow$  More compact representation than pixels
- Key idea?**

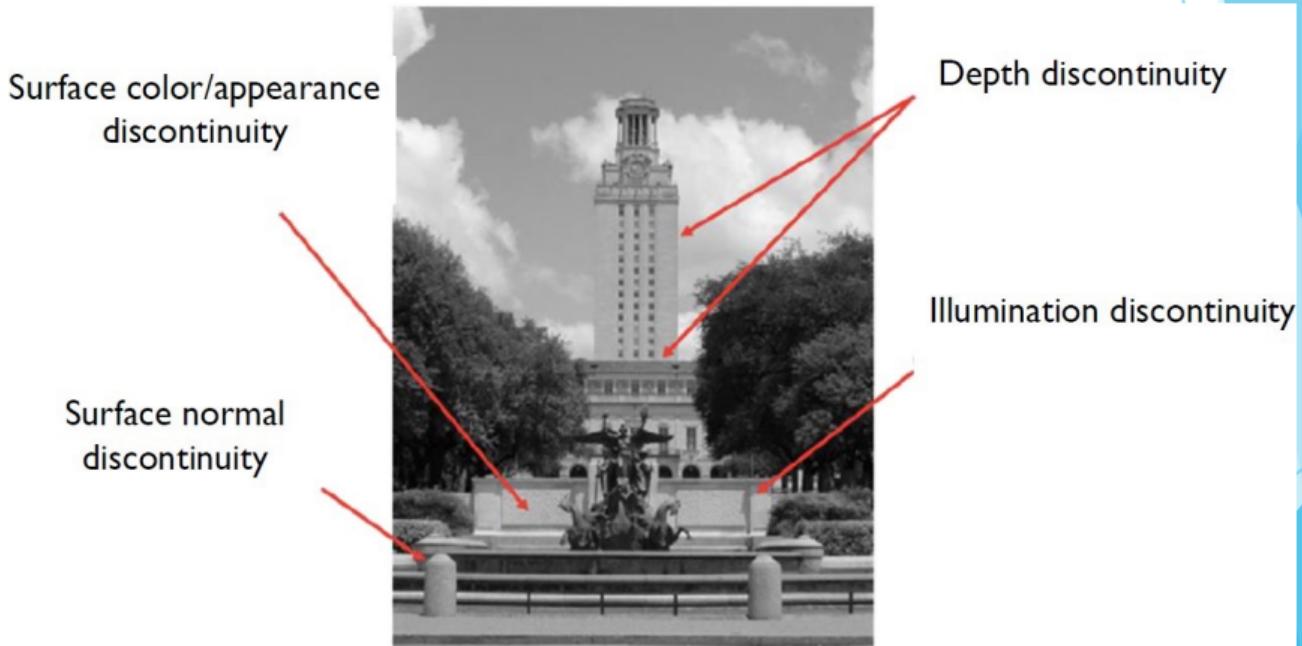
# Edge Detection



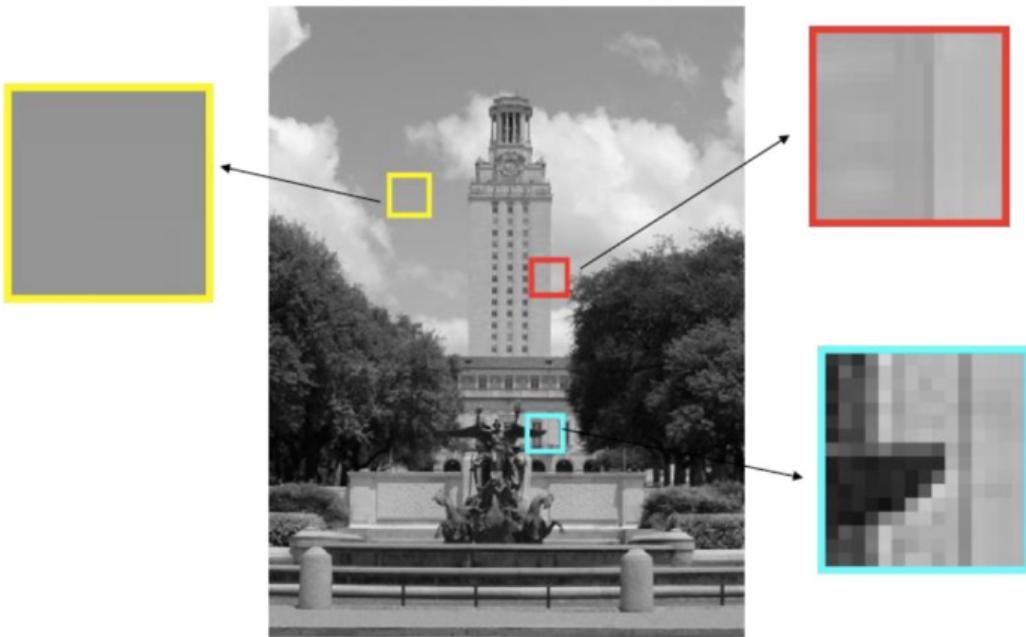
- Map image from 2D matrix of pixels to a set of curves or line segments or contours  $\Rightarrow$  More compact representation than pixels
- Key idea?** Look for strong gradients, and then post-process

# How are Edges Caused?

- Variety of factors:



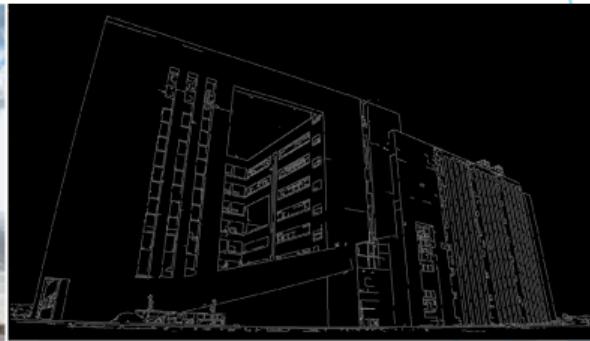
## Looking More Locally



Source: K Grauman, R Urtasun

# Why are Edges Important?

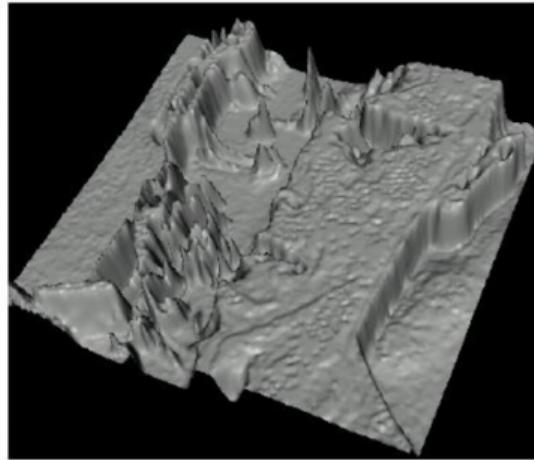
- Group pixels into objects or parts
- Allow us to track important features (e.g., corners, curves, lines).
- Cues for 3D shape
- Guiding interactive image editing



Source: Derek Hoiem

# Edges in Images as Functions

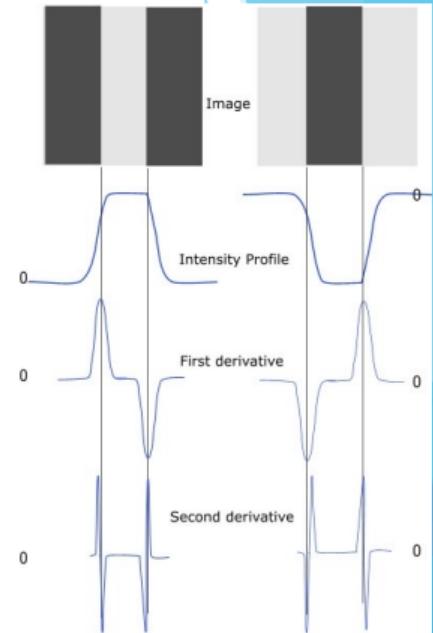
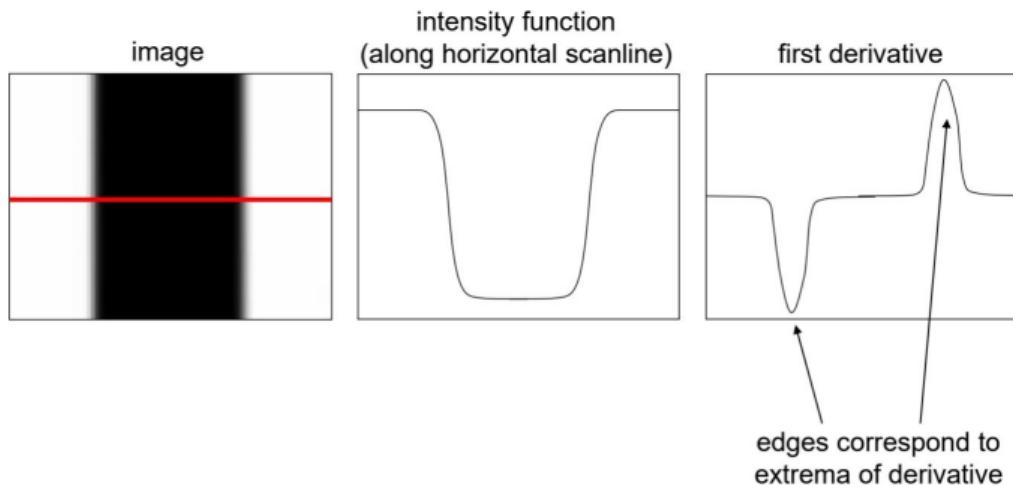
- Edges look like steep cliffs



Source: N Snavely, R Urtasun

# Derivatives and Edges

- An edge is a place of rapid change in the image intensity function



# Derivatives with Convolution

- For 2D function,  $f(x, y)$ , the partial derivative is:

$$\frac{\partial f(x, y)}{\partial x} = \lim_{\varepsilon \rightarrow 0} \frac{f(x + \varepsilon, y) - f(x, y)}{\varepsilon}$$

- For discrete data, we can approximate using finite differences:

$$\frac{\partial f(x, y)}{\partial x} \approx \frac{f(x + 1, y) - f(x, y)}{1}$$

# Derivatives with Convolution

- For 2D function,  $f(x, y)$ , the partial derivative is:

$$\frac{\partial f(x, y)}{\partial x} = \lim_{\varepsilon \rightarrow 0} \frac{f(x + \varepsilon, y) - f(x, y)}{\varepsilon}$$

- For discrete data, we can approximate using finite differences:

$$\frac{\partial f(x, y)}{\partial x} \approx \frac{f(x + 1, y) - f(x, y)}{1}$$

- To implement above as convolution, what would be the associated filter?

# Derivatives with Convolution

- For 2D function,  $f(x, y)$ , the partial derivative is:

$$\frac{\partial f(x, y)}{\partial x} = \lim_{\varepsilon \rightarrow 0} \frac{f(x + \varepsilon, y) - f(x, y)}{\varepsilon}$$

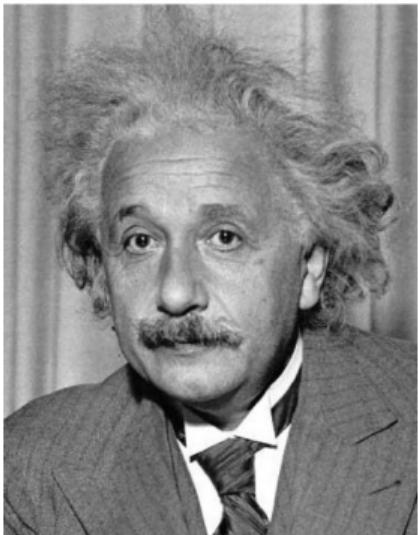
- For discrete data, we can approximate using finite differences:

$$\frac{\partial f(x, y)}{\partial x} \approx \frac{f(x + 1, y) - f(x, y)}{1}$$

- To implement above as convolution, what would be the associated filter?

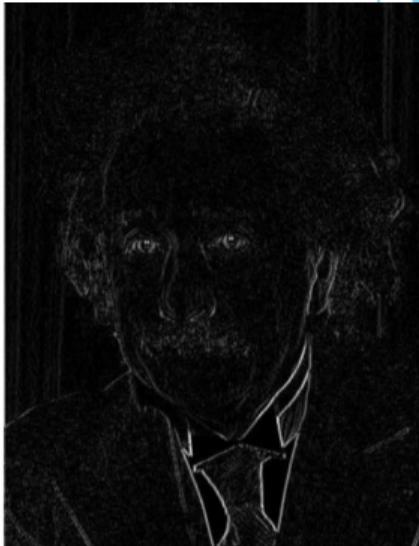


# Sobel Edge Detection Filters



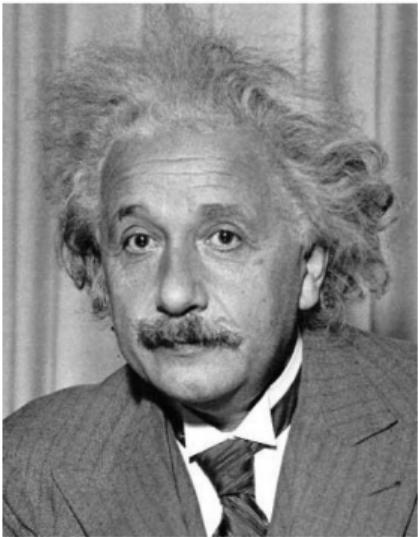
1	0	-1
2	0	-2
1	0	-1

Sobel



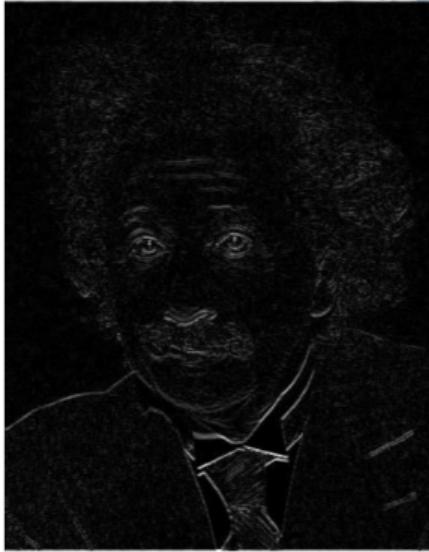
Vertical Edge  
(absolute value)

# Sobel Edge Detection Filters



1	2	1
0	0	0
-1	-2	-1

Sobel



Horizontal Edge  
(absolute value)

# Finite Difference Filters

Prewitt

$M_x$
-1 0 1
-1 0 1
-1 0 1

Sobel

$M_x$
-1 0 1
-2 0 2
-1 0 1

Roberts

$M_x$
0 1
-1 0

$M_y$

$M_y$
1 1 1
0 0 0
-1 -1 -1

$M_y$
1 2 1
0 0 0
-1 -2 -1

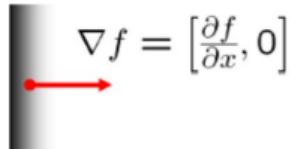
$M_y$
1 0
0 -1

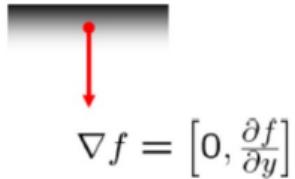
## Image Gradients

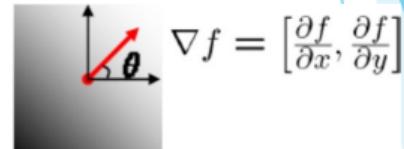
- The gradient of an image  $\nabla f = \frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}$

# Image Gradients

- The gradient of an image  $\nabla f = \begin{bmatrix} h \\ \frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \end{bmatrix}$
- The gradient points in the direction of most rapid change in intensity

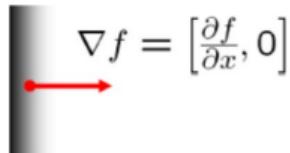

$$\nabla f = \left[ \frac{\partial f}{\partial x}, 0 \right]$$


$$\nabla f = \left[ 0, \frac{\partial f}{\partial y} \right]$$

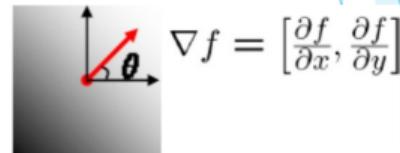

$$\nabla f = \left[ \frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

# Image Gradients

- The gradient of an image  $\nabla f = \begin{bmatrix} h \\ \frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \end{bmatrix}$
- The gradient points in the direction of most rapid change in intensity

$$\nabla f = \left[ \frac{\partial f}{\partial x}, 0 \right]$$


$$\nabla f = \left[ 0, \frac{\partial f}{\partial y} \right]$$


$$\nabla f = \left[ \frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$


- The **gradient direction** (orientation of edge normal) is given by:

$$\theta = \tan^{-1} \frac{\frac{\partial f}{\partial y}}{\frac{\partial f}{\partial x}}$$

# Image Gradients

- The gradient of an image  $\nabla f = \left[ \frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$
- The gradient points in the direction of most rapid change in intensity

$$\nabla f = \left[ \frac{\partial f}{\partial x}, 0 \right]$$
$$\nabla f = \left[ 0, \frac{\partial f}{\partial y} \right]$$
$$\nabla f = \left[ \frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

- The **gradient direction** (orientation of edge normal) is given by:

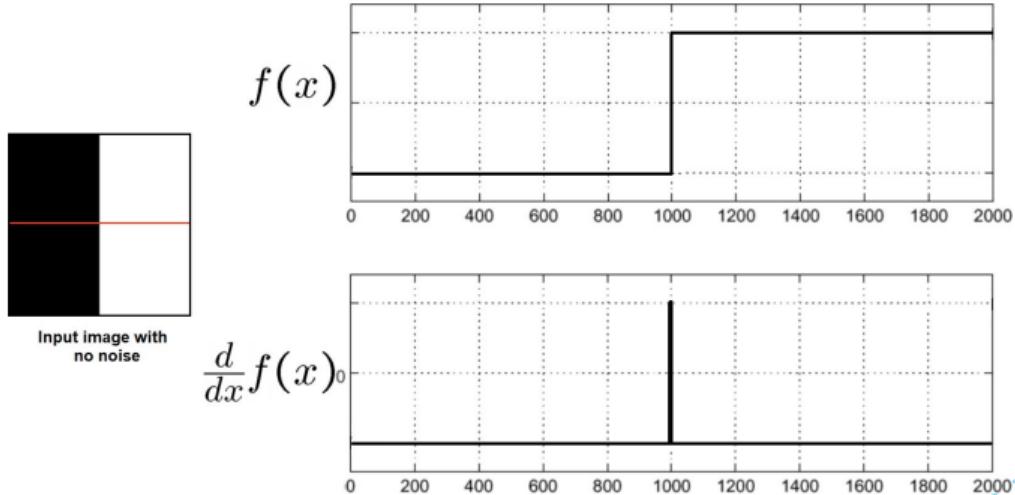
$$\theta = \tan^{-1} \frac{\frac{\partial f}{\partial y}}{\frac{\partial f}{\partial x}}$$

- The **edge strength** is given by the magnitude  $\|\nabla f\| =$

$$\sqrt{\left( \frac{\partial f}{\partial x} \right)^2 + \left( \frac{\partial f}{\partial y} \right)^2}$$

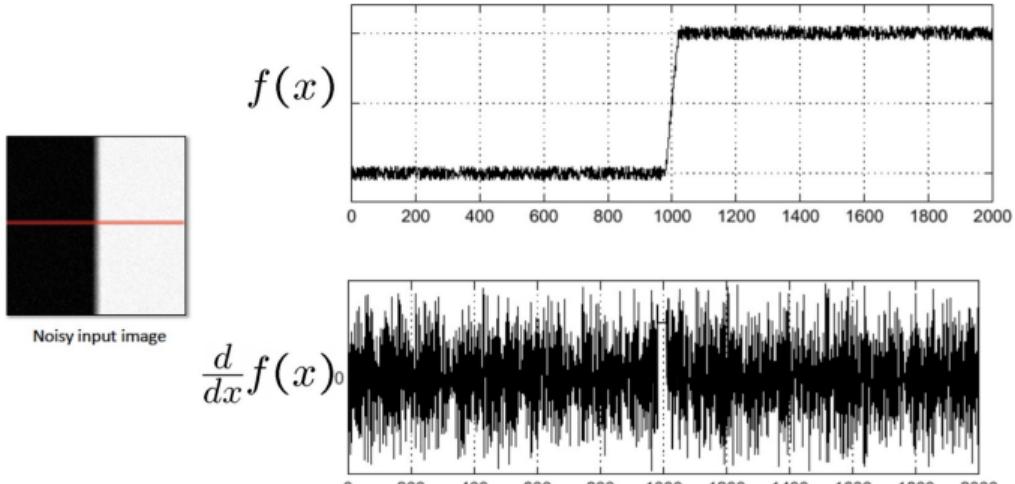
# Derivative with No Noise

- Consider a single row or column of the image
  - Plotting intensity as a function of position gives a signal



Where is the edge?

# Effect of Noise

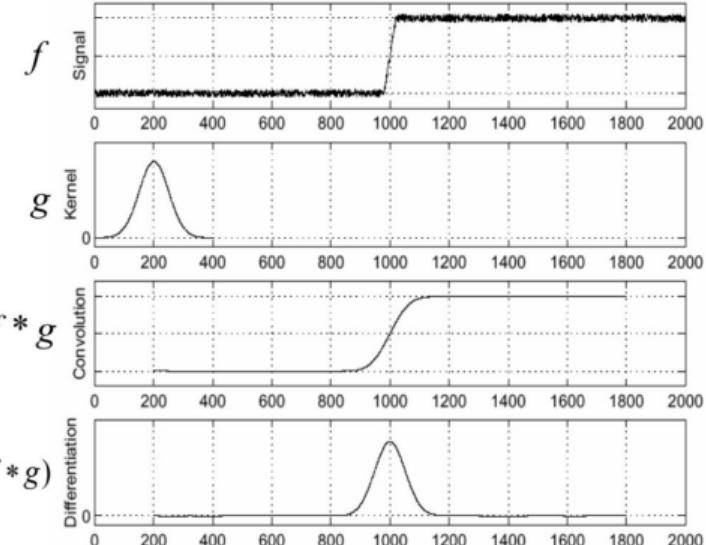


Now, where is the edge?

Source: S Seitz, K Grauman

# Effect of Noise

- Smooth first, and look for peaks in  $\frac{d}{dx}(f * g)$

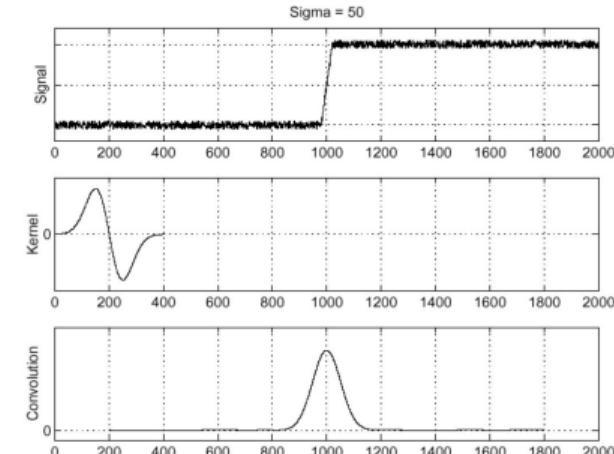


# Derivative theorem of Convolution

- Differentiation is achieved through convolution, and convolution is associative:

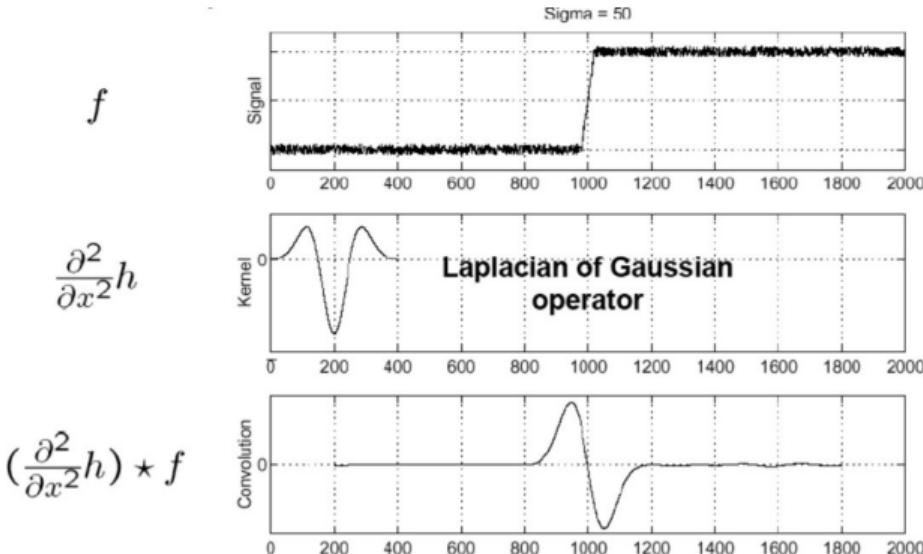
$$\frac{d}{dx}(f * g) = f * \frac{d}{dx}g = \frac{d}{dx}f * g$$

- This saves us an operation:

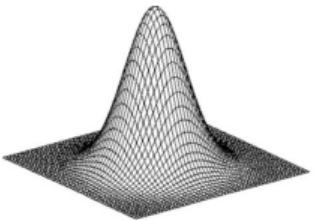


# What about Second Derivative?

- Edge by detecting **zero-crossing** of bottom graph

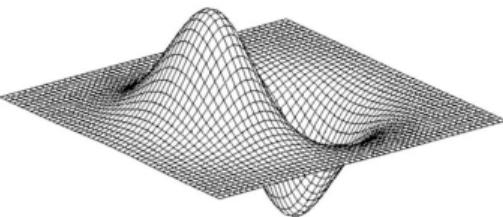


# Derivative and Laplacian of Gaussians



Gaussian

$$h_\sigma(x, y) = \frac{1}{2\pi\sigma^2} \exp^{-\frac{x^2+y^2}{2\sigma^2}}$$



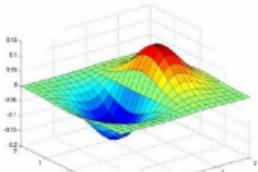
Derivative of Gaussian (x)

$$\frac{\partial}{\partial x} h_\sigma(u, v)$$

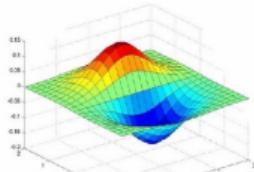


Laplacian of Gaussian

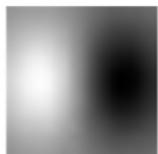
$$\nabla^2 h_\sigma(u, v)$$



x-direction



y-direction



with  $\nabla^2$  the Laplacian operator

$$\nabla^2 f = \frac{\partial^2 f}{\partial x^2} + \frac{\partial^2 f}{\partial y^2}$$

Which one finds  
horizontal/vertical  
edges?

Source: S Seitz, R Urtasun

# Compute of Gradients



X-Derivative of Gaussian



Y-Derivative of Gaussian



Gradient Magnitude

# Properties of an Edge Detector



where is the edge?

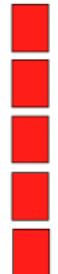


## Properties of an Edge Detector

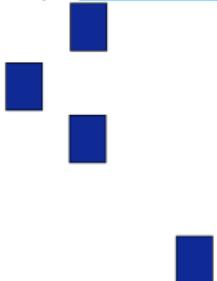
- Criteria for a good edge detector?

# Properties of an Edge Detector

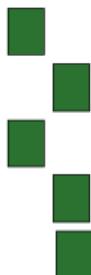
- Criteria for a good edge detector?
  - Good detection:** find all real edges, ignoring noise or other artifacts
  - Good localization:** detect edges as close as possible to true edges
  - Single response:** return one point only for each true edge point



True Edge



Poor robustness to noise



Poor localization

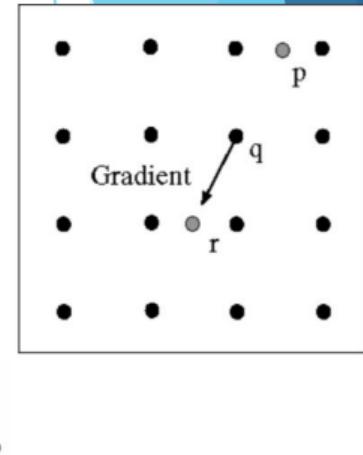
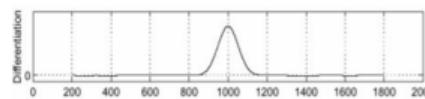
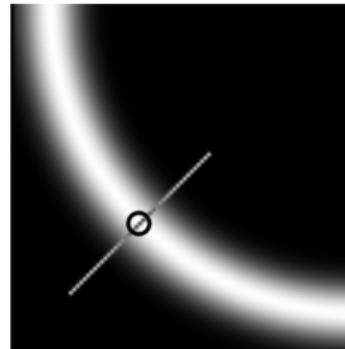


Too many responses

# Non-Maxima Suppression



where is the edge?



- Check if pixel is local maximum along gradient direction:
  - Could require checking interpolated pixels  $p$  and  $r$

## Non-Maxima Suppression



Before and after non-maxima suppression

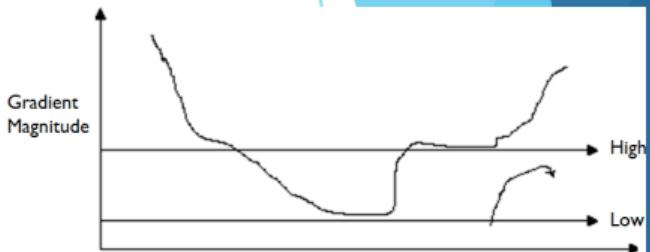
Source: Derek Hoiem

## Hysteresis Thresholding

- Check for well-connected edges. How?

# Hysteresis Thresholding

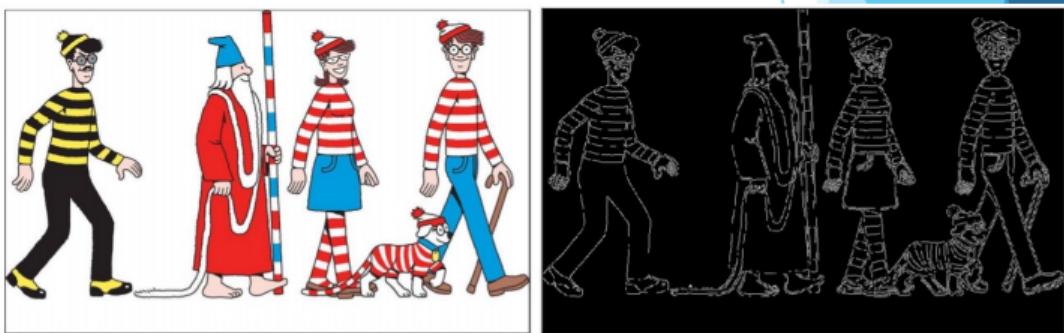
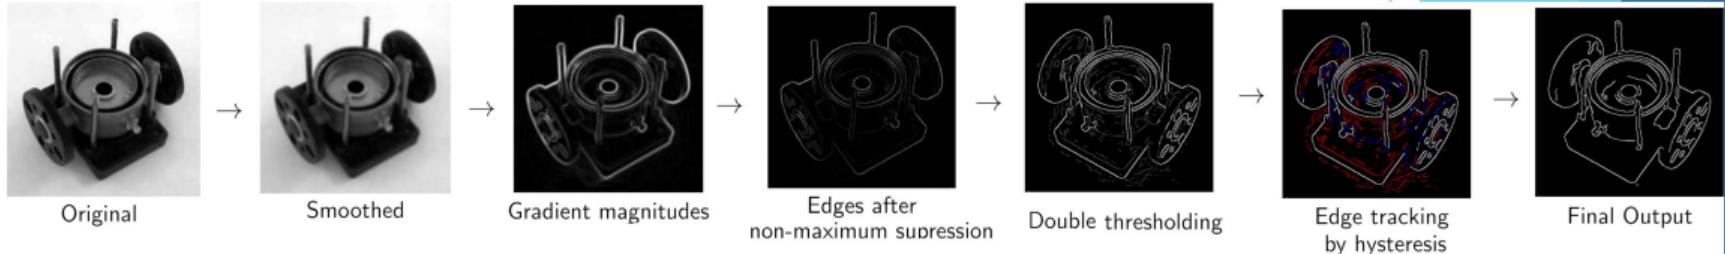
- Check for well-connected edges. How?
  - Use **hysteresis**: use a *high* threshold to start edge curves and a *low* threshold to continue them.
- How does it work?
  - If gradient at pixel  $>$  'High'  $\Rightarrow$  '**edge pixel**'
  - If gradient at pixel  $<$  'Low'  $\Rightarrow$  '**non-edge pixel**'
  - If gradient at pixel  $\geq$  'Low' and  $\leq$  'High'  $\Rightarrow$  '**edge pixel**' iff it is connected to an 'edge pixel' directly or via pixels between 'Low' and 'High'



# Canny Edge Detector

- Probably the most widely used edge detector in computer vision (Canny 1986)
- Algorithm:
  - ① Filter image with derivative of Gaussian
  - ② Find magnitude and orientation of gradient
  - ③ Non-maximum suppression
  - ④ Linking and thresholding (hysteresis):
    - Define two thresholds: low and high
    - Use the high threshold to start edge curves and the low threshold to continue them

# Canny Edge Pipeline and Examples

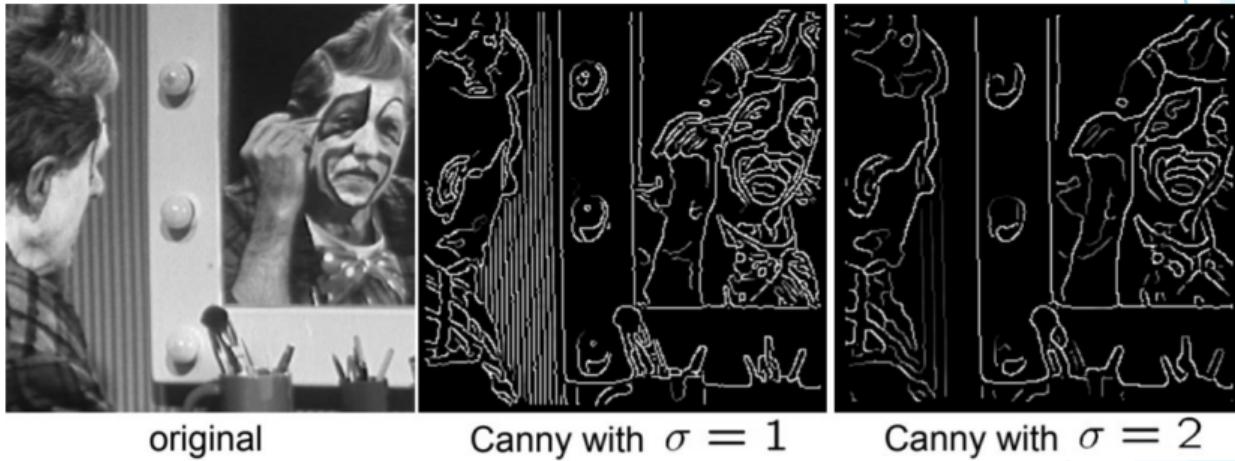


Canny Edges

Source: Prem Kalra, R Urtasun, S Fidler

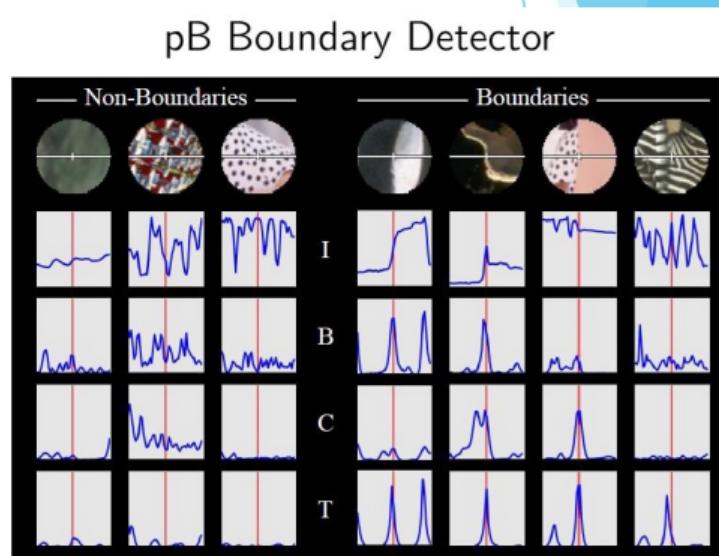
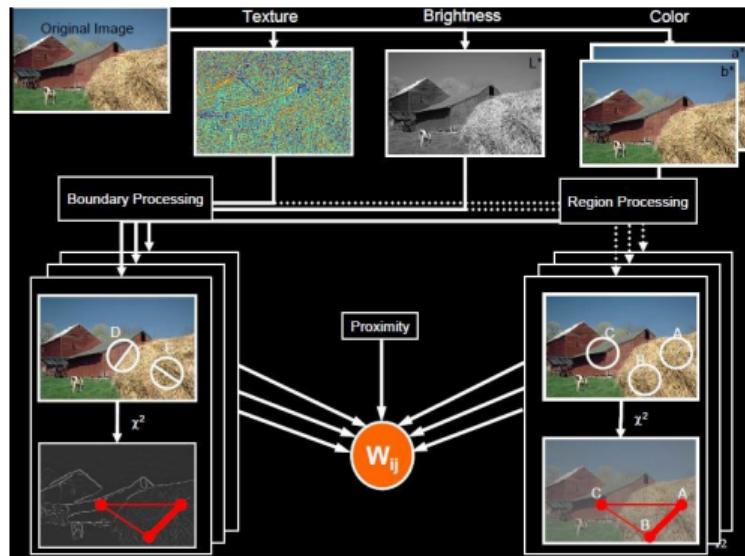
# Effect of $\sigma$ in Canny Edge Detector

- The choice of  $\sigma$  (Gaussian kernel spread/size) depends on desired behavior
  - large  $\sigma$  detects large-scale edges
  - small  $\sigma$  detects fine edges



# More Recent Methods in Edge Detection

Learning to Detect Natural Image Boundaries Using Local Brightness, Color, and Texture Cues  
(Martin et al, 2004)



Source: Derek Hoiem

# More Recent Methods in Edge Detection

Learning to Detect Natural Image Boundaries Using Local Brightness, Color, and Texture Cues  
(Martin et al, 2004)



# More Recent Methods in Edge Detection

Structured Forests for Fast Edge Detection (Dollár et al, 2013)

- Goal: quickly predict whether each pixel is an edge
- Insights
  - Predictions can be learned from training data
  - Predictions for nearby pixels should not be independent
- Solution
  - Train structured random forests to split data into patches with similar boundaries based on features
  - Predict boundaries at patch level, rather than pixel level, and aggregate (average votes)

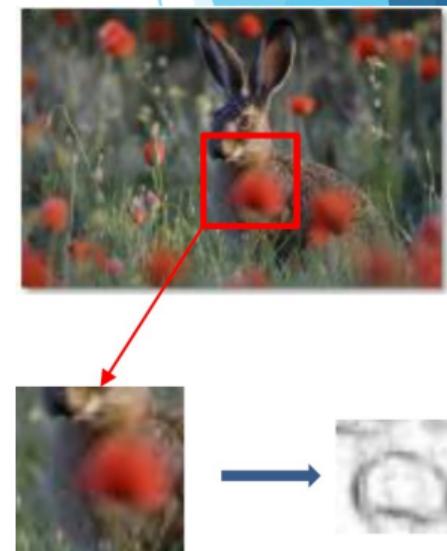


# More Recent Methods in Edge Detection

Structured Forests for Fast Edge Detection (Dollár et al, 2013)

- Algorithm

- ① Extract overlapping  $32 \times 32$  patches at three scales
- ② Features are pixel values and pairwise differences in feature maps (LUV color, gradient magnitude, oriented gradient)
- ③ Predict  $T$  boundary maps in the central  $16 \times 16$  region using  $T$  trained decision trees
- ④ Average predictions for each pixel across all patches

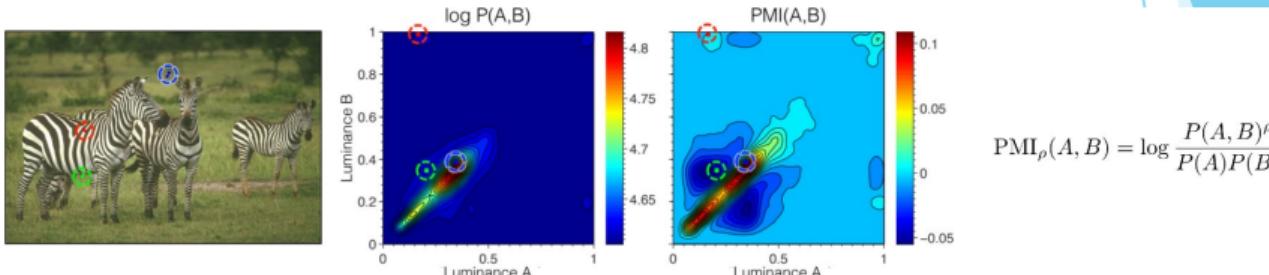


Source: Derek Hoiem

# More Recent Methods in Edge Detection

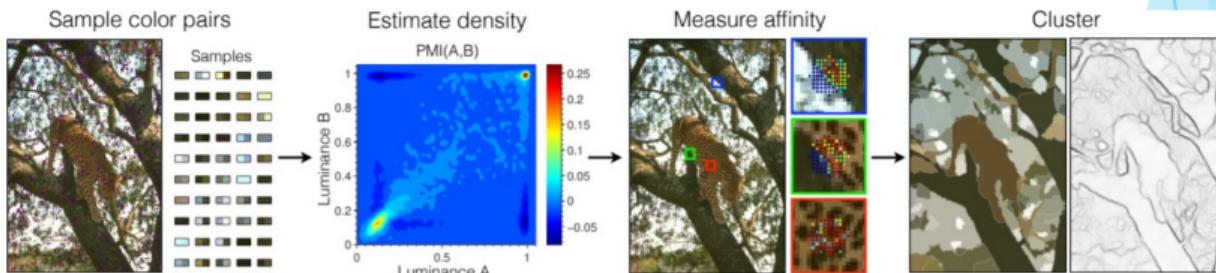
Crisp Boundary Detection using Pointwise Mutual Information (Isola et al, 2014)

- Pixel combinations that are unlikely to be together are edges



$$PMI_{\rho}(A, B) = \log \frac{P(A, B)^{\rho}}{P(A)P(B)}$$

- Algorithm Pipeline:



# More Recent Methods in Edge Detection

Crisp Boundary Detection using Pointwise Mutual Information (Isola et al, 2014)

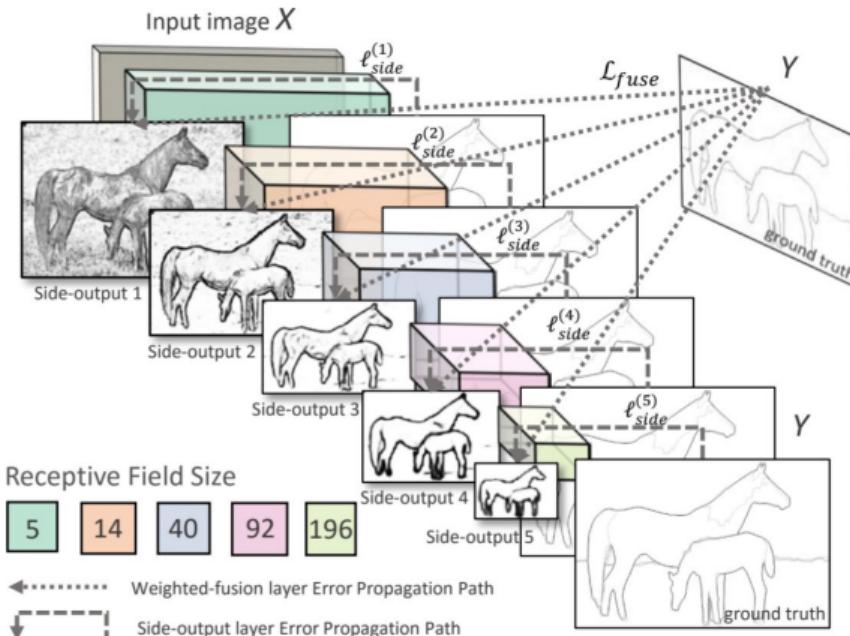


Algorithm	ODS	OIS	AP
Canny [14]	0.60	0.63	0.58
Mean Shift [36]	0.64	0.68	0.56
NCuts [37]	0.64	0.68	0.45
Felz-Hutt [38]	0.61	0.64	0.56
gPb [1]	0.71	0.74	0.65
gPb-owt-ucm [1]	0.73	0.76	0.73
SCG [9]	<b>0.74</b>	0.76	0.77
Sketch Tokens [7]	0.73	0.75	0.78
SE [8]	<b>0.74</b>	0.76	0.78
Our method – SS, color only	0.72	0.75	0.77
Our method – SS	0.73	0.76	<b>0.79</b>
Our method – MS	<b>0.74</b>	<b>0.77</b>	0.78

Evaluation on BSDS500

# More Recent Methods in Edge Detection

## Holistically Nested Edge Detection (Xie et al, 2015)

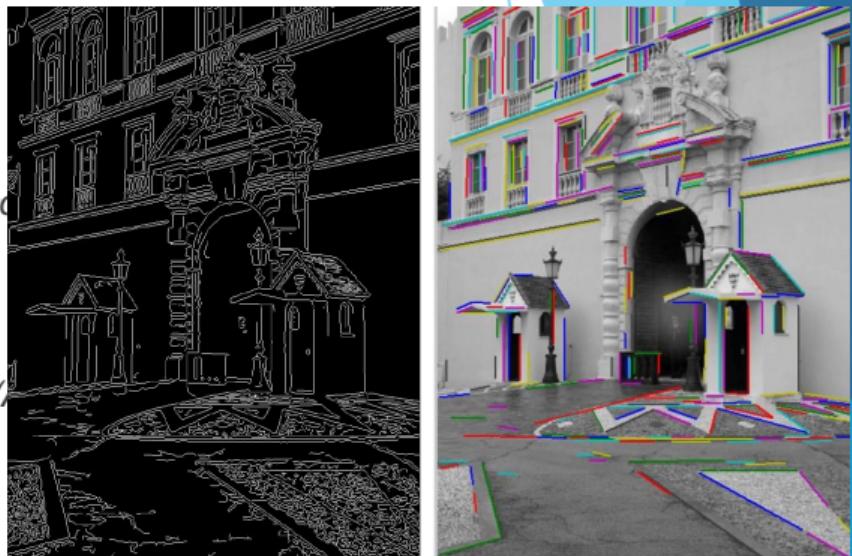


	ODS	OIS	AP	FPS
Human	.80	.80	-	-
Canny	.600	.640	.580	15
Felz-Hutt [9]	.610	.640	.560	10
BEL [5]	.660*	-	-	1/10
gPb-owt-ucm [1]	.726	.757	.696	1/240
Sketch Tokens [24]	.727	.746	.780	1
SCG [31]	.739	.758	.773	1/280
SE-Var [6]	.746	.767	.803	2.5
OEF [13]	.749	.772	.817	-
DeepNets [21]	.738	.759	.758	1/5†
N4-Fields [10]	.753	.769	.784	1/6†
DeepEdge [2]	.753	.772	.807	1/10 <sup>3</sup> †
CSCNN [19]	.756	.775	.798	-
DeepContour [34]	.756	.773	.797	1/30†
<b>HED (ours)</b>	<b>.782</b>	<b>.804</b>	<b>.833</b>	2.5†, 1/12

Source: Derek Hoiem

# Part A: Homework

- ▶ Readings
  - ▶ Chapter 2, Szeliski, *Computer Vision: Algorithms and Applications*
  
- ▶ Questions
  - ▶ How do you go from Canny edges to straight lines? (



# Part 2

From Edges to Blobs and Corners

# Review: Using Canny Edges to get Straight Lines

- Compute Canny edges
  - Compute  $\nabla_x f, \nabla_y f$  (gradients in  $x, y$  directions)
  - Compute  $\vartheta = \tan^{-1} \frac{\nabla_y f}{\nabla_x f}$
- Assign each edge to one of 8 directions. For each direction  $d$ , obtain “edgelets”:
  - Find connected components for edge pixels with directions in  $\{d - 1, d, d + 1\}$

# Review: Using Canny Edges to get Straight Lines

- Compute Canny edges
  - Compute  $\nabla_x f, \nabla_y f$  (gradients in  $x, y$  directions)
  - Compute  $\vartheta = \tan^{-1} \frac{\nabla_y f}{\nabla_x f}$
- Assign each edge to one of 8 directions. For each direction  $d$ , obtain “edgelets”:
  - Find connected components for edge pixels with directions in  $\{d - 1, d, d + 1\}$
- Compute straightness and orientation,  $\vartheta$  of edgelets using eigenvectors and eigenvalues of second moment matrix,  $M$ , of edge pixels

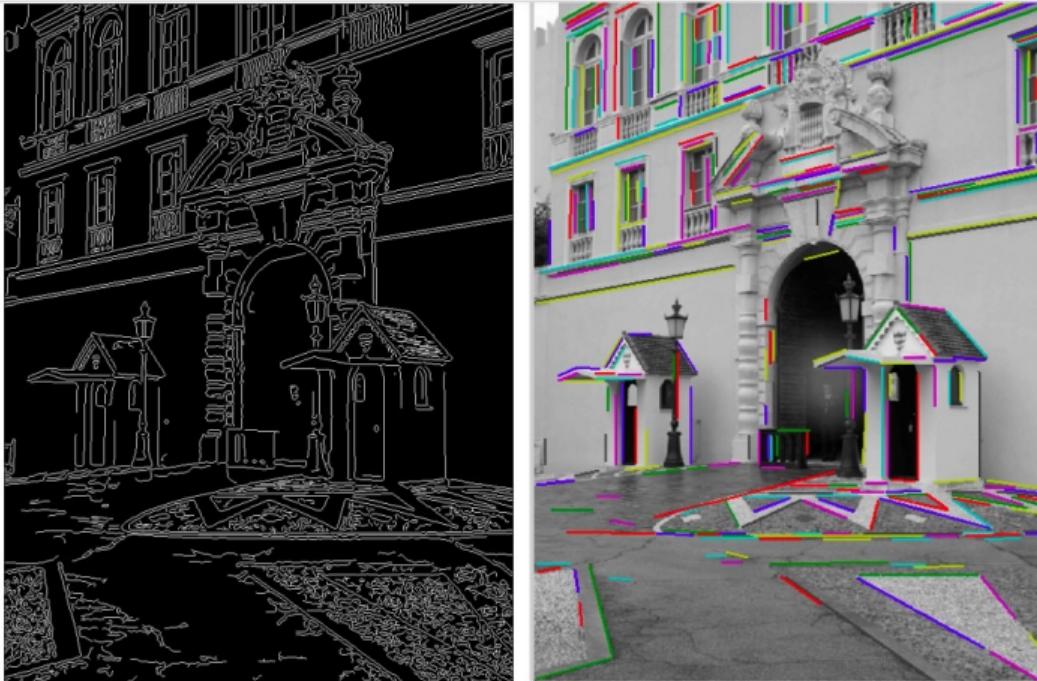
$$M = \begin{pmatrix} \sum (x - \mu_x)^2 & \sum (x - \mu_x)(y - \mu_y) \\ \sum (x - \mu_x)(y - \mu_y) & \sum (y - \mu_y)^2 \end{pmatrix} \quad [\mathbf{v}, \lambda] = \text{eig}(M)$$

$$\vartheta = \tan^{-1}(\mathbf{v}_1/\mathbf{v}_0) \text{ where } \mathbf{v}_1 \text{ is the 'larger' eigenvector;} \quad \text{straightness} = \lambda_2/\lambda_1$$

- Threshold straightness appropriately and store line segments

Credit: Derek Hoiem

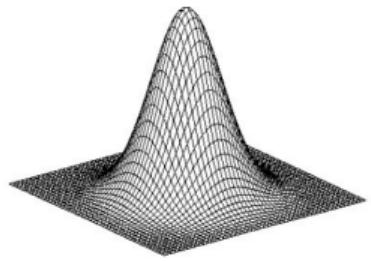
# Review: Using Canny Edges to get Straight Lines



Credit: Derek Hoiem

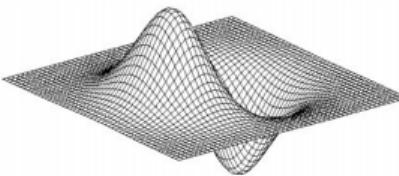
# Going Further to a Second Derivative

- What if we took the Laplacian of Gaussian?



Gaussian

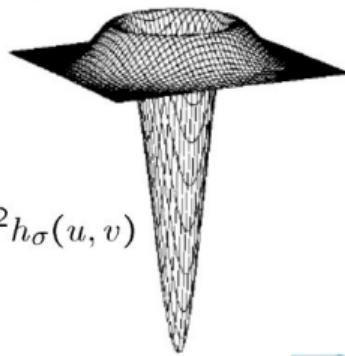
$$h_\sigma(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{2\sigma^2}}$$



Derivative of Gaussian

$$\frac{\partial}{\partial x} h_\sigma(u, v)$$

Laplacian of Gaussian



$$\text{Laplacian } \nabla^2 f = \frac{\partial^2 f}{\partial x^2} + \frac{\partial^2 f}{\partial y^2}$$

(1)

Credit: S Seitz, K Grauman

# Laplacian of Gaussian

Example of a  $3 \times 3$   
Laplacian of Gaussian  
filter:

$$\begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix}$$

# Laplacian of Gaussian



Original Image



Laplacian



Laplacian of Gaussian

# Laplacian of Gaussian



Original Image



Laplacian

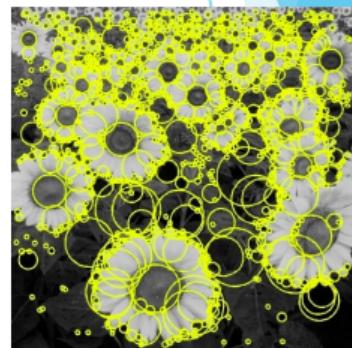
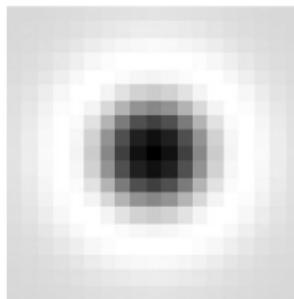
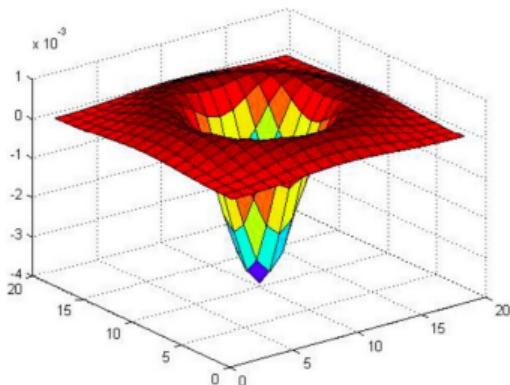


Laplacian of Gaussian

What else can LoG do?

# LoG as a Blob Detector

- Recall that convolution with a filter can be viewed as comparing a little “picture” of what you want to find against all local regions in the image.



- Observing the LoG filter matrix reveals that it is circularly symmetric. Thus it can be used for blob detection!

Credit: S Lazebnik

# From Blobs to Corners

- In the following image, what are some interesting features to choose?



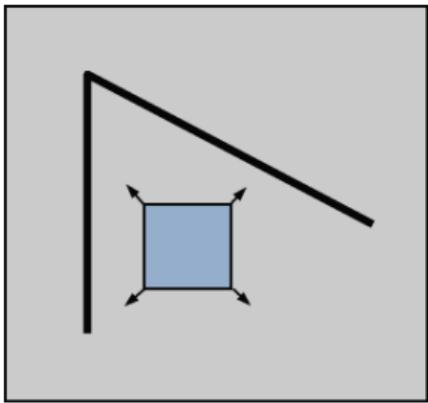
*Credit: K Grauman, R Urtasun*

# From Blobs to Corners

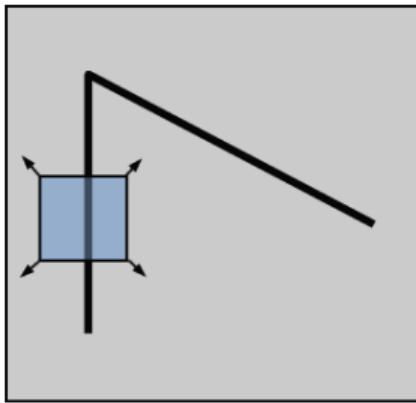
- Look for image regions that are unusual. How to define "unusual"?
- Textureless patches are nearly impossible to localize.
- Patches with **large contrast changes** (gradients) are easier to localize.
- But straight line segments at a single orientation suffer from the **aperture problem** (we'll see next slide), i.e., it is only possible to align the patches along the direction normal to the edge direction.
- Gradients in at least two (significantly) different orientations are the easiest, e.g., corners.

# From Blobs to Corners

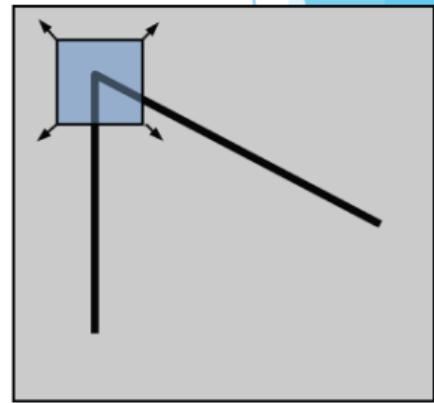
- Consider a small window of pixels. How does the window change when you shift it?



“flat” region:  
no change in all  
directions



“edge”:  
no change along the  
edge direction



“corner”:  
significant change in  
all directions

# Autocorrelation

- In the previous slide, how to quantify the "significant" change of the window?

## Part B: Homework

### Readings

Chapter 2, Szeliski, *Computer Vision: Algorithms and Applications*

### Questions: Linear Algebra review

Show that the trace of a matrix is the sum of its eigenvalues.

Show that the determinant of a matrix is the product of its eigenvalues.

# References

-  John F. Canny. "A Computational Approach to Edge Detection". In: *IEEE Transactions on Pattern Analysis and Machine Intelligence* PAMI-8 (1986), pp. 679–698.
-  David Martin, Charless Fowlkes, and Jitendra Malik. "Learning to Detect Natural Image Boundaries Using Local Brightness, Color, and Texture Cues". In: *IEEE Transactions on Pattern Analysis and Machine Intelligence* 26 (June 2004), pp. 530–49.
-  Richard Szeliski. *Computer Vision: Algorithms and Applications*. Texts in Computer Science. London: Springer-Verlag, 2011.
-  Piotr Dollár and Lawrence Zitnick. "Structured Forests for Fast Edge Detection". In: *Proceedings of the International Conference on Computer Vision*. IEEE, Dec. 2013.
-  Phillip Isola et al. "Crisp Boundary Detection Using Pointwise Mutual Information". In: *Proceedings of the European Conference on Computer Vision*. 2014.
-  Saining Xie and Zhuowen Tu. "Holistically-Nested Edge Detection". In: *International Journal of Computer Vision* 125 (2015), pp. 3–18.

## References



▶ C. Harris and M. Stephens. “A Combined Corner and Edge Detector”. In: *Proceedings of the 4th Alvey Vision Conference*. 1988, pp. 147-151.



▶ Bill Triggs. “Detecting Keypoints with Stable Position, Orientation, and Scale under Illumination Changes”. In: *Computer Vision - ECCV 2004*. Ed. by Tomáš Pajdla and Jiří Matas. Berlin, Heidelberg: Springer Berlin Heidelberg, 2004, pp. 100-113.



▶ M. Brown, R. Szeliski, and S. Winder. “Multi-image matching using multi-scale oriented patches”. In: *2005 IEEE Computer Society Conference on Computer Vision and Pattern Recognition (CVPR'05)*. Vol. 1. 2005, 510-517 vol. 1.



▶ Richard Szeliski. *Computer Vision: Algorithms and Applications*. Texts in Computer Science. London: Springer-Verlag, 2011.

▶ David Forsyth and Jean Ponce. *Computer Vision: A Modern Approach*. 2 edition. Boston: Pearson Education India, 2015.