

Vision Systems

Lecture 8

Part 1

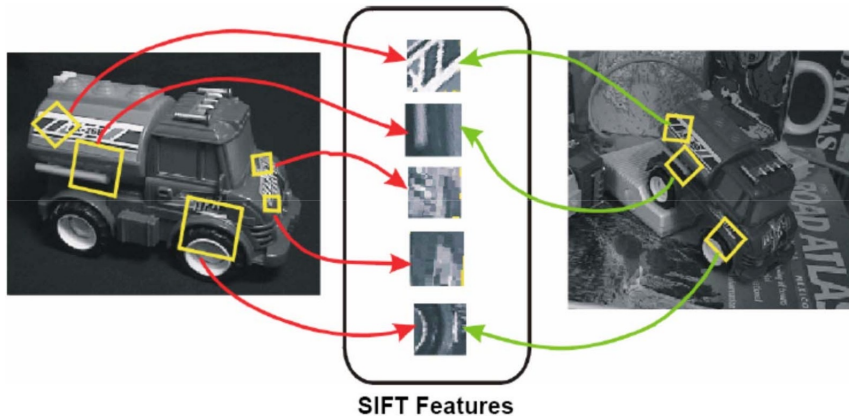
Feature Detectors: SIFT and Variants

SIFT: Scale Invariant Feature Transform

- David G. Lowe, *Distinctive Image Features from Scale-invariant Keypoints*, IJCV 2004
 - Over 50000 citations
- Transforms image data into scale-invariant coordinates
- Fundamental to many core vision problems/applications:
 - Recognition, Motion tracking, Multiview geometry

SIFT: Invariant Local Features

Image content is transformed into local feature coordinates that are invariant to translation, rotation, scale, shear.




SIFT

- **Step 1: Scale-space Extrema Detection** - Detect interesting points (invariant to scale and orientation) using DOG.
- **Step 2: Keypoint Localization** - Determine location and scale at each candidate location, and select them based on stability.
- **Step 3: Orientation Estimation** - Use local image gradients to assign orientation to each localized keypoint. Preserve orientation, scale and location for each feature.
- **Step 4: Keypoint Descriptor** - Extract local image gradients at selected scale around keypoint and form a representation invariant to local shape and illumination distortion.

SIFT: Scale-space Extrema Detection

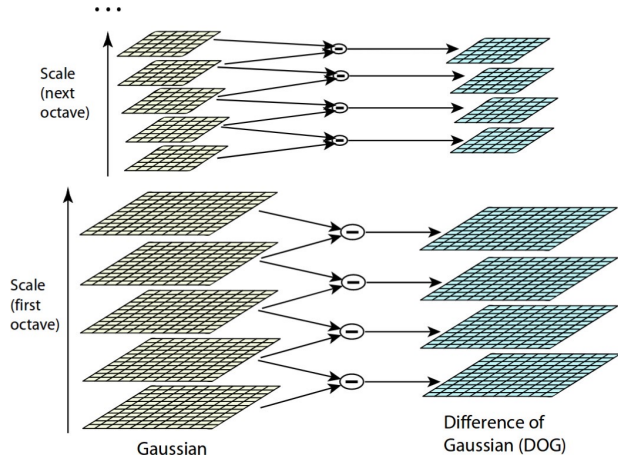
Constructing Scale Space



Credit: Ofir Pele


SIFT: Scale-space Extrema Detection

Difference of Gaussians




SIFT: Scale-space Extrema Detection

Difference of Gaussians




$$D(x, y, \sigma) = \hat{I}(x, y, \sigma) - \hat{I}(x, y, k\sigma)$$

where $\hat{I}(x, y, \sigma) = G(x, y, \sigma) * I(x, y)$



Credit: "Distinctive Image Features from Scale-Invariant Points", IJCV 2004


SIFT: Scale-space Extrema Detection



- Compare a pixel (X) with 26 pixels in current and adjacent scales (Green Circles)
- Select a pixel (X) if it is larger/smaller than all 26 pixels

Credit: "Distinctive Image Features from Scale-Invariant Points", IJCV 2004

SIFT: Scale-space Extrema Detection



- Compare a pixel (X) with 26 pixels in current and adjacent scales (Green Circles)
- Select a pixel (X) if it is larger/smaller than all 26 pixels


Credit: "Distinctive Image Features from Scale-Invariant Points", IJCV 2004

SIFT Algorithm Stages

- ▶ **Step 1: Scale-space extrema Detection** - Detect interesting points (invariant to scale and orientation) using DOG.
- ▶ **Step 2: Keypoint Localization** - Determine location and scale at each candidate location, and select them based on stability.
- ▶ **Step 3: Orientation Estimation** - Use local image gradients to assign orientation to each localized keypoint. Preserve orientation, scale and location for each feature.
- ▶ **Step 4: Keypoint Descriptor** - Extract local image gradients at selected scale around keypoint and form a representation invariant to local shape and illumination distortion them.

SIFT: Keypoint Localization

The Problem:



The Solution:

Use Taylor series expansion of the scale-space function:

SIFT: Keypoint Localization

The Solution:

- Use Taylor series expansion of the scale-space function:

$$D(\mathbf{s}_0 + \Delta \mathbf{s}) = D(\mathbf{s}_0) + \frac{\partial D}{\partial \mathbf{s}} \bigg|_{\mathbf{s}_0} \Delta \mathbf{s} + \frac{1}{2} \Delta \mathbf{s}^T \frac{\partial^2 D}{\partial^2 \mathbf{s}} \bigg|_{\mathbf{s}_0} \Delta \mathbf{s}$$

where $\mathbf{s}_0 = (x_0, y_0, \sigma_0)^T$ and $\Delta \mathbf{s} = (\delta x, \delta y, \delta \sigma)^T$

- The location of the extremum, $\hat{\mathbf{s}}$, is determined by taking the derivative of this function with respect to \mathbf{s} and setting it to zero:

$$\hat{\mathbf{s}} = - \left(\frac{\partial^2 D}{\partial^2 \mathbf{s}} \bigg|_{\mathbf{s}_0} \right)^{-1} \frac{\partial D}{\partial \mathbf{s}} \bigg|_{\mathbf{s}_0}$$

- Next, reject low contrast points and points that lie on the edge
- **Low contrast points elimination:**
 - Reject keypoint if $D(\hat{\mathbf{s}})$ is smaller than 0.03 (assuming image values are normalized in $[0,1]$)

SIFT Algorithm Stages

- ▶ **Step 1: Scale-space extrema Detection** - Detect interesting points (invariant to scale and orientation) using DOG.
- ▶ **Step 2: Keypoint Localization** - Determine location and scale at each candidate location, and select them based on stability.
- ▶ **Step 3: Orientation Estimation** - Use local image gradients to assign orientation to each localized keypoint. Preserve orientation, scale and location for each feature.
- ▶ **Step 4: Keypoint Descriptor** - Extract local image gradients at selected scale around keypoint and form a representation invariant to local shape and illumination distortion them.

SIFT: Orientation Estimation

■ Why?

SIFT: Orientation Estimation

- Why? To achieve rotation invariance

SIFT: Orientation Estimation

- Why? To achieve rotation invariance
- Use scale of point to choose correct image:

$$\hat{I}(x, y) = G(x, y, \sigma) * I(x, y)$$


- Compute gradient magnitude and orientation using finite differences:

$$m(x, y) = \sqrt{(\hat{I}(x+1, y) - \hat{I}(x-1, y))^2 + (\hat{I}(x, y+1) - \hat{I}(x, y-1))^2}$$

$$\vartheta(x, y) = \tan^{-1} \frac{(\hat{I}(x, y+1) - \hat{I}(x, y-1))}{(\hat{I}(x+1, y) - \hat{I}(x-1, y))}$$


SIFT: Orientation Estimation

- Create histogram of gradient directions, within a region around the keypoint, at selected scale:



SIFT: Orientation Estimation


- Create histogram of gradient directions, within a region around the keypoint, at selected scale:



- Histogram entries are weighted by:
 - gradient magnitude, and
 - a Gaussian function with σ equal to 1.5 times scale of the keypoint

SIFT: Orientation Estimation


- Create histogram of gradient directions, within a region around the keypoint, at selected scale:



- Histogram entries are weighted by:
 - gradient magnitude, and
 - a Gaussian function with σ equal to 1.5 times scale of the keypoint
- Select **the peak as direction of keypoint**

SIFT: Orientation Estimation

- Create histogram of gradient directions, within a region around the keypoint, at selected scale:



- Histogram entries are weighted by:
 - gradient magnitude, and
 - a Gaussian function with σ equal to 1.5 times scale of the keypoint
- Select **the peak as direction of keypoint**
- Introduce additional key points at same location if another peak is within 80% of max peak of histogram with different direction

SIFT



From 233x189 original image to 832 DoG Extrema

Credit: Mubarak Shah, University of Central Florida

SIFT



From 832 DoG Extrema to 729 keypoints after low contrast threshold

Credit: Mubarak Shah, University of Central Florida

SIFT



From 729 keypoints to 536 keypoints after testing ratio based on Hessian


Credit: Mubarak Shah, University of Central Florida

SIFT Algorithm Stages

- ▶ **Step 1: Scale-space extrema Detection** - Detect interesting points (invariant to scale and orientation) using DOG.
- ▶ **Step 2: Keypoint Localization** - Determine location and scale at each candidate location, and select them based on stability.
- ▶ **Step 3: Orientation Estimation** - Use local image gradients to assign orientation to each localized keypoint. Preserve orientation, scale and location for each feature.
- ▶ **Step 4: Keypoint Descriptor** - Extract local image gradients at selected scale around keypoint and form a representation invariant to local shape and illumination distortion.


SIFT: Keypoint Descriptor

- Compute gradient at each pixel in a 16×16 window around the detected keypoint, using the appropriate level ϵ as detected.



SIFT: Keypoint Descriptor


- Compute gradient at each pixel in a 16×16 window around the detected keypoint, using the appropriate level ϵ as detected.



- ▶ Downweight gradients by a Gaussian fall-off function (blue circle) to reduce the influence of gradients far from the center.

SIFT: Keypoint Descriptor

- Compute gradient at each pixel in a 16×16 window around the detected keypoint, using the appropriate level ϵ as detected.



- Downweight gradients by a Gaussian fall-off function (blue circle) to reduce the influence of gradients far from the center.
- In each 4×4 quadrant, compute a gradient orientation histogram using 8 orientation histogram bins.

SIFT: Keypoint Descriptor

- The resulting 128 non-negative values form a raw version of the SIFT descriptor vector.
- To reduce the effects of contrast or gain (additive variations are already removed by the gradient), the 128-D vector is normalized to unit length.
- To further make the descriptor robust to other photometric variations, values are clipped to 0.2 and the resulting vector is once again renormalized to unit length.

Credit: Raquel Urtasun, Szeliski


SIFT

- Extraordinarily robust feature detection
- Changes in viewpoint: up to about 60 degree out of plane rotation
- Changes in illumination: sometimes even day vs night (below)
- Fast and efficient - can run in real-time



Credit: Raquel Urtasun, Szeliski

SIFT: Example



Mars Rover images


Credit: Raquel Urtasun, N Snavely

SIFT: Example

Maybe, look for tiny squares...?

SIFT: Example



Maybe, look for tiny squares...



Mars Rover images with SIFT feature matches

Credit: Raquel Urtasun, N Snavely

SIFT: Invariances(Geometric Transformations)



e.g. scale,
translation,
rotation

SIFT: Invariances(Photometric Transformations)




Credit: Raquel Urtasun, Tinne Tuytelaars

SIFT Applications: Image Stitching




SIFT Applications: Image Stitching



- Detect feature points in both images.

SIFT Applications: Image Stitching



- Detect feature points in both images.
- Find corresponding pairs of feature points.

SIFT Applications: Image Stitching



- Detect feature points in both images.
- Find corresponding pairs of feature points.
- Use the pairs to align the images.

Credit: Raquel Urtasun

More Resources

If you want to learn more about SIFT

[The SIFT Keypoint Detector](#) by David Lowe

Tutorial: [SIFT \(Scale-invariant feature transform\)](#)

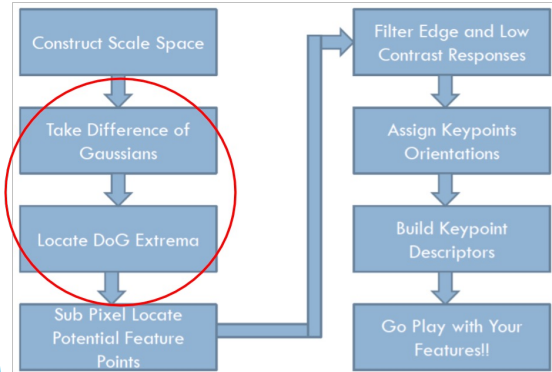
OpenCV-Python Tutorials: [Introduction to SIFT](#)

Wikipedia: [Scale-invariant feature transform](#)


[OpenSIFT: An Open-Source SIFT Library](#)

SURF: Speeded Up Robust Features

SIFT



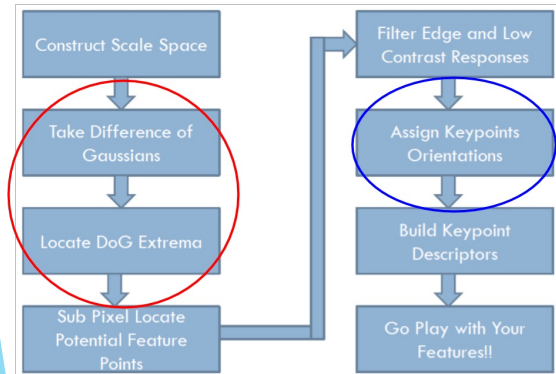
- Uses box filters instead of Gaussians to approximate Laplacians



- Uses Haar wavelets to get keypoint orientations

SURF: Speeded Up Robust Features

SIFT



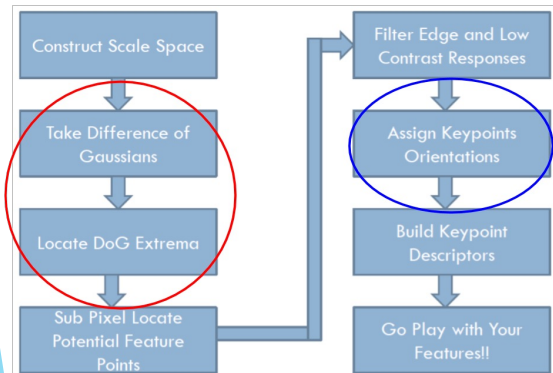
- Uses box filters instead of Gaussians to approximate Laplacians
- Uses Haar wavelets to get keypoint orientations
 - Haar wavelets are simple filters which can be used to find gradients in the x and y directions



- SURF is good at handling blur and rotation variations
- SURF is not as good as SIFT on invariance

SURF: Speeded Up Robust Features

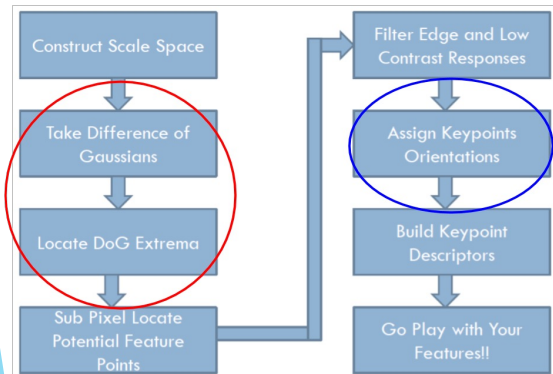
SIFT



- Uses box filters instead of Gaussians to approximate Laplacians
- Uses Haar wavelets to get keypoint orientations
- SURF is good at handling blur and rotation variations
- SURF is not as good as SIFT on invariance to illumination and viewpoint changes
- SURF is ~ 3 times faster than SIFT

SURF: Speeded Up Robust Features

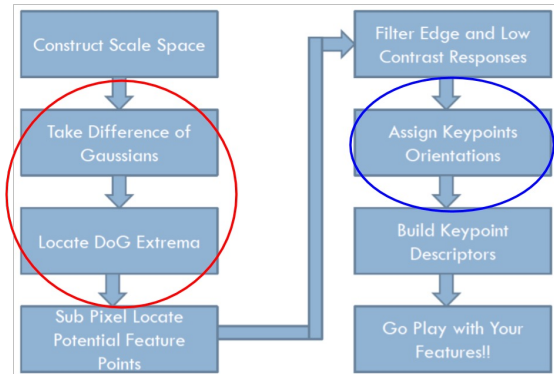
SIFT



- Uses box filters instead of Gaussians to approximate Laplacians
- Uses Haar wavelets to get keypoint orientations
- SURF is good at handling blur and rotation variations
- SURF is not as good as SIFT on invariance to illumination and viewpoint changes
- SURF is ~ 3 times faster than SIFT

SURF: Speeded Up Robust Features

SIFT




- Uses box filters instead of Gaussians to approximate Laplacians
- Uses Haar wavelets to get keypoint orientations
- SURF is good at handling blur and rotation variations
- SURF is not as good as SIFT on invariance to illumination and viewpoint changes
- SURF is ~ 3 times faster than SIFT

For more information:

- <https://medium.com/data-breach/introduction-to-surf-speeded-up-robust-features-c096c1b7c4e>
- <http://www.vision.ee.ethz.ch/~surf/>


MOPS: Making Descriptor Rotation-invariant

- **M**ultiscale **O**riented **P**atches descriptor
- Rotate patch according to its dominant gradient orientation.
- This puts the patches into a canonical orientation



Gradient Location-Orientation Histogram (GLOH)

- Variant of SIFT that uses a log-polar binning structure instead of four quadrants.
- Uses 17 spatial bins and 16 orientation bins.
- The 272D histogram is then projected onto a 128D descriptor using PCA trained on a large dataset.



Homework

Readings

- Chapter 2 and Chapter 3, Szeliski, *Computer Vision: Algorithms and Applications*
- Other links provided on respective slides
- [Multi-Scale Oriented Patches](#)
- For more information on SURF:
 - OpenCV-Python Tutorials: [Introduction to SURF](#) Wikipedia: [Speeded up robust features](#)

Labs

- *Practice the SFIT and SURF algorithms using OpenCV*

Questions

Which descriptor performs better? SIFT or MOPS?

Why is SIFT descriptor better than Harris Corner Detector?

References



David G. Lowe. "Distinctive Image Features from Scale-Invariant Keypoints". In: *Int. J. Comput. Vision* 60.2 (Nov. 2004), 91–110.



Richard Szeliski. *Computer Vision: Algorithms and Applications*. Texts in Computer Science. London: Springer-Verlag, 2011.



David Forsyth and Jean Ponce. *Computer Vision: A Modern Approach*. 2 edition. Boston: Pearson Education India, 2015.



Lazebnik, Svetlana, *CS 543 Computer Vision (Spring 2019)*. URL: <https://slazebni.cs.illinois.edu/spring19/> (visited on 06/01/2020).



Shah, Mubarak, *CAP 5415 - Computer Vision (Fall 2014)*. URL: <https://www.crcv.ucf.edu/courses/cap5415-fall-2014/> (visited on 06/01/2020).



Urtasun, Raquel, *Computer Vision (Winter 2013)*. URL: <https://www.cs.toronto.edu/~urtasun/courses/CV/cv.html> (visited on 06/01/2020).