

PyTorch

What is PYTORCH?

It's a Python based scientific computing package targeted at two use cases:

- A replacement for NumPy to use the power of GPUs
- A deep learning research platform that provides maximum flexibility and speed

PyTorch: A Python Based Framework

- A complete rewrite of Torch using Python
- With critical functions written in C/C++

PYTORCH + python™



PyTorch: A NumPy Replacement with GPU Acceleration

- Provides NumPy like capabilities
- `torch.Tensor` is similar to `numpy.ndarray`
- With support for both CPU* and GPU



Fig: Tensor operations

*NumPy only supports CPU

PyTorch: A Deep Learning Research Platform

- Allows you to solve problems using deep learning
- Caters to use cases in computer vision, text, speech, and so on

Package	Description
torch	a Tensor library like NumPy, with strong GPU support
torch.autograd	a tape based automatic differentiation library that supports all differentiable Tensor operations in torch
torch.nn	a neural networks library deeply integrated with autograd designed for maximum flexibility
torch.optim	an optimization package to be used with torch.nn with standard optimization methods such as SGD, RMSProp, LBFGS, Adam etc.
torch.multiprocessing	python multiprocessing, but with magical memory sharing of torch Tensors across processes. Useful for data loading and hogwild training.
torch.utils	DataLoader, Trainer and other utility functions for convenience

Fig: Packages in PyTorch




Fig: Object instance segmentation, a user case in computer vision

Salient Features

- Python first
- Use of Dynamic Computation Graph
- Intuitive programming model
- Easier debugging

Salient Features: Python First

- Deeply integrated with Python
- PyTorch computations run within the Python computation model
- It's imperative, just like Python
- PyTorch can be extended, just like you would extend Python

Salient Features: Use of Dynamic Computation Graph (1/2)

- PyTorch uses a Dynamic computation graph (versus Static)
- PyTorch: Computation graph gets created on the fly

Back-propagation
uses the dynamically built graph

```
from torch.autograd import Variable

x = Variable(torch.randn(1, 10))
prev_h = Variable(torch.randn(1, 20))
W_h = Variable(torch.randn(20, 20))
W_x = Variable(torch.randn(20, 10))

i2h = torch.mm(W_x, x.t())
h2h = torch.mm(W_h, prev_h.t())
next_h = i2h + h2h
next_h = next_h.tanh()

next_h.backward(torch.ones(1, 20))
```




Fig: PyTorch dynamic graph

Salient Features: Use of Dynamic Computation Graph (2/2)

- PyTorch uses a Dynamic computation graph (versus Static)
- **Dynamic** computation graphs use the **imperative** style of programming
- **Static** computation graphs use the **declarative** style of programming

```
#include <stdio.h>

int main(void)
{
    printf("hello world\n");
    return 0;
}
```

C, declarative

```
>>> print("Hello World")
Hello World
>>>
```

Python, imperative

Salient Features: Intuitive Programming Model

- Define and run
- Linear in thought
- Faster prototyping

Salient Features: Easier Debugging

- No separate virtual execution environment
- Uses host language's runtime
- Debugging Pytorch code is just like debugging Python code
- Use the same debugging tools

PyTorch: Platforms Supported

OS	<input checked="" type="radio"/> Linux	<input type="radio"/> OSX		
Package Manager	<input checked="" type="radio"/> conda	<input type="radio"/> pip	<input type="radio"/> Source	
Python	<input type="radio"/> 2.7	<input checked="" type="radio"/> 3.5	<input type="radio"/> 3.6	
CUDA	<input checked="" type="radio"/> 8	<input type="radio"/> 9.0	<input type="radio"/> 9.1	<input type="radio"/> None

Refresher: Scalar, Vector, Matrix, Tensor, Rank, and Dimension

3, 1.2

Scalar

Rank: 0 (always)
Dim: ()

[3, 4, 8]

Vector

Rank: 1 (always)
Dim: (3,)

[[1, 2, 3],
[4, 5, 6]]

Matrix

Rank: 2 (always)
Dim: (2, 3)

[[[1, 2, 3], [4, 5, 6]],
[[7, 8, 9], [10, 11, 12]]]

Tensor

Rank: 3
Dim: (2, 2, 3)

Tensors in PyTorch

- A multi-dimensional matrix containing elements of a single data type
- Part of the torch package
- Instantiated via `torch.Tensor` group of classes
- Can be stored on the CPU or GPU
- Operated on via functions available in the torch package, or via class methods
- Interoperable with NumPy array

```
>>> torch.FloatTensor([[1, 2, 3], [4, 5, 6]])  
1  2  3  
4  5  6  
[torch.FloatTensor of size 2x3]
```

Tensors in PyTorch

Data type	CPU tensor	GPU tensor
32-bit floating point	<code>torch.FloatTensor</code>	<code>torch.cuda.FloatTensor</code>
64-bit floating point	<code>torch.DoubleTensor</code>	<code>torch.cuda.DoubleTensor</code>
16-bit floating point	<code>torch.HalfTensor</code>	<code>torch.cuda.HalfTensor</code>
8-bit integer (unsigned)	<code>torch.ByteTensor</code>	<code>torch.cuda.ByteTensor</code>
8-bit integer (signed)	<code>torch.CharTensor</code>	<code>torch.cuda.CharTensor</code>
16-bit integer (signed)	<code>torch.ShortTensor</code>	<code>torch.cuda.ShortTensor</code>
32-bit integer (signed)	<code>torch.IntTensor</code>	<code>torch.cuda.IntTensor</code>
64-bit integer (signed)	<code>torch.LongTensor</code>	<code>torch.cuda.LongTensor</code>

Fig: Tensor types

Refresher: Differentiation

Model Model parameter

↓ ↓

If $y = f(x) = 2x$

then $\frac{dy}{dx} = 2$

If $y = f(x_1, x_2, \dots, x_n)$

then $\left[\frac{dy}{dx_1}, \frac{dy}{dx_2}, \dots, \frac{dy}{dx_n} \right]$

Is the gradient of y w.r.t. $[x_1, x_2, \dots, x_n]$

Function $f(x)$	Derivative $f'(x)$
a	0
x	1
$ax + b$	a
$\frac{1}{x}$	$-\frac{1}{x^2}$
x^a	ax^{a-1}
a^x	$a^x \ln a$
$\sqrt[n]{a}$	$-\frac{\sqrt[n]{a} \ln a}{x^2}$
$\log_a x$	$\frac{1}{x \ln a}$
x^x	$x^x (1 + \ln x)$
$\Gamma(x)$	$\Gamma(x)\psi(x)$

Refresher: The Computation Graph




Fig: A computation graph

Refresher: The Computation Graph

Back-propagation
uses the dynamically built graph


```
from torch.autograd import Variable

x = Variable(torch.randn(1, 10))
prev_h = Variable(torch.randn(1, 20))
W_h = Variable(torch.randn(20, 20))
W_x = Variable(torch.randn(20, 10))

i2h = torch.mm(W_x, x.t())
h2h = torch.mm(W_h, prev_h.t())
next_h = i2h + h2h
next_h = next_h.tanh()

next_h.backward(torch.ones(1, 20))
```

Fig: PyTorch dynamic graph



Variables in PyTorch

- Crucial data structure, needed for automatic differentiation
- Part of the `torch.autograd` package
- Instantiated via `torch.autograd.Variable` class
- A wrapper around a tensor object (the data)
- Holds the gradient w.r.t. it (the grad)
- Records reference to the function that created it (the creator)
- Holds the gradient of output w.r.t this tensor

autograd.Variable

data *grad*

creator

Interoperability Between PyTorch Tensors and NumPy Arrays

- `torch.Tensor` similar to `numpy.ndarray`
- 200+ operations, similar to `numpy`
- Memory pointer shared between PyTorch tensor and NumPy array
- Conversions very fast: Zero memory copy

`# -*- coding: utf-8 -*-
import numpy as np

N is batch size; D_in is input dimension;
H is hidden dimension; D_out is output dimension.
N, D_in, H, D_out = 64, 1000, 100, 10

Create random input and output data
x = np.random.randn(N, D_in)
y = np.random.randn(N, D_out)

Randomly initialize weights
w1 = np.random.randn(D_in, H)
w2 = np.random.randn(H, D_out)

learning_rate = 1e-6
for t in range(500):
 # Forward pass: compute predicted y
 h = x.dot(w1)
 h_relu = np.maximum(h, 0)
 y_pred = h_relu.dot(w2)

 # Compute and print loss
 loss = np.square(y_pred - y).sum()
 print(t, loss)

 # Backprop to compute gradients of w1 and w2 with respect to loss
 grad_y_pred = 2.0 * (y_pred - y)
 grad_w2 = h_relu.T.dot(grad_y_pred)
 grad_h_relu = grad_y_pred.dot(w2.T)
 grad_h = grad_h_relu.copy()
 grad_h[h < 0] = 0
 grad_w1 = x.T.dot(grad_h)

 # Update weights
 w1 -= learning_rate * grad_w1
 w2 -= learning_rate * grad_w2`

NumPy

`import torch

dtype = torch.FloatTensor
dtype = torch.cuda.FloatTensor # Uncomment this to run on GPU

N is batch size; D_in is input dimension;
H is hidden dimension; D_out is output dimension.
N, D_in, H, D_out = 64, 1000, 100, 10

Create random input and output data
x = torch.randn(N, D_in).type(dtype)
y = torch.randn(N, D_out).type(dtype)

Randomly initialize weights
w1 = torch.randn(D_in, H).type(dtype)
w2 = torch.randn(H, D_out).type(dtype)

learning_rate = 1e-6
for t in range(500):
 # Forward pass: compute predicted y
 h = x.mm(w1)
 h_relu = h.clamp(min=0)
 y_pred = h_relu.mm(w2)

 # Compute and print loss
 loss = (y_pred - y).pow(2).sum()
 print(t, loss)

 # Backprop to compute gradients of w1 and w2 with respect to loss
 grad_y_pred = 2.0 * (y_pred - y)
 grad_w2 = h_relu.t().mm(grad_y_pred)
 grad_h_relu = grad_y_pred.mm(w2.t())
 grad_h = grad_h_relu.clone()
 grad_h[h < 0] = 0
 grad_w1 = x.t().mm(grad_h)

 # Update weights using gradient descent
 w1 -= learning_rate * grad_w1
 w2 -= learning_rate * grad_w2`

PyTorch

Handling Datasets in PyTorch

Concepts: Dataset, Epoch, Batch, Iteration

- Dataset: A collection of training examples
- Epoch: One pass of the entire dataset through your model during training
- Batch: A subset of training examples passed through your model at a time
- Iteration: An iteration is a single pass of a batch
- Here, a pass would involve a forward and a backward propagation

Accessing Custom Vision Datasets

- Create a custom Dataset class which inherits from `torch.utils.data.Dataset`
- Use `torch.utils.data.DataLoader` to iterate through the data
- Tutorial: http://pytorch.org/tutorials/beginner/data_loading_tutorial.html

For a dataset of 1000 images



Deep Learning using PyTorch

Tasks in Computer Vision

Classification



CAT


Classification
+ Localization



CAT

Single Object

Object Detection



CAT, DOG

Instance Segmentation



HUMAN, CAR

Multiple Objects

Tasks in Natural Language Processing

Positive

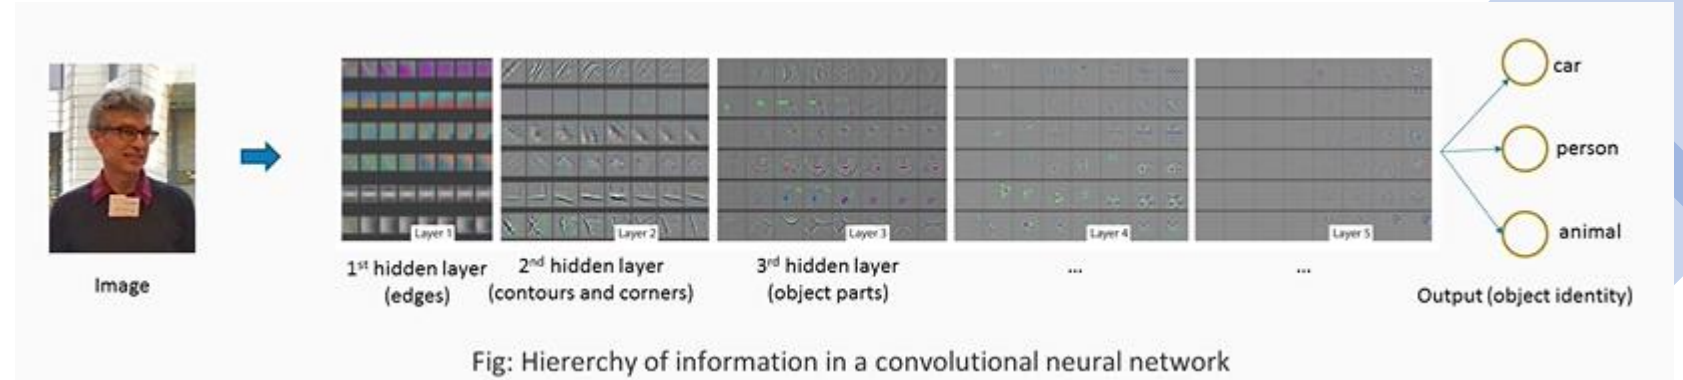

Covers pretty much everything one can require in a book to get more than an intuition for deep learning. From the math, machine learning to the modern practices in deep learning and the ongoing research in a pretty formal manner. The hardbound cover and the print quality is very nice too. 10/10 would recommend.

Sentiment Analysis



Machine Translation

What is Deep Learning?



PyTorch Packages for Deep Learning

Package	Description
torch	a Tensor library like NumPy, with strong GPU support
torch.autograd	a tape based automatic differentiation library that supports all differentiable Tensor operations in torch
torch.nn	a neural networks library deeply integrated with autograd designed for maximum flexibility
torch.optim	an optimization package to be used with torch.nn with standard optimization methods such as SGD, RMSProp, LBFGS, Adam etc.
torch.multiprocessing	python multiprocessing, but with magical memory sharing of torch Tensors across processes. Useful for data loading and hogwild training.
torch.utils	DataLoader, Trainer and other utility functions for convenience
torch.legacy(.nn/.optim)	legacy code that has been ported over from torch for backward compatibility reasons

Train Your First Neural Network


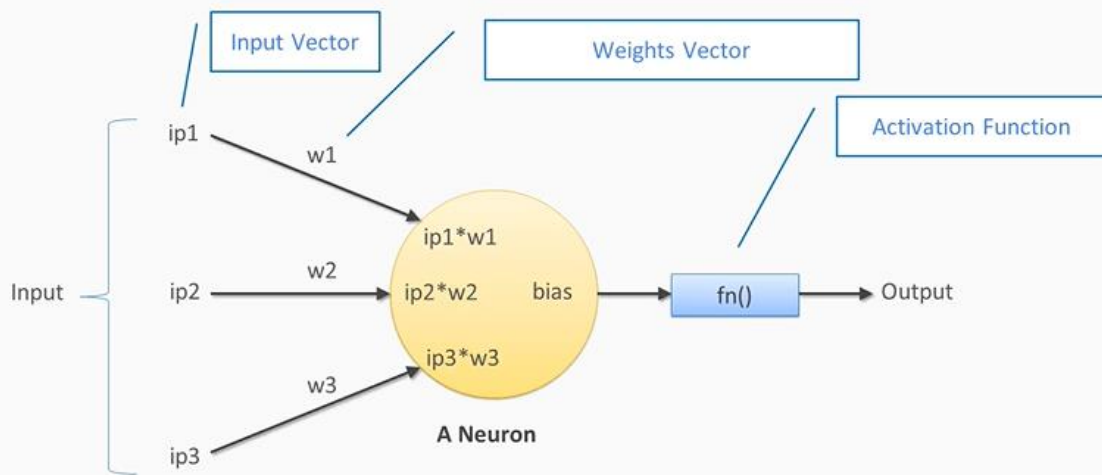


Fig: A simple neural network

Training a Neural Network with PyTorch

A Neuron




$$\text{Output} = \text{fn}(w1 * \text{ip1} + w2 * \text{ip2} + w3 * \text{ip3} + \text{bias})$$

Activation Functions


Sigmoid()

$$f(x) = \frac{1}{1 + e^{-x}}$$



Tanh()


$$f(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$




ReLU()

(Rectified Linear Unit)

$$f(x) = \max(x, 0)$$

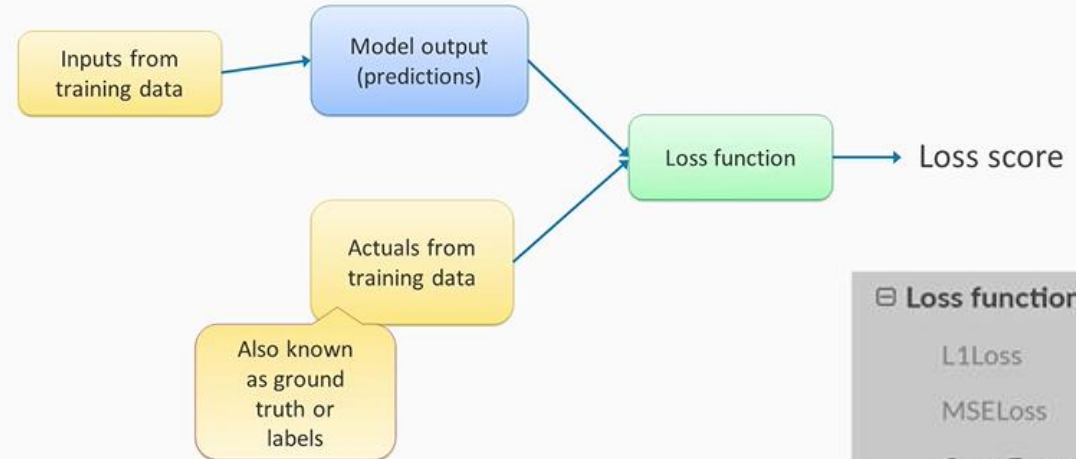


Artificial Neural Network



Hidden Layer: 3 incoming connections, 4 outgoing connections

Loss Function



Loss functions

L1Loss

MSELoss

CrossEntropyLoss

NLLLoss

PoissonNLLLoss

NLLLoss2d

KLDivLoss

BCELoss

BCEWithLogitsLoss

MarginRankingLoss

HingeEmbeddingLoss

Cross Entropy Loss

For a training example:

Class:	0 (Apple)	1 (Mango)	2 (Banana)
Prediction:	0.02	0.88	0.1

Model Predictions

Class: 1

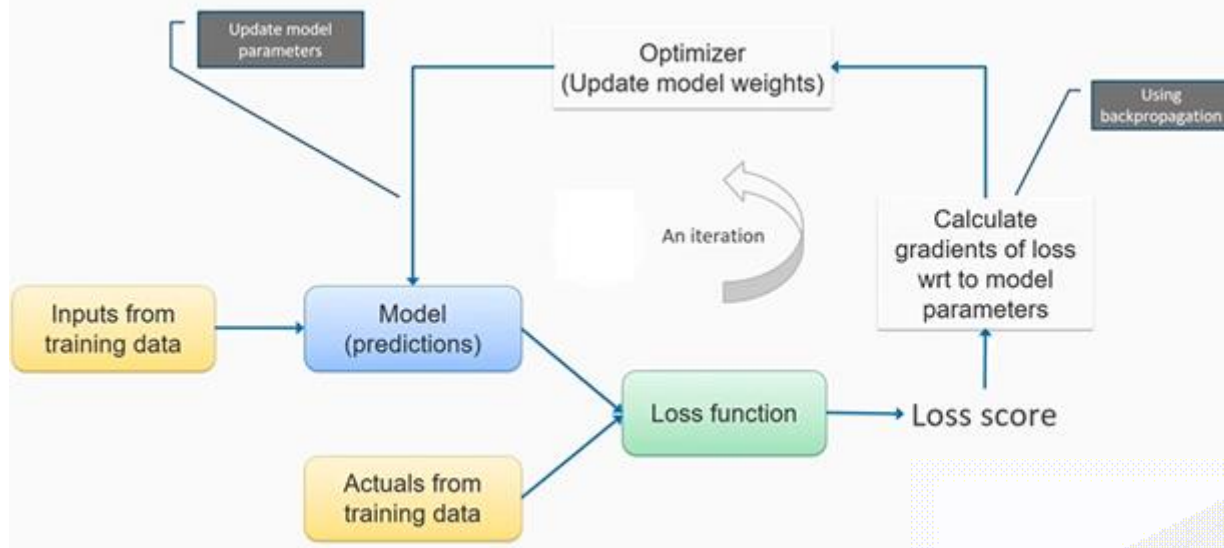
Actual/Expected



Low loss score

Optimization Problem



- An optimization problem is the problem of finding the best solution from all feasible solutions, given some criteria – Wikipedia
- Example: Which line to stand in at the supermarket



Available Optimizer

- SGD
- Adadelata
- Adagrad
- Adam
- RMSprop

```
optimizer = optim.SGD(model.parameters(), lr = 0.01, momentum=0.9)
optimizer = optim.Adam([var1, var2], lr = 0.0001)
```



The Fashion MNIST Data Set


- Is made up of Zalando's images
 - Contains 60,000 images
 - Also has a test set of 10,000 images
- Images 28 x 28 pixels
 - Items fall into 10 categories



Labels

Label	Description
0	T-shirt/top
1	Trouser
2	Pullover
3	Dress
4	Coat


Label	Description
5	Sandal
6	Shirt
7	Sport shoe
8	Bag
9	Ankle boot



Plot of ankle boots, sports shoes & sandals



RELU




Classes Overview

```
class FMNIST(nn.Module):
    def __init__(self):
        super().__init__()
        self.fc1 = nn.Linear(784, 128)
        self.fc2 = nn.Linear(128, 64)
        self.fc3 = nn.Linear(64, 10)


    def forward(self, x):
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        x = F.log_softmax(x, dim=1)

        return x
```

```
model = FMNIST()
```



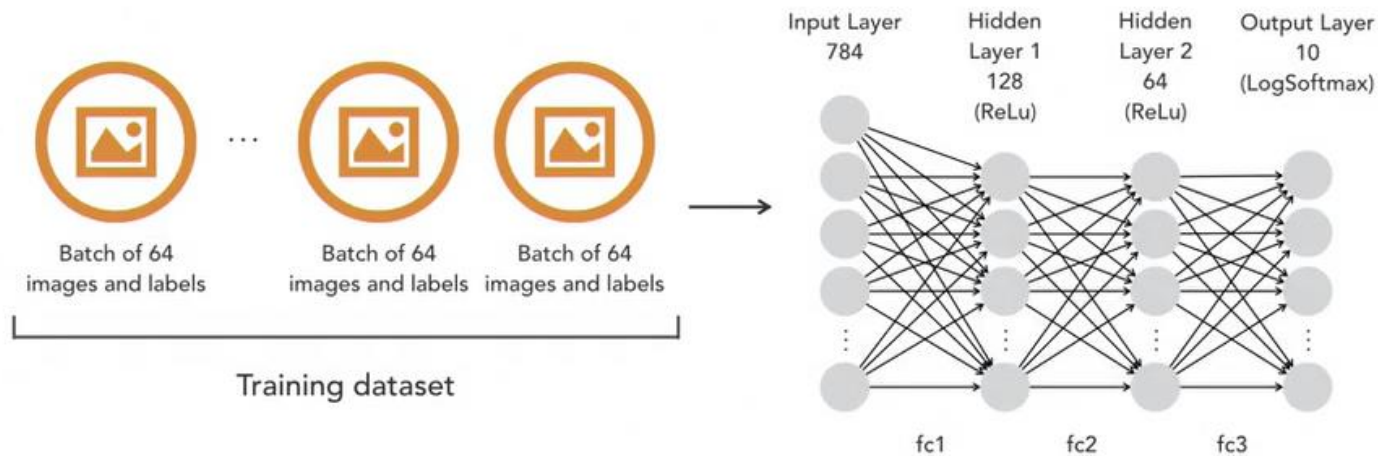
nn.Linear()



forward()

Defines the model structure, components, and order of the different layers

Training the Network



Training Steps

- Take a batch of images and targets
- Forward pass
- Calculate loss of network on batch
- Update weights of the neural network

```
from torch import optim

criterion = nn.NLLLoss()
optimizer = optim.SGD(model.parameters(), lr=0.01)

num_epochs = 3

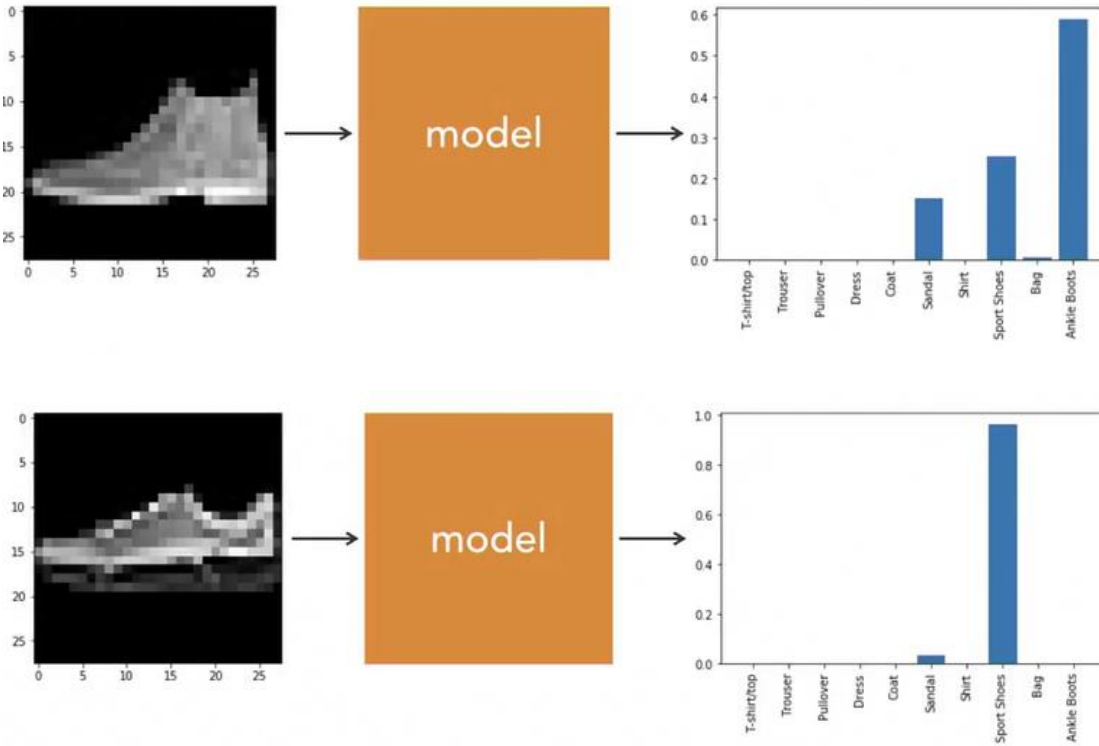
for i in range(num_epochs):
    cum_loss = 0

    for images, labels in trainloader:
        optimizer.zero_grad()
        output = model(images)
        loss = criterion(output, labels)
        loss.backward()
        optimizer.step()

        cum_loss += loss.item()

    print(f"Training loss: {cum_loss/len(trainloader)}")
```

Loss



Loss Function for Problem Types

Problem type	Last-layer activation	Loss function
Binary classification	sigmoid	Binary crossentropy
Multiclass, single-label classification	softmax	Categorical crossentropy
Multiclass, multilabel classification	sigmoid	Binary crossentropy
Regression to arbitrary values	None	MSE (mean squared error)
Regression to values between 0 and 1	sigmoid	MSE or binary crossentropy

CrossEntropyLoss

"This criterion combines `nn.LogSoftmax()` and `nn.NLLLoss()` in one single class."

CROSSENTROPYLOSS

CLASS `torch.nn.CrossEntropyLoss(weight=None, size_average=None, ignore_index=-100, reduce=None, reduction='mean')`

[SOURCE]

This criterion combines [LogSoftmax](#) and [NLLLoss](#) in one single class.

It is useful when training a classification problem with C classes. If provided, the optional argument `weight` should be a 1D *Tensor* assigning weight to each of the classes. This is particularly useful when you have an unbalanced training set.

The *input* is expected to contain raw, unnormalized scores for each class.

input has to be a *Tensor* of size either $(minibatch, C)$ or $(minibatch, C, d_1, d_2, \dots, d_K)$ with $K \geq 1$ for the K -dimensional case (described later).

This criterion expects a class index in the range $[0, C - 1]$ as the *target* for each value of a 1D tensor of size *minibatch*; if *ignore_index* is specified, this criterion also accepts this class index (this index may not necessarily be in the class range).

The loss can be described as:

$$\text{loss}(x, \text{class}) = -\log \left(\frac{\exp(x[\text{class}])}{\sum_j \exp(x[j])} \right) = -x[\text{class}] + \log \left(\sum_j \exp(x[j]) \right)$$

or in the case of the `weight` argument being specified:

$$\text{loss}(x, \text{class}) = \text{weight}[\text{class}] \left(-x[\text{class}] + \log \left(\sum_j \exp(x[j]) \right) \right)$$

The losses are averaged across observations for each minibatch. If the `weight` argument is specified then this is a weighted average:

$$\text{loss} = \frac{\sum_{i=1}^N \text{loss}(i, \text{class}[i])}{\sum_{i=1}^N \text{weight}[\text{class}[i]]}$$

Softmax

CLASS `torch.nn.Softmax(dim=None)`

[SOURCE]

Applies the Softmax function to an n-dimensional input Tensor rescaling them so that the elements of the n-dimensional output Tensor lie in the range [0,1] and sum to 1.

Softmax is defined as:

$$\text{Softmax}(x_i) = \frac{\exp(x_i)}{\sum_j \exp(x_j)}$$

Shape:

- Input: (*) where * means, any number of additional dimensions
- Output: (*), same shape as the input

Returns

a Tensor of the same dimension and shape as the input with values in the range [0, 1]

Parameters

dim (*int*) – A dimension along which Softmax will be computed (so every slice along dim will sum to 1).

• NOTE

This module doesn't work directly with NLLLoss, which expects the Log to be computed between the Softmax and itself. Use `LogSoftmax` instead (it's faster and has better numerical properties).

Examples:

```
>>> m = nn.Softmax(dim=1)
>>> input = torch.randn(2, 3)
>>> output = m(input)
```

Softmax2d

CLASS `torch.nn.Softmax2d`

[SOURCE]

Applies SoftMax over features to each spatial location.

When given an image of Channels x Height x Width, it will apply Softmax to each location (Channels, h_i , w_j)

Shape:

- Input: (N, C, H, W)
- Output: (N, C, H, W) (same shape as input)

Softmax

"Applies the Softmax function to an n-dimensional input Tensor rescaling them so that the elements of the n-dimensional output Tensor lie in the range [0,1] and sum to 1."

$$\text{Softmax}(x_i) = \frac{\exp(x_i)}{\sum_j \exp(x_j)}$$

CLASS `torch.nn.LogSoftmax(dim=None)`

[SOURCE] [↗](#)

Applies the $\log(\text{Softmax}(x))$ function to an n-dimensional input Tensor. The LogSoftmax formulation can be simplified as:

$$\text{LogSoftmax}(x_i) = \log\left(\frac{\exp(x_i)}{\sum_j \exp(x_j)}\right)$$

Shape:

- Input: (*) where * means, any number of additional dimensions
- Output: (*), same shape as the input

Parameters

dim (int) – A dimension along which LogSoftmax will be computed.

Returns

a Tensor of the same dimension and shape as the input with values in the range $[-\inf, 0)$

LogSoftmax

“Applies the $\log(\text{Softmax}(x))$ function to an n-dimensional input Tensor. The LogSoftmax formulation can be simplified as:”

$$\text{LogSoftmax}(x_i) = \log\left(\frac{\exp(x_i)}{\sum_j \exp(x_j)}\right)$$

```
[docs]def log_softmax(input, dim=None, _stacklevel=3, dtype=None):
    # type: (Tensor, Optional[int], int, Optional[int]) -> Tensor
    r"""Applies a softmax followed by a logarithm.
```

While mathematically equivalent to $\log(\text{softmax}(x))$, doing these two operations separately is slower, and numerically unstable. This function uses an alternative formulation to compute the output and gradient correctly.

See :class:`~torch.nn.LogSoftmax` for more details.

Arguments:

input (Tensor): input
dim (int): A dimension along which log_softmax will be computed.
dtype (:class:`~torch.dtype`, optional): the desired data type of returned tensor. If specified, the input tensor is casted to :attr:`dtype` before the operation is performed. This is useful for preventing data type overflows. Default: None.

"""

```
if dim is None:
    dim = _get_softmax_dim('log_softmax', input.dim(), _stacklevel)
if dtype is None:
    ret = input.log_softmax(dim)
else:
    ret = input.log_softmax(dim, dtype=dtype)
return ret
```

```
softshrink = _add_docstr(torch._C._nn.softshrink, r"""
softshrink(input, lambd=0.5) -> Tensor
```

Applies the soft shrinkage function elementwise

See :class:`~torch.nn.Softshrink` for more details.

NLLLoss (Negative Log Likelihood Loss)

```
from torch import optim

criterion = nn.NLLLoss()
optimizer = optim.SGD(model.parameters(), lr=0.01)

num_epochs = 3

for i in range(num_epochs):
    cum_loss = 0
    for images, labels in trainloader:
        images = images.view(images.shape[0], -1)

        optimizer.zero_grad()

        output = model(images)
        loss = criterion(output, labels)
        loss.backward()
        optimizer.step()

        cum_loss += loss.item()

    print(f"Training loss: {cum_loss/len(trainloader)}")
```

```
class FMNIST(nn.Module):
    def __init__(self):
        super().__init__()
        self.fc1 = nn.Linear(784, 784)
        self.fc2 = nn.Linear(784, 784)
        self.fc3 = nn.Linear(784, 10)

    def forward(self, x):
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        x = F.log_softmax(x, dim=1)

        return x

model = FMNIST()
```

CrossEntropyLoss

```
from torch import optim

criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(model.parameters(), lr=0.01)

num_epochs = 3

for i in range(num_epochs):
    cum_loss = 0
    for images, labels in trainloader:
        images = images.view(images.shape[0], -1)

        optimizer.zero_grad()

        output = model(images)
        loss = criterion(output, labels)
        loss.backward()
        optimizer.step()

        cum_loss += loss.item()

    print(f"Training loss: {cum_loss/len(trainloader)}")
```


```
class FMNIST(nn.Module):
    def __init__(self):
        super().__init__()
        self.fc1 = nn.Linear(784, 784)
        self.fc2 = nn.Linear(784, 784)
        self.fc3 = nn.Linear(784, 10)

    def forward(self, x):
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)


        return x

model = FMNIST()
```


Our goal is to try
and determine some
way to adjust the
weights so that the
loss is minimized.




Autograd



Forward Pass



Backward Pass



Why Do You Need an Optimizer?

- The loss function determines the quantity that will be minimized during training
- The optimizer determines how the network will be updated based on the loss function; it implements a variant of stochastic gradient descent (SGD)

Training Steps

Take a batch of samples and targets

Forward pass

Calculate loss

Backward pass

Update weights of the parameter

- $w = w - \text{learning rate} \times [\text{gradient of loss w.r.t. } w]$

Zero Out Gradients for Each Epoch

`optimizer.zero_grad()`

Debugging

h (help) [<i>command</i>]	print help about <i>command</i>
n (next)	execute current line of code, go to next line
c (continue)	continue executing the program until next breakpoint, exception, or end of the program
s (step into)	execute current line of code; if a function is called, follow execution inside the function
l (list)	print code around the current line
w (where)	show a trace of the function call that led to the current line
p (print)	print the value of a variable
q (quit)	leave the debugger
b (break) [<i>lineno</i> <i>function</i> [, <i>condition</i>]]	set a breakpoint at a given line number or function, stop execution there if <i>condition</i> is fulfilled
cl (clear)	clear a breakpoint
! (execute)	execute a python command
<enter>	repeat last command

CPU vs. GPU

- CPU – general purpose computing
- GPU – intensive computational calculations

Moving from CPU to GPU

1. At the top of the notebook, specify the CUDA device.
2. Transfer neural network to the GPU.
3. Send all inputs and targets to the GPU.